


HPCSE I

Monte Carlo integration

Integrating a function

- Convert the integral to a discrete sum

$$\int_a^b f(x)dx = \frac{b-a}{N} \sum_{i=1}^N f\left(a + i\frac{b-a}{N}\right) + O(1/N)$$



- Higher order integrators:

- Trapezoidal rule:

$$\int_a^b f(x)dx = \frac{b-a}{N} \left(\frac{1}{2} f(a) + \sum_{i=1}^{N-1} f\left(a + i\frac{b-a}{N}\right) + \frac{1}{2} f(b) \right) + O(1/N^2)$$

- Simpson's rule:


$$\int_a^b f(x)dx = \frac{b-a}{3N} \left(f(a) + \sum_{i=1}^{N-1} (3 - (-1)^i) f\left(a + i\frac{b-a}{N}\right) + f(b) \right) + O(1/N^4)$$

High dimensional integrals

- Simpson rule with M points per dimension
 - one dimension the error is $O(M^{-4})$
 - d dimensions we need $N = M^d$ points
the error is order $O(M^{-4}) = O(N^{-4/d})$
- An order- n scheme in 1 dimension is order- n/d in d dimensions!
- In a statistical mechanics model with N particles we have $6N$ -dimensional integrals ($3N$ positions and $3N$ momenta).
- Integration becomes extremely inefficient!


Ulam: the Monte Carlo Method

- What is the probability to win in Solitaire?
 - Ulam's answer: play it 100 times, count the number of wins and you have a pretty good estimate



Throwing stones into a pond

- How can we calculate π by throwing stones?
- Take a square surrounding the area we want to measure:



- Choose M pairs of random numbers (x, y) and count how many points (x, y) lie in the interesting area

Monte Carlo integration

- Consider an integral

$$\langle f \rangle = \int_{\Omega} f(\vec{x}) d\vec{x} / \int_{\Omega} d\vec{x}$$

- Instead of evaluating it at equally spaced points evaluate it at M points x_i chosen randomly in Ω :

$$\langle f \rangle \approx \frac{1}{M} \sum_{i=1}^M f(\vec{x}_i)$$

- The error is statistical:

$$\Delta = \sqrt{\frac{\text{Var } f}{M}} \propto M^{-1/2}$$

$$\text{Var } f = \langle f^2 \rangle - \langle f \rangle^2$$

- In $d > 8$ dimensions Monte Carlo scales better than Simpson!

Estimating the error

- The expectation value of random variables X and Y is linear:

$$E[aX + bY] = aE[X] + E[Y]$$

- The expectation value of the mean

$$\bar{X} = \frac{1}{N} \sum_{i=1}^N X_i$$

$$E[\bar{X}] = E\left[\frac{1}{N} \sum_{i=1}^N X_i\right] = \frac{1}{N} \sum_{i=1}^N E[X_i] = \frac{1}{N} \sum_{i=1}^N E[X] = E[X]$$

- \Rightarrow the mean is a “true estimator”

Estimating the error

- The sampling error is the rms (root mean square) deviation

$$\begin{aligned}(\Delta X)^2 &= E\left[(\bar{X} - E[X])^2\right] = E\left[\left(\frac{1}{N} \sum_{i=1}^N (X_i - E[X])\right)^2\right] \\&= E\left[\frac{1}{N^2} \sum_{i,j=1}^N (X_i - E[X])(X_j - E[X])\right] \\&= \frac{1}{N^2} \sum_{i,j=1}^N (E[X_i X_j] - E[X]^2) \\&= \frac{1}{N^2} \sum_{i=1}^N (E[X_i^2] - E[X]^2) = \frac{1}{N} (E[X^2] - E[X]^2) = \frac{\text{Var } X}{N}\end{aligned}$$

- We used that samples are uncorrelated:

$$E[X_i X_j] = E[X_i] E[X_j] \text{ for } i \neq j$$

Monte Carlo data analysis

- Recipe for Monte Carlo integration and data analysis of uncorrelated data:
 - Draw random points x and evaluate the function, to get random variables $X=f(x)$.
 - Store the number, sums, and sum of squares

$$N, \sum_{i=1}^N X_i, \sum_{i=1}^N X_i^2$$

- Calculate the mean as an estimate of the expectation value:

$$E[X] \approx \bar{X} = \frac{1}{N} \sum_{i=1}^N X_i$$


- Estimate the variance from the mean of squares, and from it the error

$$\text{Var } X = E[X^2] - E[X]^2 \approx \frac{N}{N-1} \left(\overline{X^2} - \bar{X}^2 \right) = \frac{N}{N-1} \left(\frac{1}{N} \sum_{i=1}^N X_i^2 - \bar{X}^2 \right)$$

$$\Rightarrow \Delta X \approx \sqrt{\frac{1}{N-1} \left(\overline{X^2} - \bar{X}^2 \right)}$$

Generating random numbers

- Real random numbers are hard to obtain
 - classical or thermal chaos (atmospheric noise)
 - quantum mechanics
- Commercial products: quantum random number generators
 - based on photons and semi-transparent mirror
 - 4 Mbit/s from a USB device, too slow for most MC simulations



<http://www.idquantique.com/>

Pseudo Random numbers

- Are generated by an algorithm
- Not random at all, but completely deterministic
- Look nearly random however when algorithm is not known and may be good enough for our purposes
- Never trust pseudo random numbers however!

Linear congruential generators

- are of the simple form $x_{n+1}=f(x_n)$
- A reasonably good choice is the GGL generator

$$x_{n+1} = (ax_n + c) \bmod m$$

with $a = 16807, c = 0, m = 2^{31}-1$

- quality is sensitive to a good choice of a, c, m
- Periodicity is a problem with such 32-bit generators
 - The sequence repeats identically after $2^{31}-1$ iterations
 - With 500 million numbers per second that is just 4 seconds!
 - **Should not be used anymore except to seed other generators**

Lagged Fibonacci generators (LFG)

$$x_n = x_{n-p} \otimes x_{n-q} \bmod M$$

- Good choices are (p,q) with $(2281,1252)$, $(9689,5502)$, $(44497,23463)$ with $M = 2^{32}$ or $M = 2^{31}-1$ and either addition or xor as operations
- Seed blocks usually generated by linear congruential generators
- Have very long periods since large block of seeds
- A very fast generator that can be parallelized!

More advanced generators

- As well-established generators fail new tests, better and better generators get developed
- Mersenne twister (Matsumoto & Nishimura, 1997)
<http://www.math.sci.hiroshima-u.ac.jp/~m-mat/MT/emt.html>
- Well generator (Panneton and L'Ecuyer, 2004)
<http://www.iro.umontreal.ca/~panneton/WELLRNG.html>
- However, they cannot be easily parallelized or vectorized

Creating independent sequences

- **Leapfrog:** elements of the sequence are given in a cyclic fashion. For n cores, element i is given to core $i \bmod n$.
- **Sequential splitting:** the cores use non-overlapping sequential blocks.
- **Independent sequences:** the lagged Fibonacci generators can produce independent streams, depending on the seed.
 - See the SPRNG library and papers by Mascagni for seeding the lagged Fibonacci generator
- **Stochastic seeding:** use a good generator to fill all n seed blocks and hope for the best

Independent sequences in the LFG

- With $M=2^m$, the LFG has a maximum cycle length of $(2^q-1)2^{m-1}$
- This is less than the theoretically possible cycle length of $2^{mq} - 1$
- There exist $2^{(q-1)(m-1)}$ cycles of maximal length
- We assign each thread to a different cycle by choosing the seed block in a smart way. This seeding is implemented in the SPRNG library (<http://sprng.cs.fsu.edu/>)
- The Mersenne twister and WELL generator are smarter: they have the maximum possible cycle length. The disadvantage is that no independent sequences can be generated.

Are these numbers really random?

- No!
- Are they random enough?
 - Maybe?
- Statistical tests for distribution and correlations
- Are these tests enough?
 - No! Your calculation could depend in a subtle way on hidden correlations!
- What is the ultimate test?
 - Run your simulation with various random number generators and compare the results

Marsaglia's diehard tests

- **Birthday spacings:** Choose random points on a large interval. The spacings between the points should be asymptotically Poisson distributed. The name is based on the birthday paradox.
- **Overlapping permutations:** Analyze sequences of five consecutive random numbers. The 120 possible orderings should occur with statistically equal probability.
- **Ranks of matrices:** Select some number of bits from some number of random numbers to form a matrix over $\{0,1\}$, then determine the rank of the matrix. Count the ranks.
- **Monkey tests:** Treat sequences of some number of bits as "words". Count the overlapping words in a stream. The number of "words" that don't appear should follow a known distribution. The name is based on the infinite monkey theorem.
- **Count the 1s:** Count the 1 bits in each of either successive or chosen bytes. Convert the counts to "letters", and count the occurrences of five-letter "words".
- **Parking lot test:** Randomly place unit circles in a 100 x 100 square. If the circle overlaps an existing one, try again. After 12,000 tries, the number of successfully "parked" circles should follow a certain normal distribution.

Marsaglia's diehard tests (cont.)

- **Minimum distance test:** Randomly place 8,000 points in a 10,000 x 10,000 square, then find the minimum distance between the pairs. The square of this distance should be exponentially distributed with a certain mean.
- **Random spheres test:** Randomly choose 4,000 points in a cube of edge 1,000. Center a sphere on each point, whose radius is the minimum distance to another point. The smallest sphere's volume should be exponentially distributed with a certain mean.
- **The squeeze test:** Multiply 2^{31} by random floats on $[0,1)$ until you reach 1. Repeat this 100,000 times. The number of floats needed to reach 1 should follow a certain distribution.
- **Overlapping sums test:** Generate a long sequence of random floats on $[0,1)$. Add sequences of 100 consecutive floats. The sums should be normally distributed with characteristic mean and sigma.
- **Runs test:** Generate a long sequence of random floats on $[0,1)$. Count ascending and descending runs. The counts should follow a certain distribution.
- **The craps test:** Play 200,000 games of craps, counting the wins and the number of throws per game. Each count should follow a certain distribution.


Non-uniform random numbers

- we found ways to generate pseudo random numbers u in the interval $[0,1[$
- How do we get other uniform distributions?
 - uniform x in $[a,b[$: $x = a + (b-a) u$
- Other distributions:
 - Inversion of integrated distribution
 - Rejection method

Non-uniform distributions

- How can we get a random number x distributed with $f(x)$ in the interval $[a, b[$ from a uniform random number u ?
- Look at probabilities:

$$P[x < y] = \int_a^y f(x) dx =: F(y) \equiv P[u < F(y)]$$
$$\Rightarrow x = F^{-1}(u)$$



- This method is feasible if the integral can be inverted easily
 - exponential distribution $f(x) = \lambda \exp(-\lambda x)$
 - can be obtained from uniform by $x = -1/\lambda \ln(1-u)$


Normally distributed numbers

- The normal distribution $f(x) = \frac{1}{\sqrt{2\pi}} \exp(-x^2)$
- cannot easily be integrated in one dimension but can be easily integrated in 2 dimensions!
- We can obtain two normally distributed numbers from two uniform ones (Box-Muller method)

$$n_1 = \sqrt{-2 \ln(1 - u_1)} \sin(2\pi u_2)$$


$$n_2 = \sqrt{-2 \ln(1 - u_1)} \cos(2\pi u_2)$$

Rejection method (von Neumann)




- Look for a simple distribution h that bounds $f: f(x) < \lambda h(x)$
 - Choose an h -distributed number x
 - Choose a uniform random number $0 \leq u < 1$
 - Accept x if $u < f(x) / \lambda h(x)$,
otherwise reject x and get a new pair (x, u)
- Needs a good guess h to be efficient

Sharply peaked functions



- In many cases a function is large only in a tiny region
- Lots of time wasted in regions where the function is small
- The sampling error is large since the variance is large

Importance sampling



- Choose points not uniformly but with probability $p(x)$:

$$\langle f \rangle = \left\langle \frac{f}{p} \right\rangle_p := \int_{\Omega} \frac{f(\vec{x})}{p(\vec{x})} p(\vec{x}) d\vec{x} \bigg/ \int_{\Omega} d\vec{x}$$

- The error is now determined by $\text{Var } f/p$
- Find p similar to f and such that p -distributed random numbers are easily available