

The background of the slide features a solid blue horizontal band that spans the entire width of the image, positioned in the middle. Above and below this band are light gray areas.

**HPCSE I**

**Multithreading**



# Threads and Processes

- Process – execution sequence within the OS, i.e. a program
  - Relatively expensive to create
  - Independent resources, state (by default)
  - Immune to many concurrency issues
- Thread – execution sequence within the process
  - `main()` is the first thread
  - Cheap to create
  - Shared resources, state
  - Difficult to use correctly




# Concurrency Tradeoffs

- Pros
  - Compute faster
  - “Unblocking” – get work done while waiting for events to occur outside the CPU
- Cons
  - Synchronization overhead
  - Programming discipline – problems abound!
  - Harder to reason about
  - Harder to debug
- Don't use threads unless you can't avoid it!
- However, we can't avoid it and thus have to learn it.



# Spawning and joining threads

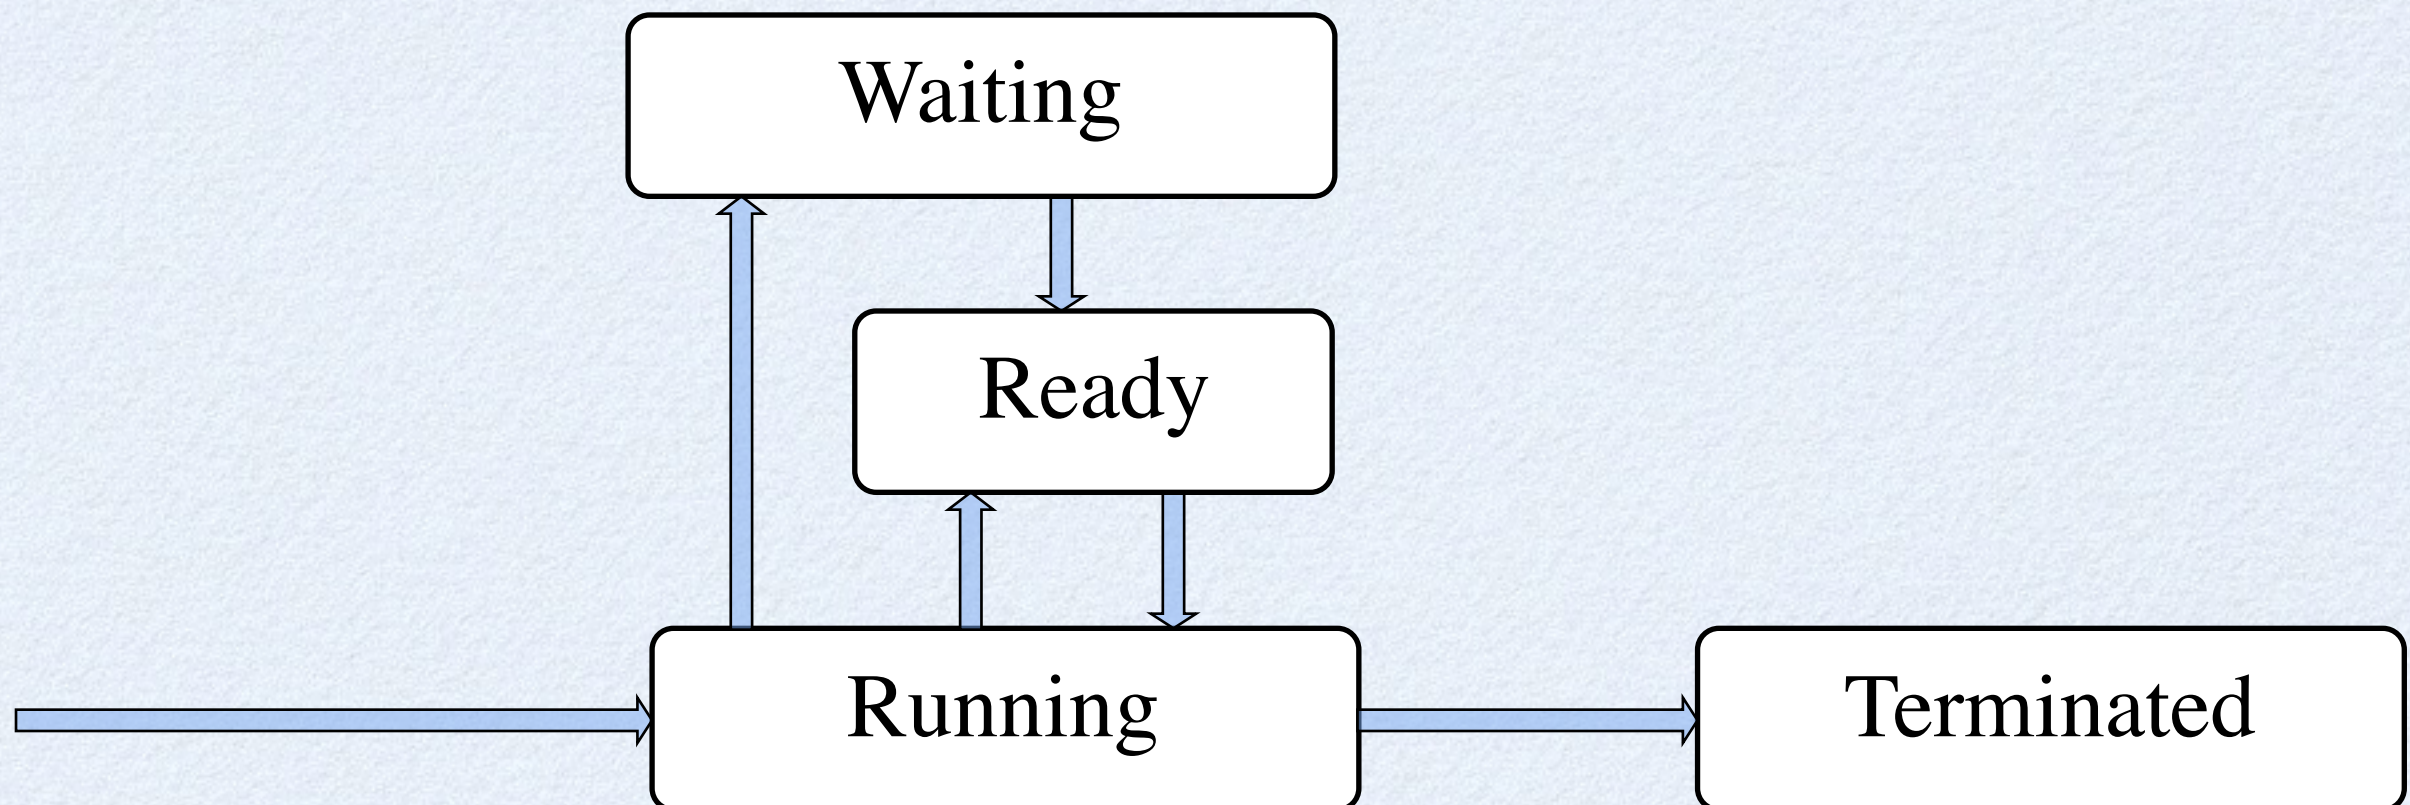
- During execution of a multithreaded program threads get spawned and joined dynamically:





# Thread States

- **Running** – Currently executing.  $\# \text{running threads} \leq \# \text{CPU cores}$
- **Ready** – Prepared to run whenever processor time can be allocated to it. Not waiting.
- **Blocked** (or waiting) – Paused until some resource (other than processor) is allocated to it.
- **Terminated** – Finished execution but OS resources not yet deallocated.





# Threading libraries

- Most operating systems have some kind of native threading library (pthreads on Unix, Win32 threads, ....). To achieve portable codes we want to use a cross-platform standard
- The **C++11** standard contains a cross-platform threading library
  - supported by g++-4.7, clang++-4.0 and MSVC11
- On other compilers
  - Use **Boost.Threads**, the predecessor to C++11 threads on many compilers
  - **Intel Thread Building Blocks (TBB)** with Intel C++, based on draft version of the C++11 standard with its own extensions
- We will use the **C++11** subset supported by most modern compilers



# Compiling C++11 codes

- With **g++** use `-pthread` or you will get runtime errors:
  - **g++ -std=c++11 -pthread ...**
- With **clang** specify that we want to use the clang standard library and not the gnu version:
  - For clang 4.0: **clang++ -std=c++11 -stdlib=libc++ ...**
  - For clang 3.x: **clang++ -std=c++0x -stdlib=libc++ ...**
- Does anyone want to use MSVC compilers on Windows?
- Example sources are available by git:
  - `git clone https://gitlab.phys.ethz.ch/hpcse\_hs14/lecture.git`



# Launching and joining threads

- A thread is launched by passing a function (and optionally its arguments) to the thread constructor:

```
std::thread t (foo, arg1);
```

- We can also use C++11 lambda functions:

```
std::thread t ([] () { std::cout << "Hello world!\n"; });
```

- Threads are joined calling the join function:

```
t.join();
```



# std::thread (abridged)

```
class thread
{
public:
    thread() noexcept;
    ~thread();

    void swap(thread& x);
    thread& operator=(thread&&) noexcept;
    // ...move support but noncopyable...

    typedef platform-specific-type
    native_handle_type;
    native_handle_type native_handle();
```

```
    // launch
    template <class F>
    explicit thread(F f);

    template <class F, class A1, class A2, ...>
    thread(F f, A1 a1, A2 a2, ...);

    void join();

    bool joinable() const; // still attached?
    void detach();

    static unsigned hardware_concurrency();

    class id;
    id get_id() const noexcept;

    static void yield();
    static void sleep(const system_time& xt);
};
```



# Movable/Noncopyable Types

- Can't copy/assign from lvalues
  - `std::thread x, y; x = y; // error!`
- Can place in C++11 containers...
  - `std::vector<std::thread> v(10); // ok!`
- Can “copy”/assign from rvalues
  - `pool[3] = std::thread(f);`
- Can pass to/return from functions
  - `std::thread t = make_thread();`
  - `do_something(make_thread());`
- Can swap
  - `x.swap(y);`
  - `swap(x, y);`



# Detaching and destroying threads

- The detach() member function let's the thread run on, but the object no longer refers to it.
- Destroying running threads is different in C++11 and Boost:
  - **Boost** silently detaches a joinable (still running) thread
  - **C++11** terminates the program if the thread is joinable
  - The reason is that a detached thread might be a security hole or bug



# Example: Integration using Simpson's rule

## simpson.hpp

```
#include <cassert>
#include <functional>

inline double
simpson(std::function<double(double)> f,
        double a, double b, unsigned int N)
{
    assert (b>=a);
    assert (N!=0u);

    double h=(b-a)/N;

    // boundary values
    double result = ( f(a) + 4*f(a+h/2) +
                     f(b) ) / 2.0;

    // values between boundaries
    for ( unsigned int i = 1; i <= N-1; ++i )
        result += f(a+i*h) + 2*f(a+(i+0.5)*h);

    return result * h / 3.0;
}
```

## simpson\_serial.cpp

```
#include "simpson.hpp"
#include <cmath>
#include <iostream>

double func(double x)
{
    return x * std::sin(x);
}

int main()
{
    double a;    // lower bound of integration
    double b;    // upper bound of integration
    unsigned int nsteps; // number of subintervals

    // read the parameters
    std::cin >> a >> b >> nsteps;

    // print the result
    std::cout << simpson(func,a,b,nsteps)
               << std::endl;

    return 0;
}
```



# Simpson's rules using two threads

simpson\_threaded1.cpp

```
#include "simpson.hpp"
#include <cmath>
#include <iostream>
#include <thread>

double func(double x) { return x * std::sin(x); }

int main()
{
    double a;           // lower bound of integration
    double b;           // upper bound of integration
    unsigned int nsteps; // number of subintervals for integration

    std::cin >> a >> b >> nsteps;

    double result1; // the integral of the first half

    // spawn a thread for the first half of the interval
    std::thread t( [&] () { result1 = simpson(func,a,a+(b-a)/2.,nsteps/2);} );

    // locally integrate the second half
    double result2 = simpson(func,a+(b-a)/2.,b,nsteps/2);

    t.join(); // wait for the thread to join
    std::cout << result1 + result2 << std::endl;

    return 0;
}
```



# Futures

- This worked but was clumsy:
  - we needed to declare a variable to hold the return type
  - we needed to create a (lambda) function to fill it
- Futures hold future return values of a function called asynchronously in a thread.
  - we use the future to specify the return value
  - we can access its value after the asynchronous call finishes



# Simpson's rules using a future

simpson\_threaded2.cpp

```
#include "simpson.hpp"
#include <cmath>
#include <iostream>
#include <thread>
#include <future>

double func(double x) { return x * std::sin(x); }

int main()
{
    double a;           // lower bound of integration
    double b;           // upper bound of integration
    unsigned int nsteps; // number of subintervals for integration
    std::cin >> a >> b >> nsteps;

    // create a packaged task
    std::packaged_task<double> pt(std::bind(simpson, func, a, a+(b-a)/2., nsteps/2));
    std::future<double> fi = pt.get_future(); // get the future return value
    std::thread t (std::move(pt));          // launch the thread

    double result2 = simpson(func, a+(b-a)/2., b, nsteps/2);

    fi.wait(); // wait for the task to finish and the future to be ready

    std::cout << result2 + fi.get() << std::endl;
    t.join();

    return 0;
}
```



# Simpson's rules using a future

simpson\_threaded2.cpp

```
#include "simpson.hpp"
#include <cmath>
#include <iostream>
#include <thread>
#include <future>

double func(double x) { return x * std::sin(x); }

int main()
{
    double a;           // lower bound of integration
    double b;           // upper bound of integration
    unsigned int nsteps; // number of subintervals for integration
    std::cin >> a >> b >> nsteps;

    // create a packaged task
    std::packaged_task<double> pt(std::bind(simpson, func, a, a+(b-a)/2., nsteps/2));
    std::future<double> fi = pt.get_future(); // get the future return value
    std::thread t (std::move(pt));          // launch the thread

    double result2 = simpson(func, a+(b-a)/2., b, nsteps/2);

    std::cout << result2 + fi.get() << std::endl; // fi.get() will wait for the result

    t.join();

    return 0;
}
```



# Asynchronous function calls

- even simpler are explicit asynchronous calls

```
std::future<double> fi = std::async(simpson,a,b,n);  
  
std::cout << fi.get() << std::endl; // get blocks automatically
```

- but this might or might not run in a new thread.
- better might be to force an asynchronous call:

```
std::future<double> fi=std::async(std::launch::async, simpson,a,b,n);  
  
std::cout << fi.get() << std::endl;
```



# Simpson's rules using asynchronous calls

simpson\_threaded4.cpp

```
#include "simpson.hpp"
#include <cmath>
#include <iostream>
#include <thread>
#include <future>

double func(double x) { return x * std::sin(x); }

int main()
{
    double a;           // lower bound of integration
    double b;           // upper bound of integration
    unsigned int nsteps; // number of subintervals for integration
    std::cin >> a >> b >> nsteps;

    // even easier: launch an asynchronous function call
    // force it to be asynchronous and thus in a separate thread
    std::future<double> fi = std::async(std::launch::async, simpson, func, a, a+(b-a)/2., nsteps/2);

    // locally integrate the second half
    double result = simpson(func, a+(b-a)/2., b, nsteps/2);

    std::cout << result + fi.get() << std::endl;

    return 0;
}
```



# The running thread in C++11

- Information about the thread is in the namespace `std::this_thread`:

```
namespace std {  
    namespace this_thread {  
        thread::id get_id() noexcept;  
  
        void yield() noexcept;  
  
        template <class Clock, class Duration>  
        void sleep_until(const chrono::time_point<Clock, Duration> t) noexcept;  
  
        template <class Rep, class Period>  
        void sleep_for(const chrono::duration<Rep, Period>& t) noexcept;  
    }  
}
```



# Calculating $\pi$ through a series

```
#include <vector>
#include <iostream>
#include <thread>
#include <numeric>
#include <iomanip>

// sum terms [i-j) of the power series for
// pi/4
void sumterms(long double& sum,
              std::size_t i, std::size_t j)
{
    sum = 0.0;

    for (std::size_t t = i; t < j; ++t)
        sum += (1.0 - 2* (t % 2)) / (2*t + 1);
}
```

```
int main()
{
    // decide how many threads to use
    std::size_t const nthreads = std::max(1u,
                                           std::thread::hardware_concurrency());

    std::vector<std::thread> threads(nthreads);
    std::vector<long double> results(nthreads);

    unsigned long const nterms = 100000000;
    long double const step = (nterms+0.5l) / nthreads;

    for (unsigned i = 0; i < nthreads; ++i)
        threads[i] = std::thread(
            sumterms, std::ref(results[i]),
            i * step, (i+1) * step
        );

    for (std::thread& t : threads)
        t.join();

    long double pi = 4 * std::accumulate(
        results.begin(), results.end(), 0.);

    std::cout << "pi=" << std::setprecision(18)
              << pi << std::endl;

    return 0;
}
```

$$\arctan 1 = 1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \cdots = \frac{\pi}{4}$$



# But why keep so many return values?

```
#include <vector>
#include <iostream>
#include <thread>
#include <numeric>
#include <iomanip>

// sum terms [i-j) of the power series for
// pi/4
void sumterms(long double& sum,
              std::size_t i, std::size_t j)
{
    for (std::size_t t = i; t < j; ++t)
        sum += (1.0 - 2* (t % 2)) / (2*t + 1);
}
```

```
int main()
{
    // decide how many threads to use
    std::size_t const nthreads = std::max(1u,
                                           std::thread::hardware_concurrency());

    std::vector<std::thread> threads(nthreads);
    // let us just use a single result
    long double result=0.;

    unsigned long const nterms = 100000000;
    long double const step = (nterms+0.5l) / nthreads;

    for (unsigned i = 0; i < nthreads; ++i)
        threads[i] =std::thread(
            sumterms, std::ref(result),
            i * step, (i+1) * step
        );

    for (std::thread& t : threads)
        t.join();

    std::cout << "pi=" << std::setprecision(18)
              << 4.*result << std::endl;

    return 0;
}
```

Do you see a problem?



# Threading and race conditions

- In threaded programs, we must stop other threads from looking (or touching) data that we need. Requires cooperation!
- A concurrent program that doesn't control visibility of broken invariants has a **race condition**
- From the point of view of threading, even an int has an invariant that is broken during mutation



# Thread Safety: Serializing Access

- Basic mechanism: **mutex**
- Associated with some shared mutable data, which may be
  - as small as a char
  - as large as a list<vector<string> > (or larger)
- At any time, a mutex is either **locked by one thread** or **unlocked**.
- When a thread asks to lock a mutex
  - If the mutex is unlocked, it becomes locked and the thread proceeds
  - If the mutex is locked, the thread is blocked until the lock is released and reallocated to the locking thread.
- Protocol – threads agree to:
  - acquire a lock on the mutex before accessing the data
  - release the lock when done accessing the data



# Locks

- Movable/noncopyable. Expresses ownership of a thread
- Forgetting to unlock a mutex will cause the next thread that locks it to wait forever
- Use RAI (resource acquisition is initialization) lock objects to eliminate this problem:
  - Constructor locks (acquires) the mutex
  - Destructor unlocks (releases) it
  - Very safe!
- Note: one lock object should never be accessed by multiple threads!



# Safety through mutex and lock\_guard

```
#include <vector>
#include <iostream>
#include <thread>
#include <numeric>
#include <iomanip>

// sum terms [i-j) of the power series for
// pi/4
void sumterms(
    std::pair<long double, std::mutex>& result,
    std::size_t i, std::size_t j
)
{
    long double sum=0.;

    for (std::size_t t = i; t < j; ++t)
        sum += (1.0 - 2* (t % 2)) / (2*t + 1);

    std::lock_guard<std::mutex> l (result.second);
    result.first += sum;
}
```

```
int main()
{
    // decide how many threads to use
    std::size_t const nthreads = std::max(1u,
        std::thread::hardware_concurrency());

    std::vector<std::thread> threads(nthreads);
    // let us just use a single result
    std::pair<long double, std::mutex> result;
    result.first = 0.;

    unsigned long const nterms = 100000000;
    long double const step = (nterms+0.5l) / nthreads;

    for (unsigned i = 0; i < nthreads; ++i)
        threads[i] = std::thread(
            sumterms, std::ref(result),
            i * step, (i+1) * step
        );

    for (std::thread& t : threads)
        t.join();

    // no need to lock here
    std::cout << "pi=" << std::setprecision(18)
        << 4.*result << std::endl;

    return 0;
}
```

Now we are safe



# Example: garbled I/O

garbledio.cpp

```
#include <iostream>
#include <thread>
#include <vector>

void printer( int n )
{
    for ( int i = 0; i < 100; ++i)
        std::cout << "do not garble thread " << n << ": " << i << std::endl;
}

int main()
{
    std::vector<std::thread> threads;

    for (int n = 1; n < 10; ++n)
        threads.push_back(std::thread(printer, n));

    for (std::thread& t : threads)
        t.join();
}
```



# Example: synchronized I/O

## syncedio.cpp

```
#include <iostream>
#include <thread>
#include <vector>

std::mutex io_mutex; // global

struct sync
{
    sync( std::ostream& os )
    : os(os)
    , lock(io_mutex) {}

    template <class T>
    std::ostream& operator<<(T const& x)
    {
        return os << x;
    }

private:
    std::ostream& os;
    std::lock_guard<std::mutex> lock;
};
```

```
void printer( int n )
{
    for ( int i = 0; i < 100; ++i) {
        sync(std::cout)
        << "do not garble thread "
        << n << ": " << i << std::endl;
    }
}

int main()
{
    std::vector<std::thread> threads;

    for (int n = 1; n < 10; ++n)
        threads.push_back(std::thread(printer, n));

    for (std::thread& t : threads)
        t.join();
}
```



# Mutexes and Locks

- We have four basic mutex types
  - **mutex**
  - **recursive\_mutex**: allows multiple locking by the same thread
  - **timed\_mutex**: allows time-outs in lock attempts
  - **recursive\_timed\_mutex**: both of the above
  - We need to use a timed mutex for timed locks
- Lock types:
  - **lock\_guard**
  - **unique\_lock**



# unique\_lock

- The unique\_lock is more flexible and allows deferring the lock
  - `unique_lock<mutex> l(m);` // locks the lock
  - `unique_lock<mutex> l(m, std::adopt_lock);` // adopts the lock state
  - `unique_lock<mutex> l(m, std::defer_lock);` // does not lock yet
  - `unique_lock<mutex> l(m, std::try_to_lock);` // tries to lock
  - `unique_lock<mutex> l(m, abs_time);` // tries to lock, with timeout
- And it has some important functions:
  - `l.owns_lock();` // returns whether it is locked
  - `if (l) ...` // tests whether locked



# unique\_lock (continued)

- It can be locked later:
  - `l.try_lock();` // tries to lock and returns whether is succeeded
  - `l.try_lock_for(rel_time);` // tries to lock with timeout
  - `l.try_lock_until(abs_time);` // tries to lock with timeout
  - `l.lock();` // locks the lock
  - `l.unlock();`
  - `std::lock(l1,l2); std::lock(l1,l2,l3); ...` // lock multiple locks at the same time
- The timed locks and time specification are slightly different in Boost.Thread .



# Example: Pairwise Associations

```
struct collab
{
    collab() : partner(0) {}
    ~collab() { decouple(); }

    void couple(collab* new_partner);
    void decouple();

private:
    collab* partner;
    std::mutex gate;
};

typedef std::lock_guard<std::mutex> guard;

struct lock2
{
    lock2(std::mutex& a, std::mutex& b)
        : l0( a ),
          l1( b )
    {}
    guard l0, l1;
};
```

```
void collab::couple(collab* other)
{
    decouple();
    other->decouple();
    lock2 g(gate, other->gate);
    if (partner || other->partner) return;
    partner = other;
    other->partner = this;
}

void collab::decouple()
{
    collab* cur;
    {
        guard g0(gate);
        cur = partner;
        if (!cur) return;
    }

    lock2 g(gate, cur->gate);
    if (partner != cur) return;
    partner = 0;
    cur->partner = 0;
}
```



# Deadlock: The Deadly Embrace

- Once you have synchronization, you can also have **deadlock**
- Scenario:
  - Mutexes 1 and 2, unlocked
  - Thread A locks mutex 1                      A still running
  - Thread B locks mutex 2                      B still running
  - Thread A locks mutex 2                      A waits (for B)
  - Thread B locks mutex 1                      B waits (for A)
- Solution:
  - We need to lock both in the same order
  - Introduce an ordering on mutexes, e.g. by address
  - Easier solution: use **std::lock()** function with multiple mutexes



# We need to lock both at the same time

```
struct collab
{
    collab() : partner(0) {}
    ~collab() { decouple(); }
    void couple(collab* new_partner);
    void decouple();
private:
    collab* partner;
    std::mutex gate;
};

typedef std::unique_lock<std::mutex> guard;
```

```
void collab::couple(collab* other)
{
    decouple();
    other->decouple();
    guard g1(gate,defer_lock);
    guard g2(other->gate,defer_lock);
    std::lock(g1,g2); // lock both simultaneously
    if (partner || other->partner) return;
    partner = other;
    other->partner = this;
}

void collab::decouple()
{
    collab* cur;
    {
        guard g0(gate);
        cur = partner;
        if (!cur) return;
    }

    guard g1(gate,defer_lock());
    guard g2(cur->gate,defer_lock());
    std::lock(g1,g2); // lock both simultaneously
    if (partner != cur) return;
    partner = 0;
    cur->partner = 0;
}
```



# std::call\_once

- “Once routines”
  - Executed once, no matter how many invocations
  - No invocation will complete until the one execution finishes
  - Typical use: initialization of static and function-static data
- Protocol:
  - Declare a global (namespace scope) `once_flag` for each once routine
  - Invoke the once routine indirectly by passing its address and `once_flag` to `call_once`.

```
std::once_flag printonce_flag;

void printonce() { std::cout << "This should be printed only once\n"; }

int main()
{
    std::vector<std::thread> threads;
    for (int n = 0; n < 10; ++n)
        threads.push_back(
            std::thread( [&]() {std::call_once(printonce_flag, printonce);} ));

    for (std::thread& t : threads)
        t.join();

    return 0;
}
```