

OpenMP and MPI: Summary

High Performance Computing for Science and Engineering I

December 12, 2014

Sources

- OpenMP specifications at www.openmp.org
 - OpenMP 3.1 (2011): C/C++, Fortran and Examples
- OpenMP tutorial:
 - <https://computing.llnl.gov/tutorials/openMP/>
- MPI tutorial:
 - <https://computing.llnl.gov/tutorials/mpi/>

OpenMP

- Compilation
- OpenMP function calls
- OpenMP environment variables
- OpenMP directives (pragmas)

Compilation

- Compile and link using the **-fopenmp** option on the GNU compiler:
 - `g++ -fopenmp openmp1.cpp`
- You will need to adapt a Makefile in order to compile your code
- Useful flag: **-Wall** (enables all the warnings)
 - `g++ -Wall -fopenmp openmp1.cpp`
- We might need to add C++11 support
 - `g++ -Wall --std=c++11 -fopenmp openmp1.cpp`

OpenMP function calls

- Do not forget to include `<omp.h>`
 - All functions defined in this header file
- List of functions
 - `int omp_get_thread_num()`
 - `int omp_get_num_threads()`
 - `void omp_set_dynamic()`
 - `int omp_get_dynamic()`
 - `void omp_set_nested()`
 - `int omp_get_nested()`
 - `double omp_get_wtime()`

OpenMP function calls

<code>void omp_set_num_threads(int n)</code>	Sets the number of threads to be used.
<code>int omp_get_num_threads()</code>	Gets the number of currently running threads.
<code>int omp_get_max_threads()</code>	Gets the maximum number of threads that can be used for one parallel region.
<code>int omp_get_thread_num()</code>	Get the id of the calling thread
<code>int omp_get_thread_limit ()</code>	Gets the maximum number of threads used for nested parallel region.
<code>int omp_get_num_procs()</code>	Gets the number of processors available
<code>int omp_in_parallel()</code>	Returns true if called from within a parallel region
<code>void omp_set_dynamic(int n)</code>	Sets dynamic adjustment of threads with the given number as maximum. This overrides the environment variable.
<code>int omp_get_dynamic()</code>	Returns true if dynamic scheduling is enables
<code>double omp_get_wtime()</code>	A portable wallclock timing routine, returns time in seconds. The time is not synchronized across threads to be fast.
<code>double omp_get_wtick()</code>	Returns the number of seconds between successive clock ticks

OpenMP environment variables

- OpenMP provides the following environment variables for controlling the execution of parallel code
 - OMP_NUM_THREADS: max threads to use during execution
 - OMP_PROC_BIND: thread binding to cores
 - OMP_DYNAMIC: dynamic adjustment of number of threads
 - OMP_NESTED: support of nested parallelism
- Examples
 - `export OMP_NUM_THREADS=4`
 - `export OMP_DYNAMIC=FALSE`

OpenMP directives

- Parallel regions
 - `#pragma omp parallel`
- Synchronization
 - `#pragma omp master`
 - `#pragma omp single`
 - `#pragma omp critical`
 - `#pragma omp barrier`
- Work-sharing
 - `#pragma omp section`
 - `#pragma omp for`

omp parallel

```
#include <omp.h>

main () {

    int nthreads, tid;

    /* Fork a team of threads with each thread having a private tid variable */
    #pragma omp parallel private(tid)
    {

        /* Obtain and print thread id */
        tid = omp_get_thread_num();
        printf("Hello World from thread = %d\n", tid);

        /* Only master thread does this */
        if (tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }

    } /* All threads join master thread and terminate */

}
```

omp parallel

```
int A, B, C;  
A = B = C = 1;  
#pragma omp parallel private(B) firstprivate(C)  
{  
    // code  
}
```

Inside the parallel region

- A is shared between threads and equal to 1
- B, C are private to each thread
- B is not initialized
- C is equal to 1

After the parallel region

- The values of B and C cannot be determined

omp master

```
#include <omp.h>

main () {

    int nthreads, tid;

    /* Fork a team of threads with each thread having a private tid variable */
    #pragma omp parallel private(tid)
    {

        /* Obtain and print thread id */
        tid = omp_get_thread_num();
        printf("Hello World from thread = %d\n", tid);

        /* Only master thread does this */
        #pragma omp master
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }

    } /* All threads join master thread and terminate */

}
```

omp single

```
#include <omp.h>

main () {

    int nthreads, tid;

    /* Fork a team of threads with each thread having a private tid variable */
    #pragma omp parallel private(tid)
    {

        /* Obtain and print thread id */
        tid = omp_get_thread_num();
        printf("Hello World from thread = %d\n", tid);

        /* Only one thread does this */
        #pragma omp single
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }

    } /* All threads join master thread and terminate */

}
```


omp critical

```
#include <omp.h>

main()
{

    int x;
    x = 0;

    #pragma omp parallel shared(x)
    {

        #pragma omp critical
        x = x + 1;

    } /* end of parallel section */

}
```

omp critical (2)

```
#include <omp.h>
```

```
int x; // what is the difference if I put it here?
```

```
main()  
{  
  //int x;  
  x = 0;
```

```
#pragma omp parallel shared(x)  
  {  
  
    #pragma omp critical  
    x = x + 1;  
  
  } /* end of parallel section */  
}
```


OpenMP barrier

- OpenMP 3.1 specs, Section 2.8.3, pages 70-71:
- Summary
 - The barrier construct specifies an explicit barrier at the point at which the construct.
 - The barrier directive may not be used in place of the statement following an if, while, do, switch, or label.
- Description ...
- Restrictions ...

omp barrier + single vs master

```
#pragma omp parallel  
{
```

```
    do_many_things();
```

```
    #pragma omp single
```

```
    {
```

```
        exchange_boundaries();
```

```
    }
```

```
    do_many_other_things();
```

```
}
```

only one thread enters

implicit barrier



```
#pragma omp parallel  
{
```

```
    do_many_things();
```

```
    #pragma omp master
```

```
    {
```

```
        exchange_boundaries();
```

```
    }
```

```
    #pragma barrier
```

```
    do_many_other_things();
```

```
}
```

only thread 0 enters

omp for

```
#include <omp.h>
#define CHUNKSIZE 100
#define N      1000

main ()
{

    int i, chunk;
    float a[N], b[N], c[N];

    /* Some initializations */
    for (i=0; i < N; i++)
        a[i] = b[i] = i * 1.0;

    chunk = CHUNKSIZE;

    #pragma omp parallel shared(a,b,c,chunk) private(i)
    {

        #pragma omp for schedule(dynamic,chunk) nowait
        for (i=0; i < N; i++)
            c[i] = a[i] + b[i];

    } /* end of parallel section */

}
```

omp sections

```
include <omp.h>
#define N      1000

main ()
{
int i;
float a[N], b[N], c[N], d[N];

/* Some initializations */
for (i=0; i < N; i++) { a[i] = i * 1.5; b[i] = i + 22.35; }

#pragma omp parallel shared(a,b,c,d) private(i)
{
    #pragma omp sections nowait
    {

        #pragma omp section
        for (i=0; i < N; i++)
            c[i] = a[i] + b[i];

        #pragma omp section
        for (i=0; i < N; i++)
            d[i] = a[i] * b[i];

    } /* end of sections */

} /* end of parallel section */
}
```

omp sections (1/2)

```
//sequential code
V = alpha();
W = beta();
X = gamma(V, W);
Y = delta();
printf("%f\n", epsilon(X,Y));
```

```
#pragma omp parallel sections
{
    #pragma omp section
    V = alpha();

    #pragma omp section
    W = beta();

    #pragma omp section
    Y = delta();
}
X = gamma(V, W);
printf("%f\n", epsilon(X,Y));
```

let's assume 2 threads

omp sections (2/2)

```
#pragma omp parallel
{

    #pragma omp sections
    {
        #pragma omp section
        V = alpha();

        #pragma omp section
        W = beta();
    }

    #pragma omp sections
    {
        #pragma omp section
        X = gamma(V, W);

        #pragma omp section
        Y = delta();
    }
}
printf("%f\n", epsilon(X,Y));
```

omp sections

```
void XAXIS();  
void YAXIS();  
void ZAXIS();
```

```
void a9()  
{  
    #pragma omp parallel sections  
    {  
        #pragma omp section  
        XAXIS();  
        #pragma omp section  
        YAXIS();  
        #pragma omp section  
        ZAXIS();  
    }  
}
```

- Quiz: What if sections is not available?

omp reduction

```
#include <omp.h>
```

```
main () {  
    int n, chunk;  
    float a[100], b[100], result;
```

```
    /* Some initializations */  
    n = 100;  
    chunk = 10;  
    result = 0.0;  
    for (int i=0; i < n; i++)  
    {  
        a[i] = i * 1.0;  
        b[i] = i * 2.0;  
    }
```

```
    #pragma omp parallel for schedule(static,chunk) reduction(+:result)  
    for (int i=0; i < n; i++)  
        result = result + (a[i] * b[i]);
```

```
    printf("Final result= %f\n",result);
```

```
}
```

Be careful with the reductions

- avoid false sharing
- minimize synchronization

MPI

- Compilation and Execution
 - mpic++ and mpiexec
- MPI function calls
 - Initialization
 - Message passing
 - blocking and nonblocking
 - Collective communications

Compilation and Execution

- First, you need to load the MPI environment
 - `module load mpi/mpich-x86_64`
- Compile
 - `mpic++ -Wall --std=c++11 mpi1.cpp`
- Execute with 4 processes
 - `mpiexec -n 4 ./a.out`
- Useful option: `-prepend-rank` (equivalently `-l`)
 - `mpiexec -prepend-rank -n 4 ./a.out`

MPI runtime environment

```
#include <stdio.h>
#include <mpi.h>

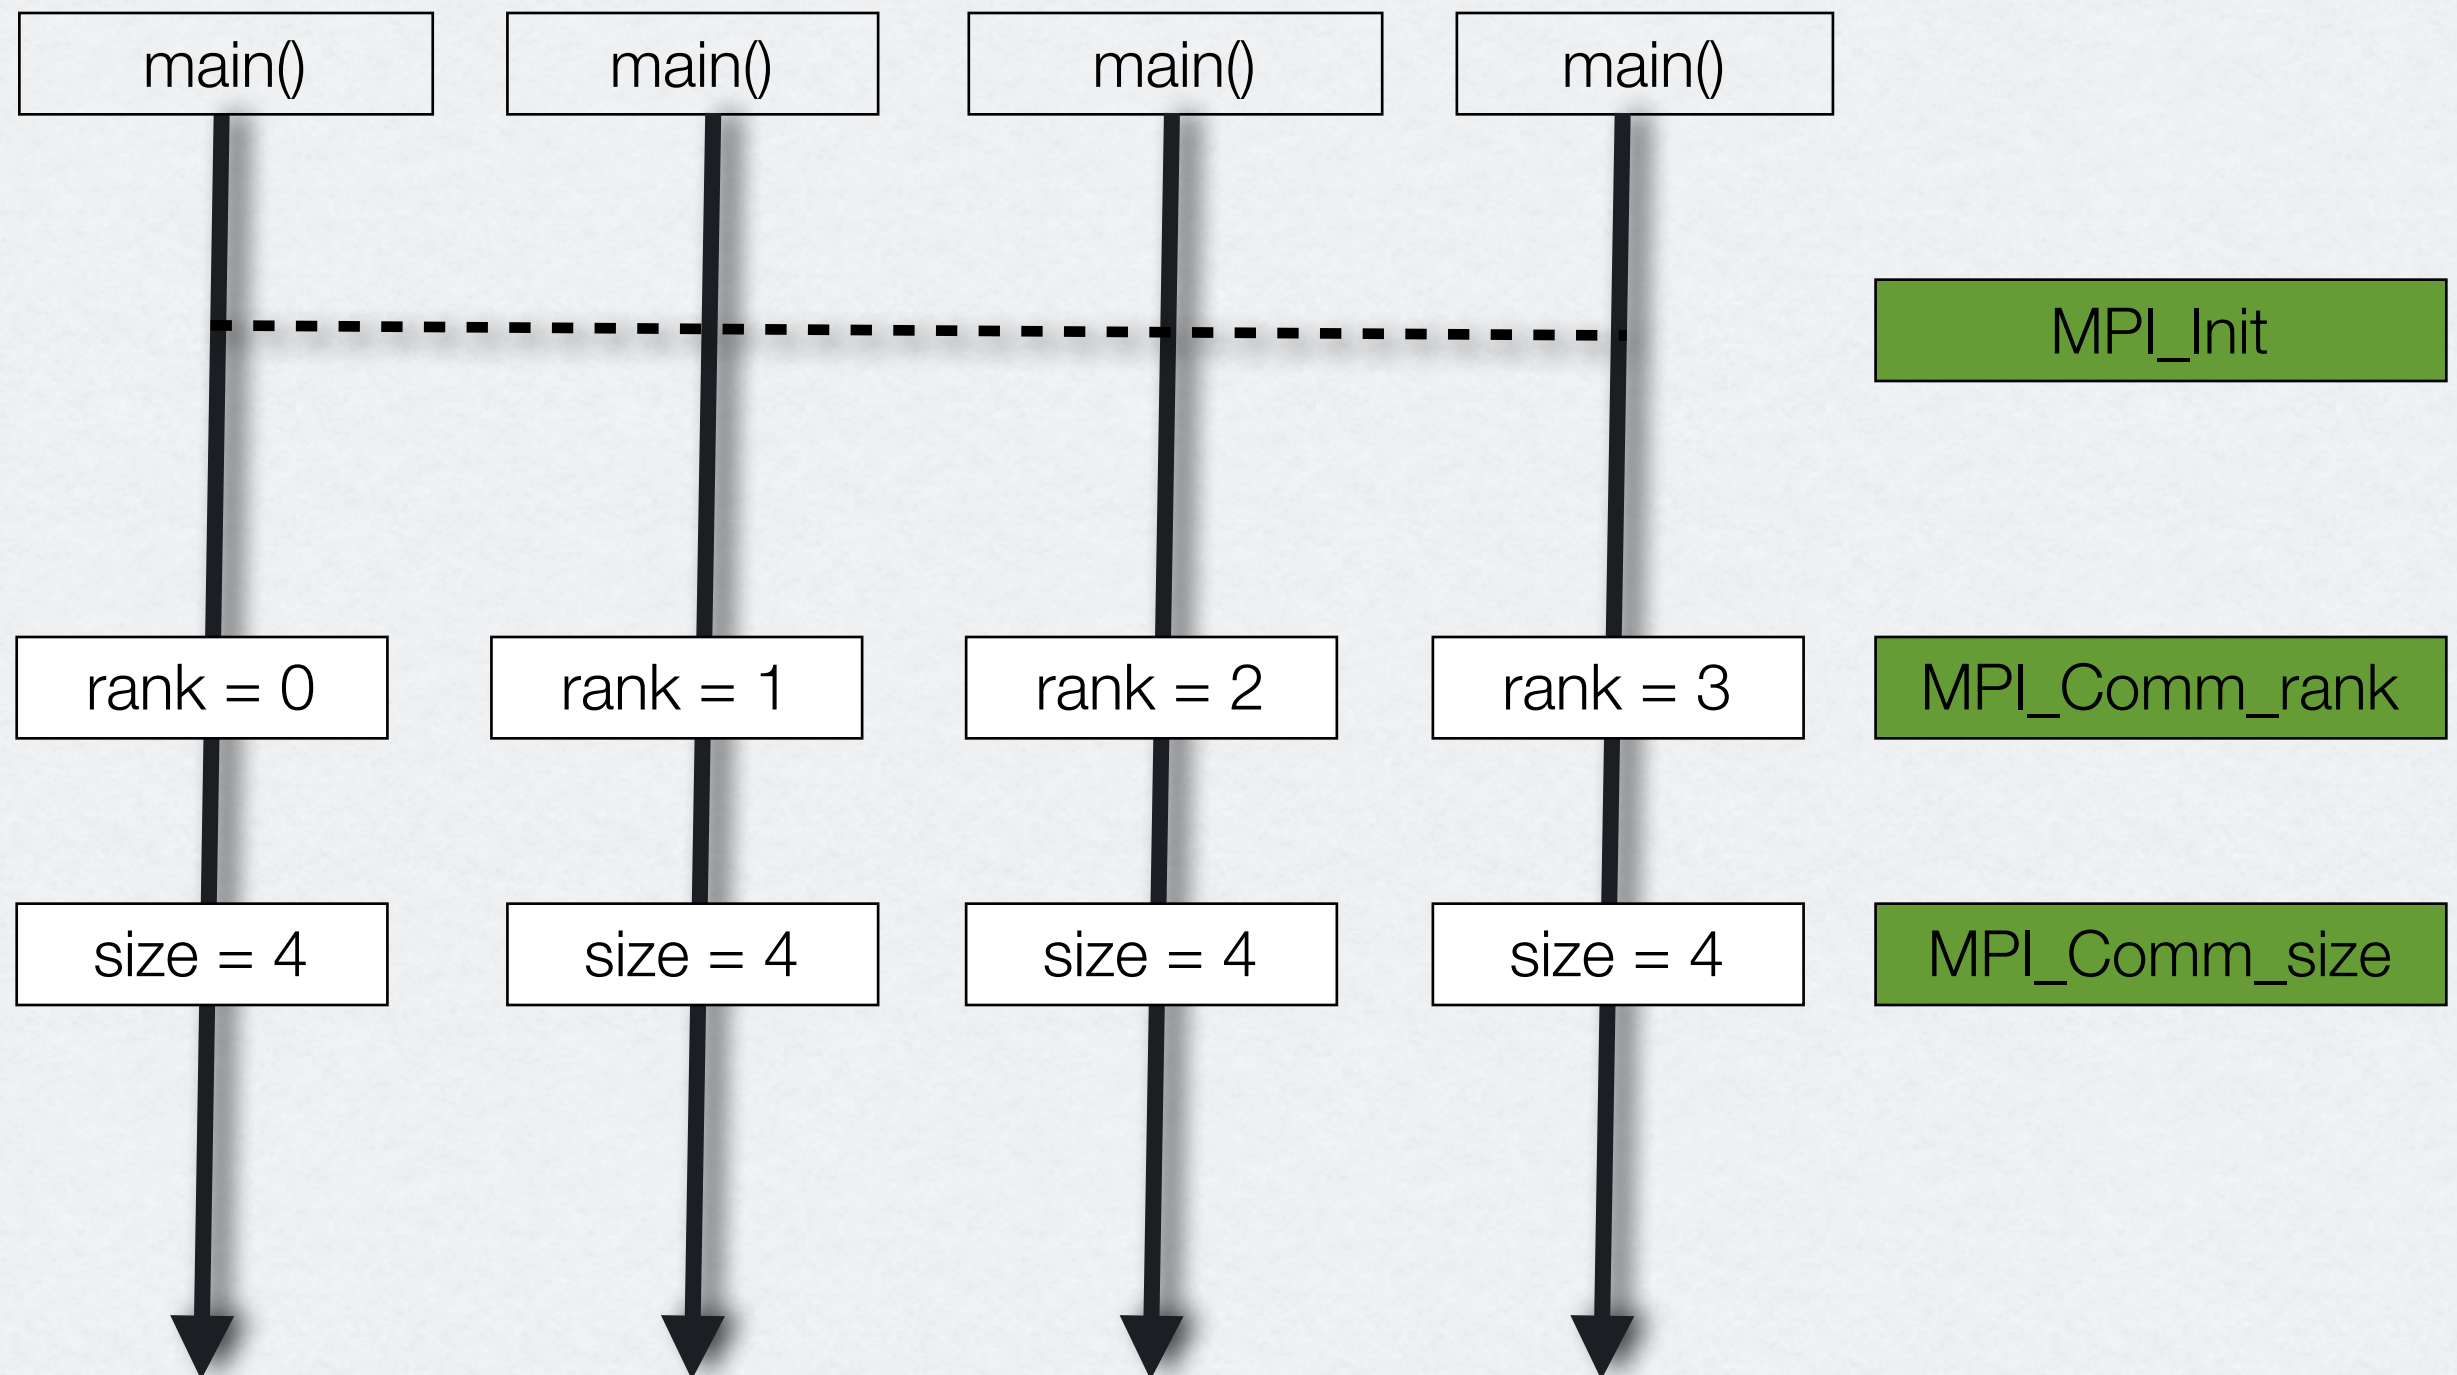
int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv); // initialize the environment

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    printf("Hello from process %d of %d\n", rank, size);

    MPI_Finalize(); // cleanup
    return 0;
}
```


SPMD execution model

- The mpirun / mpiexec utility (spawner) starts the executable on the target cores



Point-to-Point Communication

- Messages are sent and received through MPI_Send and MPI_Recv calls

```
int MPI_Send(void* buf, int count, MPI_Datatype type,  
             int dest, int tag, MPI_Comm comm);  
  
int MPI_Recv(void* buf, int count, MPI_Datatype type,  
             int source, int tag, MPI_Comm comm,  
             MPI_Status* status)
```

- An MPI_Recv matches a message sent by MPI_Send if tag, source and dest match
- MPI_ANY_TAG, MPI_ANY_SOURCE can be used for MPI_Recv

Blocking communication

```
#include <mpi.h>

int main(int argc, char *argv[])
{
    int rank;
    MPI_Status status;

    MPI_Init(&argc, &argv); // initialize the environment
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if (rank == 0) {
        int x = 33;
        MPI_Send(&x, 1, MPI_INT, 1, 123, MPI_COMM_WORLD);
    }
    if (rank == 1) {
        int y;
        MPI_Recv(&y, 1, MPI_INT, 0, 123, MPI_COMM_WORLD, &status);
    }
    MPI_Finalize();
    return 0;
}
```


Non-blocking communication

```
#include <mpi.h>
#include <stdio.h>

main(int argc, char *argv[]) {
    int numtasks, rank, next, prev, buf[2], tag1=1, tag2=2;
    MPI_Request reqs[4];
    MPI_Status stats[2];

    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    prev = rank-1;
    next = rank+1;
    if (rank == 0) prev = numtasks - 1;
    if (rank == (numtasks - 1)) next = 0;

    MPI_Irecv(&buf[0], 1, MPI_INT, prev, tag1, MPI_COMM_WORLD, &reqs[0]);
    MPI_Irecv(&buf[1], 1, MPI_INT, next, tag2, MPI_COMM_WORLD, &reqs[1]);

    MPI_Isend(&rank, 1, MPI_INT, prev, tag2, MPI_COMM_WORLD, &reqs[2]);
    MPI_Isend(&rank, 1, MPI_INT, next, tag1, MPI_COMM_WORLD, &reqs[3]);

    { /* do some work */ }

    MPI_Waitall(4, reqs, stats);

    MPI_Finalize();
}
```

MPI_Barrier

- Synchronization operation. Creates a barrier synchronization in a group. Each task, when reaching the MPI_Barrier call, blocks until all tasks in the group reach the same MPI_Barrier call. Then all tasks are free to proceed.

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv); // initialize the environment

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Barrier(MPI_COMM_WORLD);
    printf("Hello from process %d of %d\n", rank, size);

    MPI_Finalize(); // cleanup
    return 0;
}
```

MPI_Bcast


- Data movement operation. Broadcasts (sends) a message from the process with rank "root" to all other processes in the group

MPI_Bcast

Broadcasts a message to all other processes of that group

```
count = 1;  
source = 1;  
MPI_Bcast(&msg, count, MPI_INT, source, MPI_COMM_WORLD);
```

broadcast originates in task 1



MPI_Reduce

- Collective computation operation. Applies a reduction operation on all tasks in the group and places the result in one task.

MPI_Reduce

Perform and associate reduction operation across all tasks in the group and place the result in one task

count = 1;

dest = 1;

result will be placed in task 1

MPI_Reduce(sendbuf, recvbuf, count, MPI_INT, MPI_SUM,
dest, MPI_COMM_WORLD);

task 0

task 1

task 2

task 3

1

2

3

4

← sendbuf (before)

10

← recvbuf (after)


MPI_Allreduce

- Collective computation operation + data movement. Applies a reduction operation and places the result in all tasks in the group.
- Equivalent to MPI_Reduce + MPI_Bcast

MPI_Allreduce

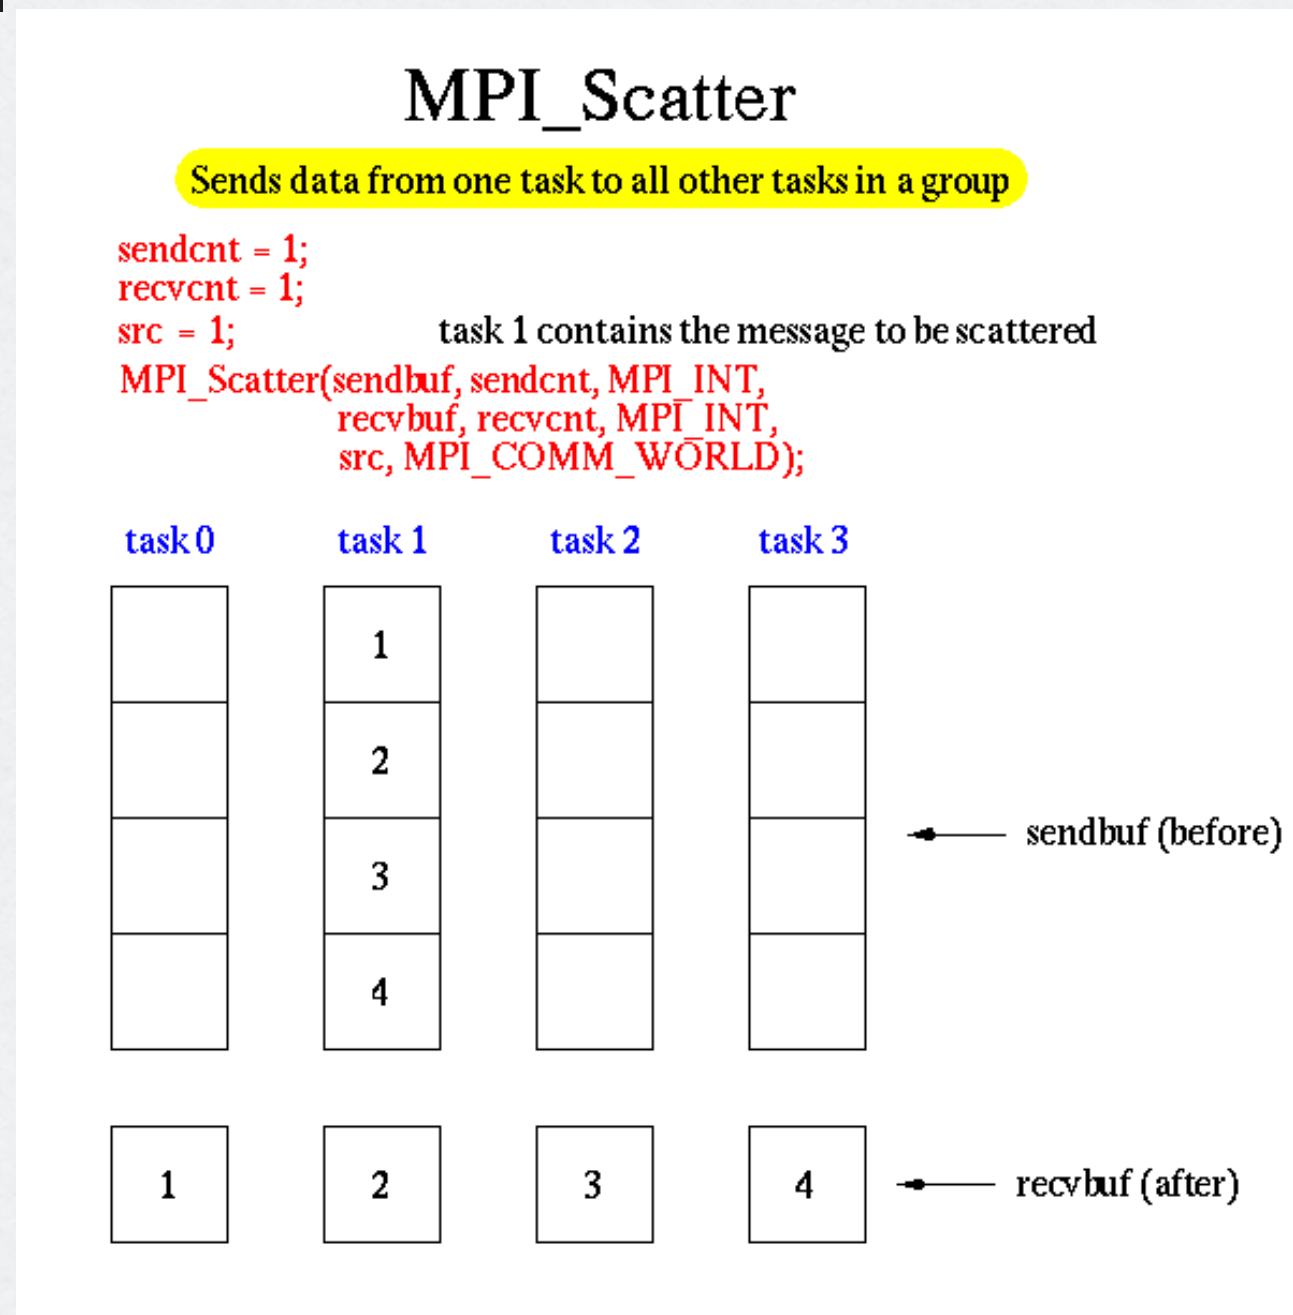
Perform and associate reduction operation across all tasks in the group and place the result in all tasks

```
count = 1;  
MPI_Allreduce(sendbuf, recvbuf, count, MPI_INT, MPI_SUM,  
              MPI_COMM_WORLD);
```



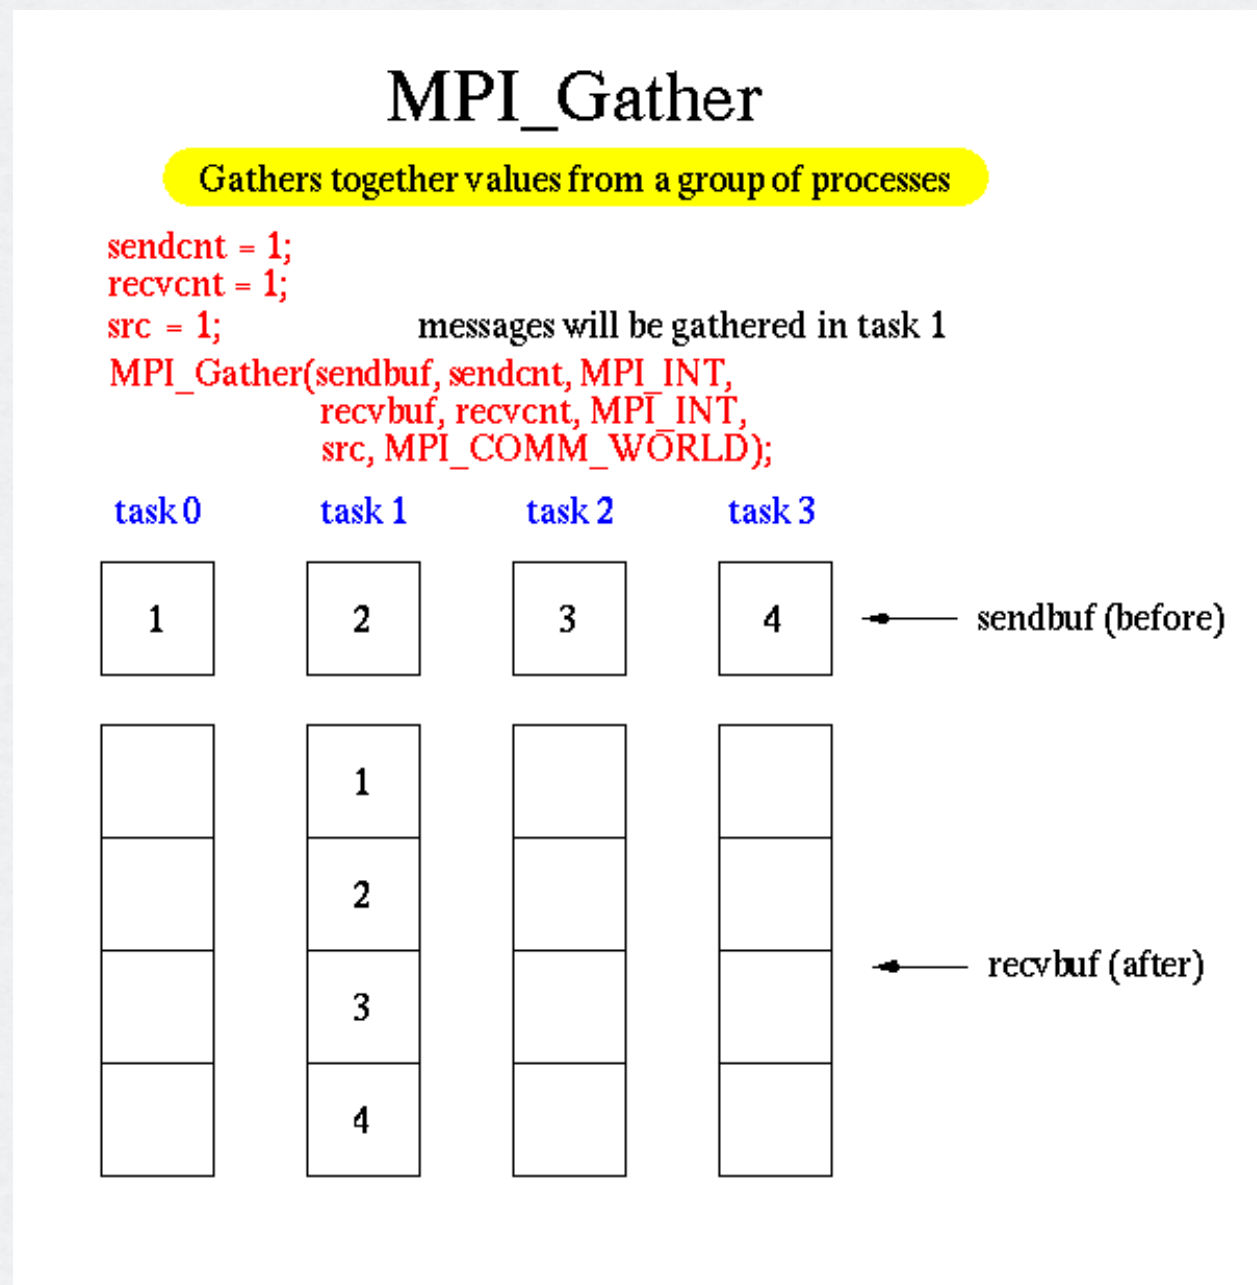
MPI_Scatter

- Data movement operation. Distributes distinct messages from a single source task to each task in the group.



MPI_Gather

- Data movement operation. Gathers distinct messages from each task in the group to a single destination task (reverse of MPI_Scatter)



Measuring time

```
#include <mpi.h>
#include <stdio.h> // printf
#include <unistd.h> // sleep

int main( int argc, char *argv[] )
{
    double t1, t2;

    MPI_Init(&argc, &argv);
    t1 = MPI_Wtime();
    sleep(2);
    t2 = MPI_Wtime();
    printf("Elapsed time =%f seconds\n", t2-t1);
    MPI_Finalize();
    return 0;
}
```

Implementation of MPI_Bcast

```
#include <mpi.h>

int main(int argc , char **argv)
{
    int size, rank;
    double data;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    srand48(rank);

    for (int k = 0; k < 10; k++) {
        if (!rank) data = drand48();

        MPI_Bcast(&data, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);
        printf("Step %d: I am Process %d Data = %f\n", k, rank, data);
    }
    MPI_Finalize();
}
```


Implementation of MPI_Bcast

```
#include <mpi.h>

int main(int argc , char **argv)
{
    int size, rank;
    double data;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    srand48(rank);

    for (int k = 0; k < 10; k++) {
        if (!rank) data = drand48();

        if (!rank) {
            for (i = 1; i < size; i++)
                MPI_Send(&data, 1, MPI_DOUBLE, i, 0, MPI_COMM_WORLD);
        } else {
            MPI_Recv(&data, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
        }

        printf("Step %d: I am Process %d Data = %f\n", k, rank, data);
    }
    MPI_Finalize();
}
```

Next semester: MPI + OpenMP

```
#include <stdio.h>
#include <mpi.h>
#include <omp.h>

int main(int argc, char *argv[])
{
    int rank, nprocs;

    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&nprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);

    omp_set_num_threads(2);
    #pragma omp parallel
    {
        printf("Hello, world.  I am %d of %d  thread=%d\n", rank,
               nprocs, omp_get_thread_num());
    }

    MPI_Finalize();
    return 0;
}
```