

DERIVATION OF THE  $\vec{U}$ - $\omega$  EULER EQUATION  
FROM THE  $\vec{U}$ - $P$  FORM (BLACKBOARD)

Euler equations in velocity-pressure form

$$\begin{array}{l} \textcircled{1} \left\{ \frac{D\vec{U}}{Dt} = \frac{\partial \vec{U}}{\partial t} + (\vec{U} \cdot \nabla) \vec{U} = - \frac{1}{\rho} \nabla P \quad (\text{Momentum equation}) \right. \\ \textcircled{2} \left. \nabla \cdot \vec{U} = 0 \quad (\text{Continuity equation - incompressibility}) \right. \end{array}$$

$\nabla \times \textcircled{1}$ :

$$\begin{aligned} \nabla \times \frac{\partial \vec{U}}{\partial t} + \nabla \times [(\vec{U} \cdot \nabla) \vec{U}] &= \nabla \times \left( - \frac{1}{\rho} \nabla P \right) \quad \text{(curl of gradient of a scalar field)} \\ \frac{\partial \omega}{\partial t} + \nabla \times \left[ \frac{1}{2} \nabla(\vec{U} \cdot \vec{U}) - \vec{U} \times (\nabla \times \vec{U}) \right] &= \quad \text{(ii)} \quad \nabla \times \vec{U} = \omega \\ &= \frac{\partial \omega}{\partial t} + (\vec{U} \cdot \nabla) \omega - (\omega \cdot \nabla) \vec{U} + \cancel{\omega(\nabla \cdot \vec{U})} - \cancel{\vec{U}(\nabla \cdot \omega)} = 0 \\ &\quad \nabla \cdot (\nabla \times \vec{U}) = 0 \end{aligned}$$

$$\Rightarrow \boxed{\frac{\partial \omega}{\partial t} + (\vec{U} \cdot \nabla) \omega = (\omega \cdot \nabla) \vec{U}}$$

Further simplification in 2D:  $(\omega \cdot \nabla) \vec{U} = 0$

$$\Rightarrow \boxed{\frac{\partial \omega}{\partial t} + (\vec{U} \cdot \nabla) \omega = 0} \quad \text{2D}$$

## RECOVERING OF VELOCITIES FROM VORTICITIES

$$\nabla \times (\omega = \nabla \times \vec{U})$$

$$\nabla \times \omega = \nabla \times (\nabla \times \vec{U}) = \nabla(\nabla \cdot \vec{U}) - \nabla^2 \vec{U}$$

$$\Rightarrow \boxed{\Delta \vec{U} = -\nabla \times \omega}$$

Poisson Equation

$$\begin{aligned} \text{OR} \quad \omega &= \nabla \times \vec{U} = \nabla \times (\nabla \times \psi) \\ &= \nabla(\nabla \cdot \psi) - \nabla^2 \psi \end{aligned}$$

$$\begin{aligned} \text{OR} \quad \omega &= \nabla \times \vec{U} = \nabla \times (\nabla \times \psi) \\ &= \nabla(\nabla \cdot \psi) - \nabla^2 \psi \end{aligned}$$

$$\Rightarrow \boxed{\Delta \psi = -\omega}$$

Poisson Equation  
(scalar)

$$\rightarrow \vec{U} = \nabla \times \psi$$

GREEN'S FUNCTION SOLUTION FOR THE DIFFUSION EQUATION

$$\frac{\partial u}{\partial t} = D \frac{\partial^2 u}{\partial x^2} \quad (10)$$

the Green's function solution satisfies

$$\left[ D \frac{\partial^2}{\partial x^2} - \frac{\partial}{\partial t} \right] G = -\delta(x)\delta(t)$$

the Fourier Transform  $\hat{G} = \int_{-\infty}^{\infty} G e^{-ikx} dx$   
yields:

$$\left[ Dk^2 + \frac{\partial}{\partial t} \right] \hat{G} = \delta(t)$$

this is now an ODE which has solution

$$\hat{G} = H(t) e^{-tDk^2} \quad \begin{cases} H: \text{Heaviside function} \\ \frac{\partial H(t)}{\partial t} = \delta(t) \end{cases}$$

test:

$$\frac{\partial \hat{G}}{\partial t} = \delta(t) - Dk^2 e^{-tDk^2}$$

Taking the inverse FT of  $\hat{G}$ , we obtain

$$\begin{aligned} G &= \frac{1}{2\pi} \int_{-\infty}^{\infty} H(t) e^{-tDk^2} e^{ikx} dk = \\ &= \frac{H(t)}{2\pi} \int_{-\infty}^{\infty} e^{-tD(k^2 - \frac{ikx}{2D})} dk = \\ &= \frac{H(t)}{2\pi} \int_{-\infty}^{\infty} e^{-tD(k - \frac{ix}{2D})^2 - \frac{x^2}{4Dt}} dk = \\ &= \frac{H(t)}{2\pi} e^{-\frac{x^2}{4Dt}} \int_{-\infty}^{\infty} e^{-tD(k - \frac{ix}{2D})^2} dk = \\ &= \frac{H(t)}{\sqrt{4\pi t + D}} e^{-\frac{x^2}{4Dt}} \end{aligned}$$

$\Rightarrow$  Gaussian!

$\rightarrow$  we are interested in  $t > 0$ ,  
thus

$$G = \frac{1}{\sqrt{4\pi t + D}} e^{-\frac{x^2}{4Dt}}$$

This is a short range interaction!

$$\int_{-\infty}^{\infty} e^{-ak^2} dk = \sqrt{\frac{\pi}{a}}$$

# Cell-Lists

High Performance Computing for Science and Engineering

# Diffusion as N-Body Problem

The fundamental solution to the diffusion equation

$$\frac{\partial u}{\partial t} = D \Delta u$$

is given by the Green's function

$$G(\mathbf{x}, t) = \left( \frac{1}{4\pi Dt} \right)^{d/2} e^{-\frac{\mathbf{x} \cdot \mathbf{x}}{4Dt}}$$

which is a  $d$ -dimensional Gaussian function



fast decay

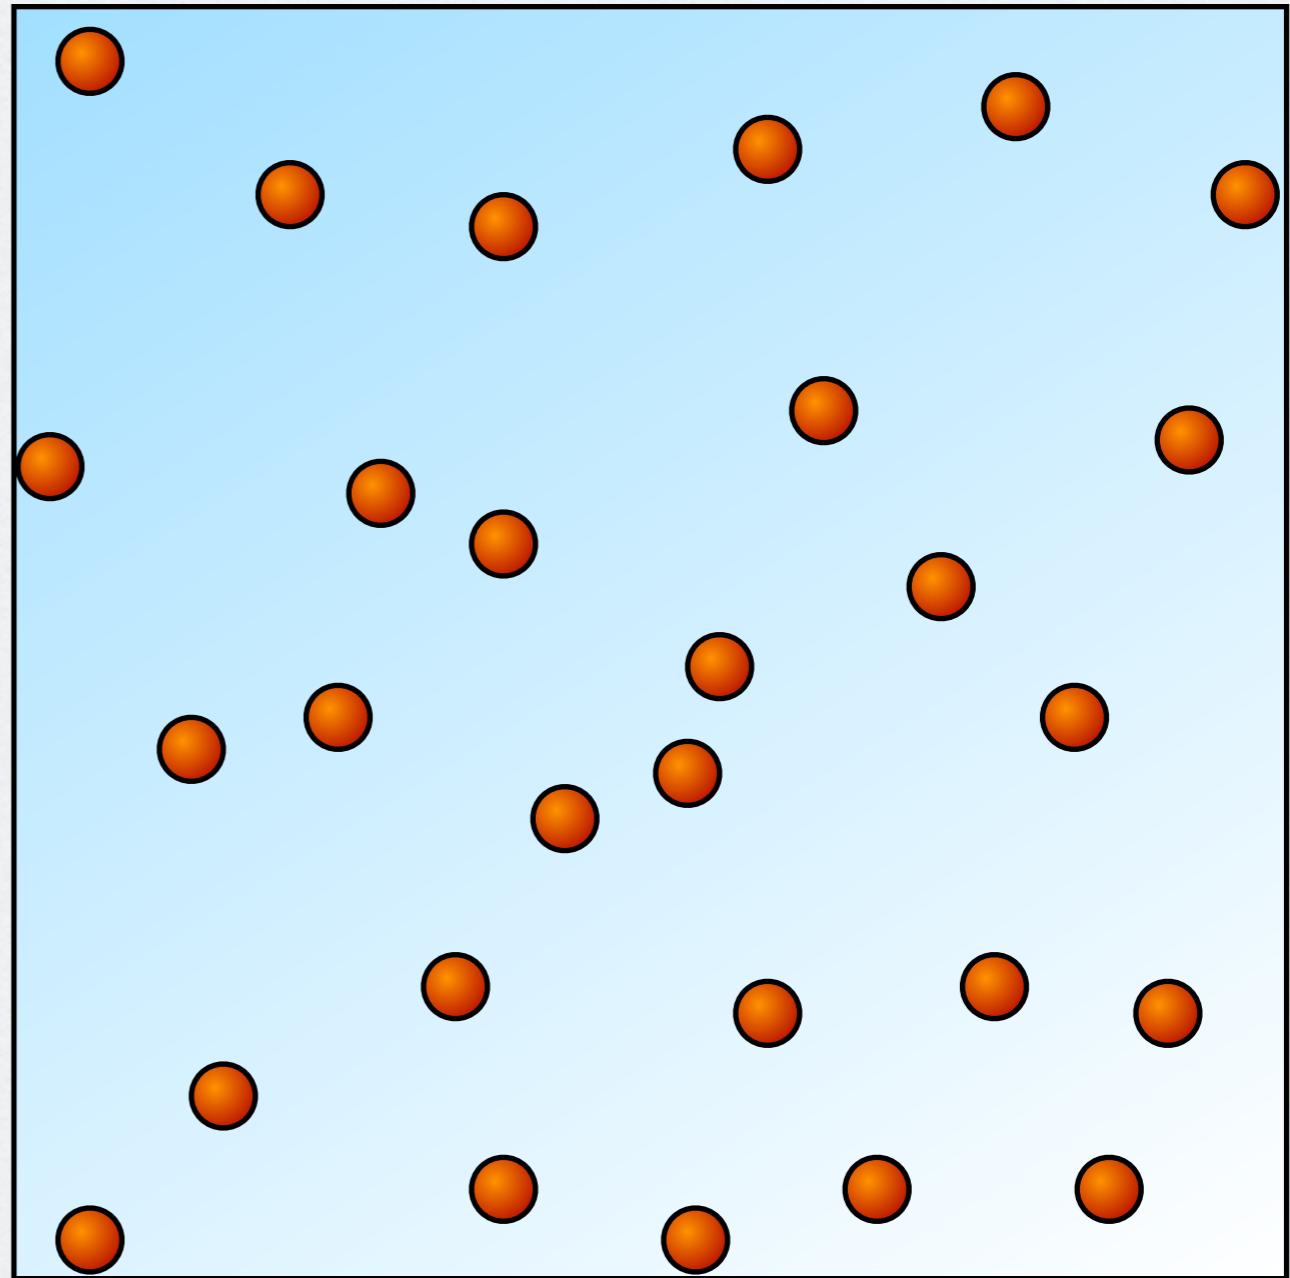
**short-range interactions**

# N-Body Problem in MD

$$\frac{d\mathbf{x}_i}{dt} = \mathbf{v}_i$$

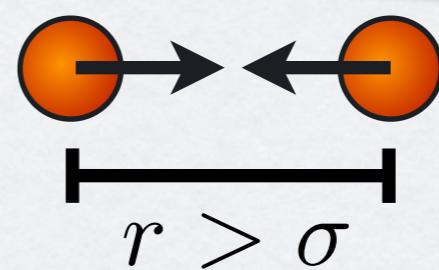
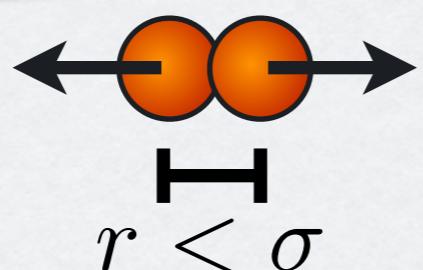
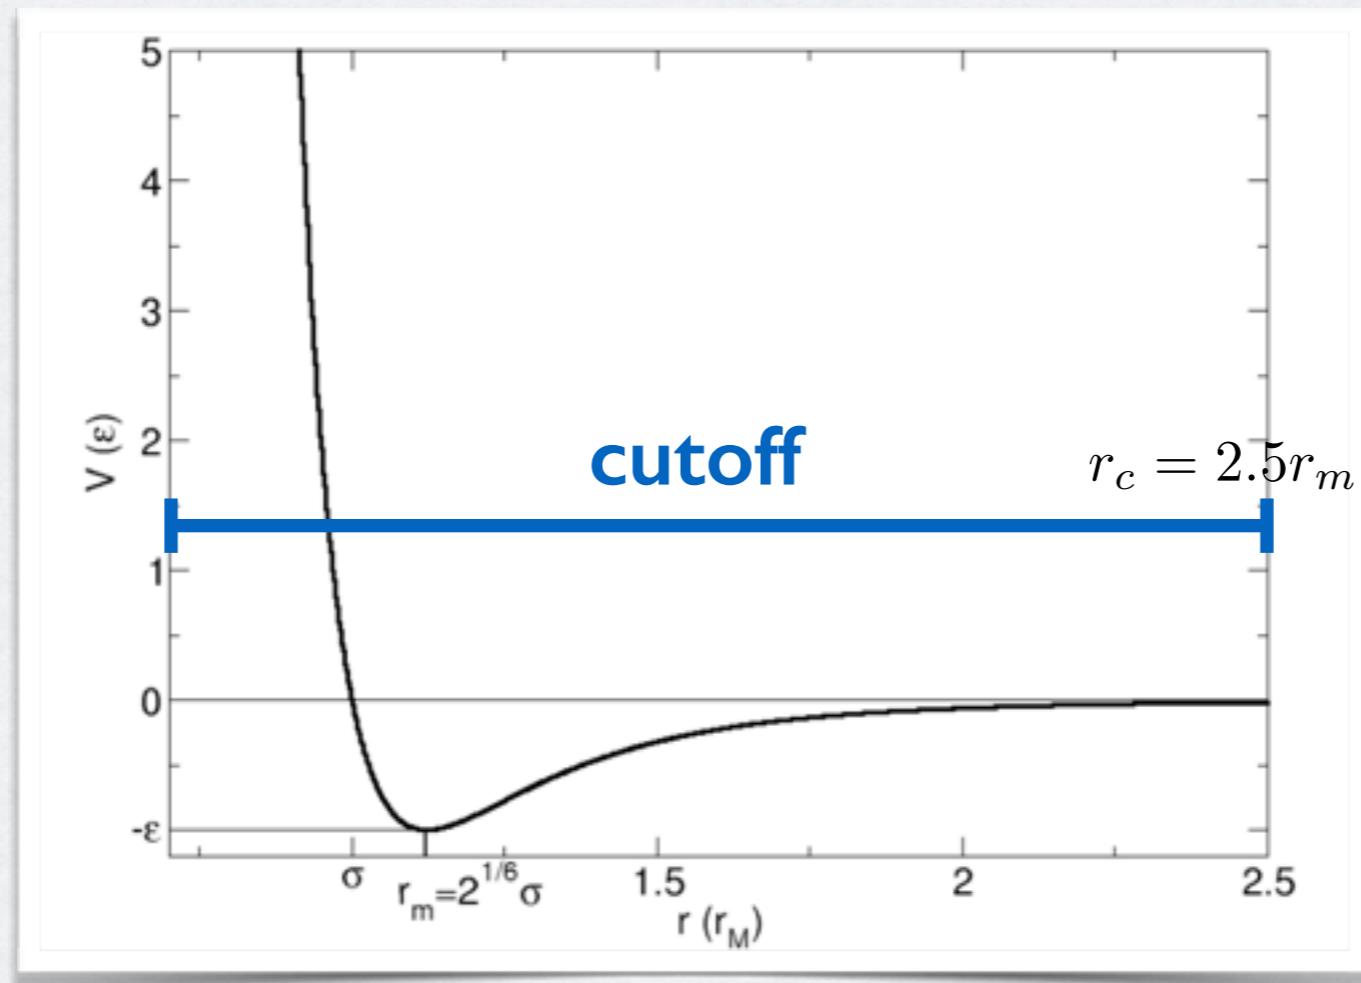
$$m_i \frac{d\mathbf{u}_i}{dt} = \sum_{j=1; j \neq i}^N \mathbf{F}_{ij}$$

$$\mathbf{F}_{ij} = -\frac{\mathbf{x}_i - \mathbf{x}_j}{||\mathbf{x}_i - \mathbf{x}_j||} \frac{dV}{dr}$$

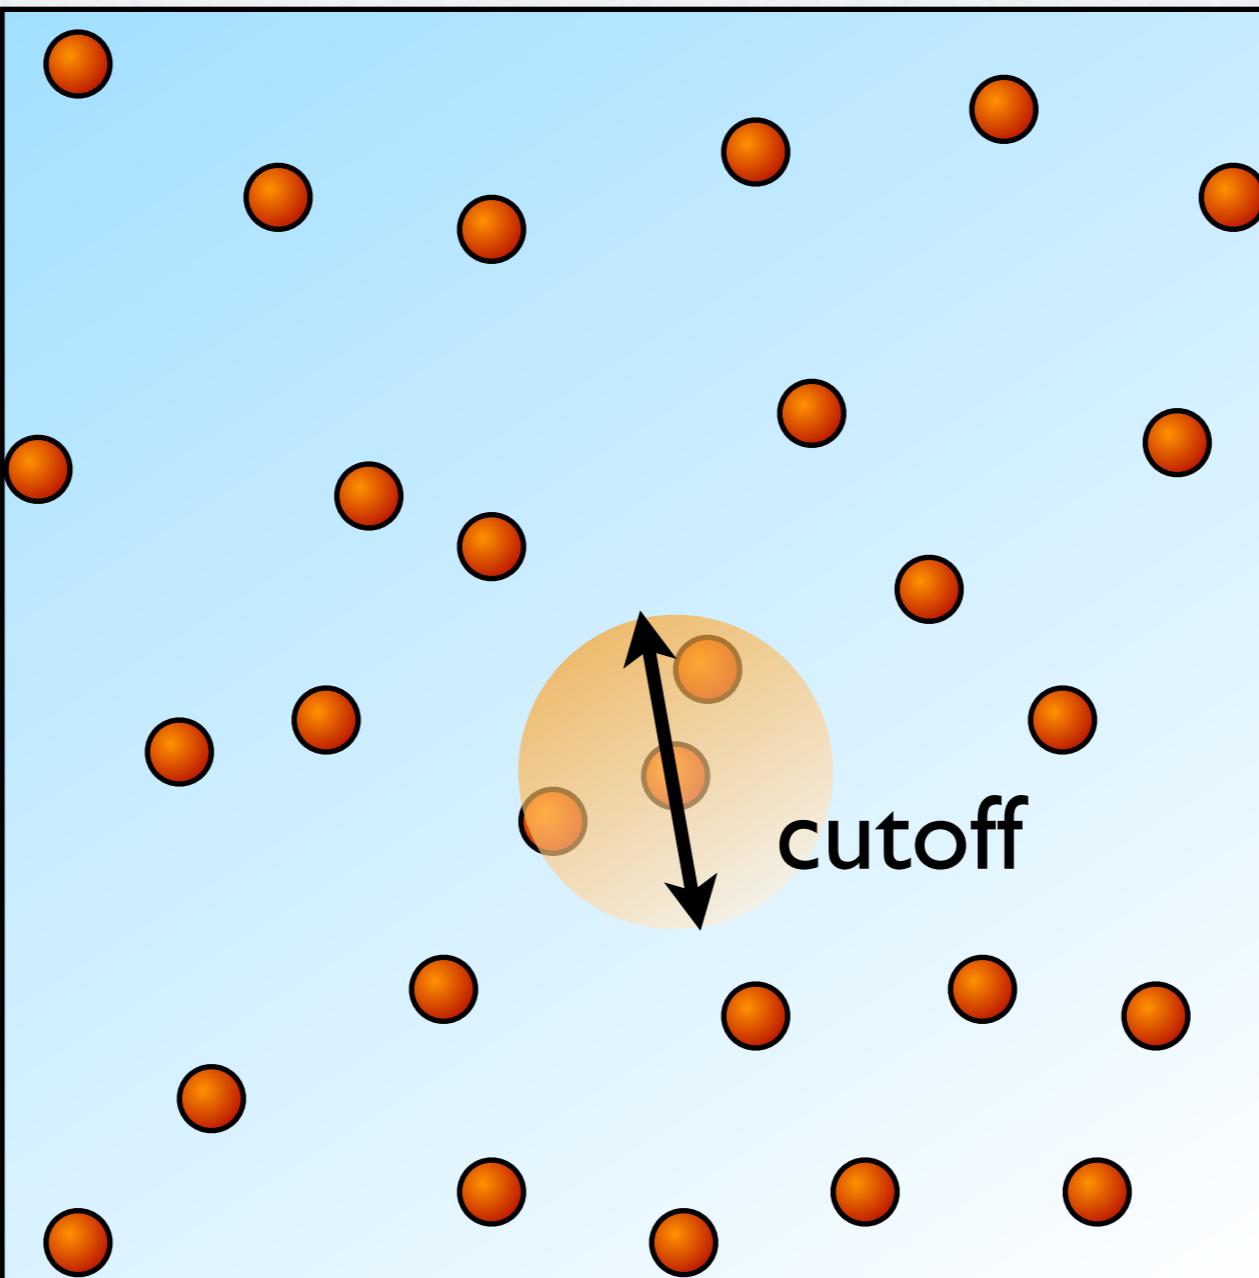


# The Lennard-Jones Potential

$$V_c(r) = \begin{cases} 4\epsilon \left[ \left(\frac{\sigma}{r}\right)^{12} - \left(\frac{\sigma}{r}\right)^6 \right] - V(r_c) & \text{for } r < r_c, \\ 0 & \text{else,} \end{cases}$$



# Local Interactions



Can we take advantage of this to avoid  $\mathcal{O}(N^2)$  operations?

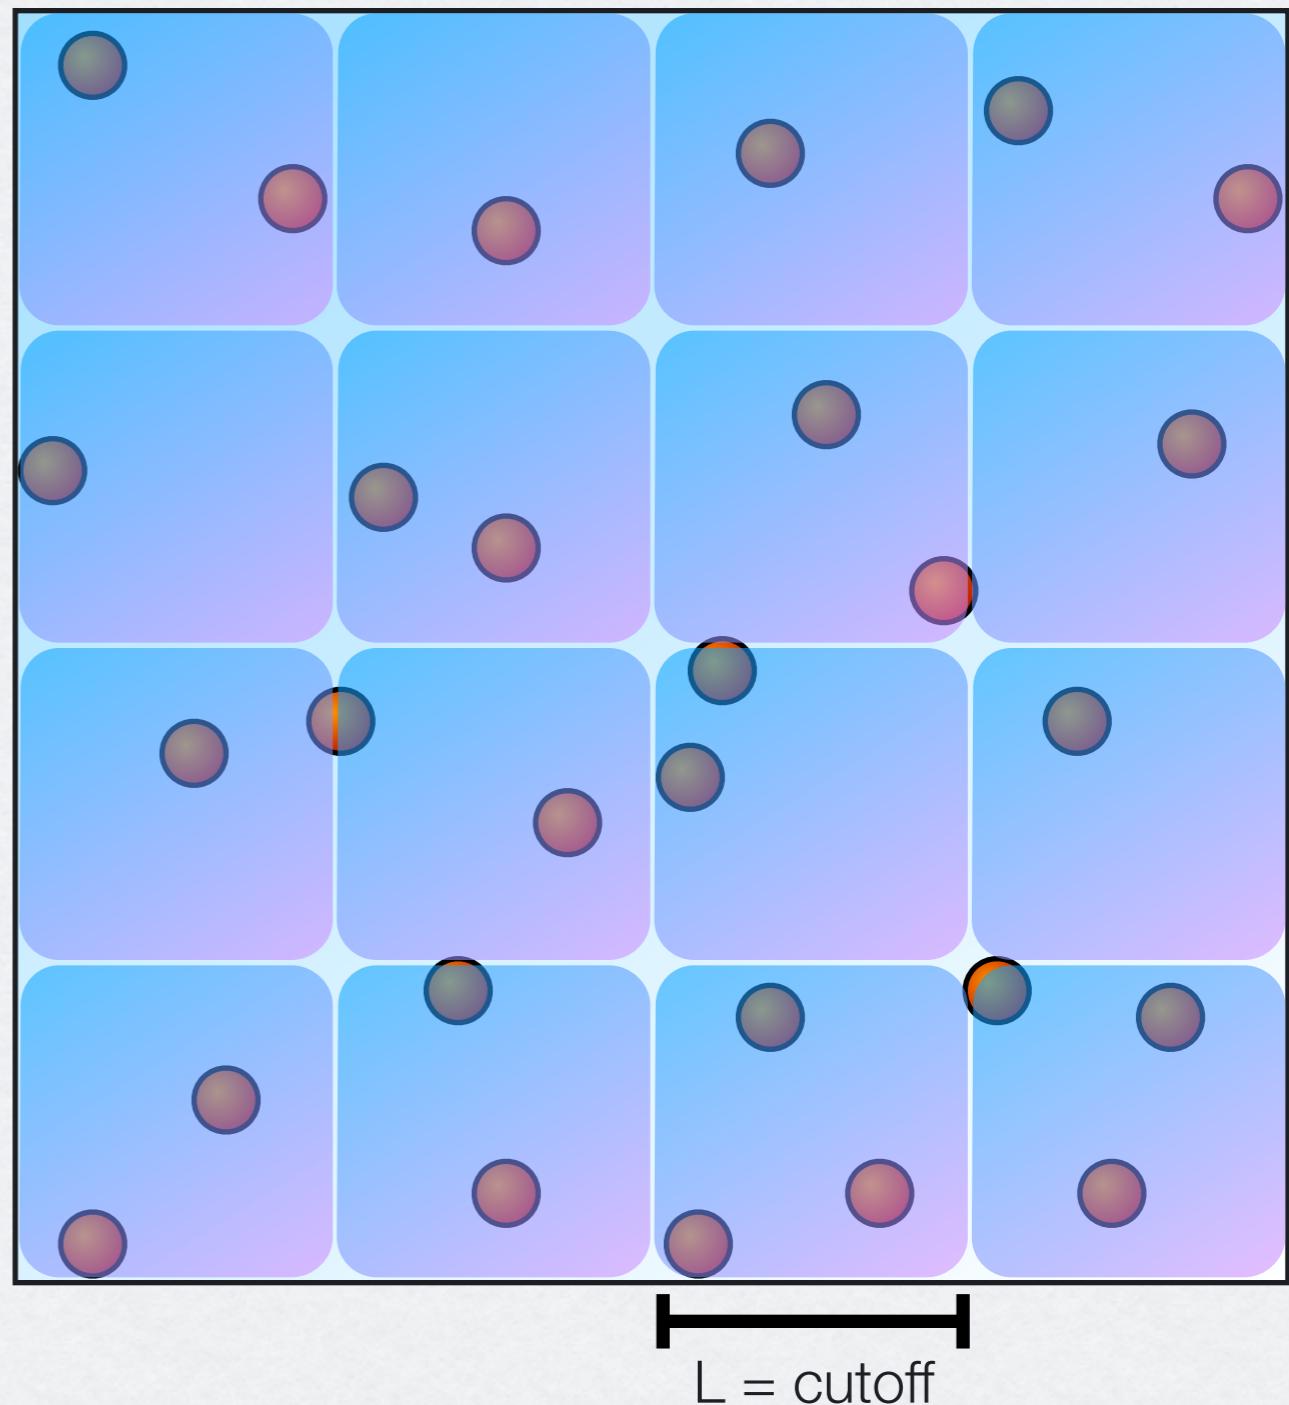
# Idea - Subdivide Domain

## Domain subdivision in *cells*:

- Cell size is the cutoff
- Cells contain particles

## Particles:

- Spread into the cells according to their position



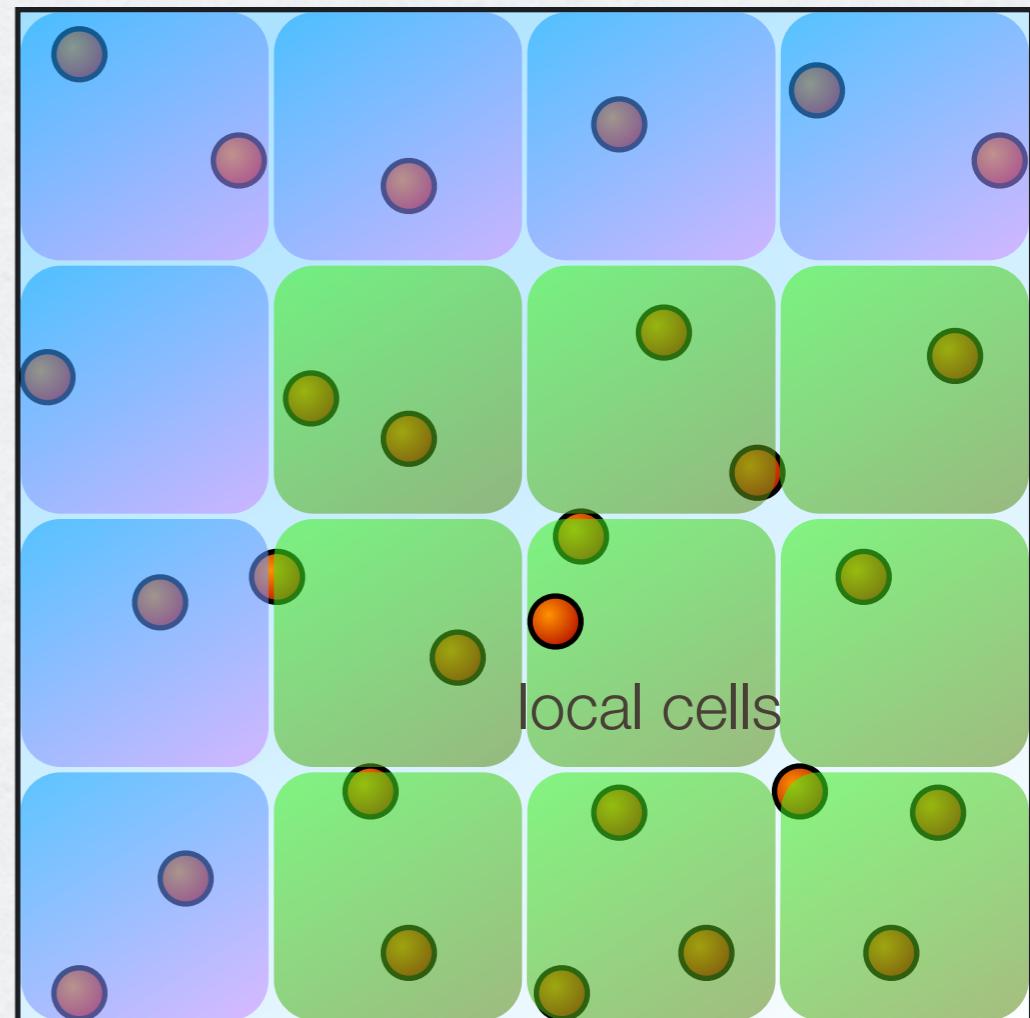
# Computing Interactions

Interactions: only local cells contribute

Cost of N interactions (2D):

- $c=N/m$  (average #particles/cell)
- assume constant  $c$  w.r.t.  $N$
- cost for one particle:  $9c$
- total:  $9cN \rightarrow \mathcal{O}(N)$
- algorithmic improvement:

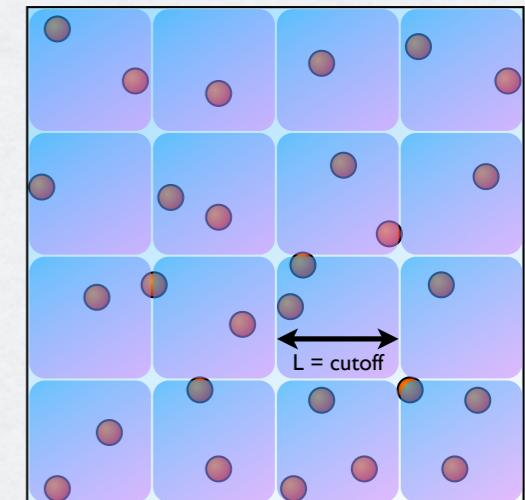
$$\frac{N^2}{9cN} = \frac{N}{9c}$$



m cells

# Cell List Construction

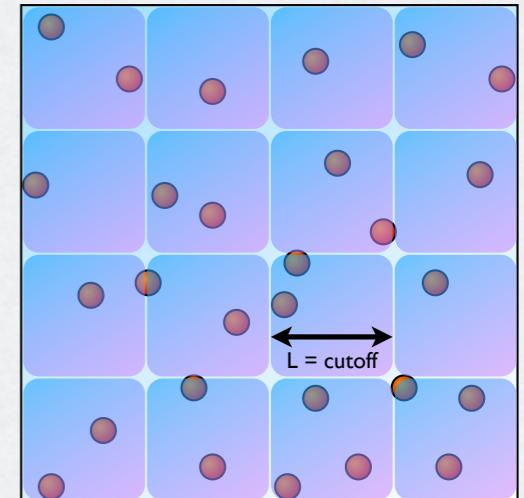
HOW TO FIND THE PARTICLE NEAREST NEIGHBORS?



# Cell List Construction

## HOW TO FIND THE PARTICLE NEAREST NEIGHBORS?

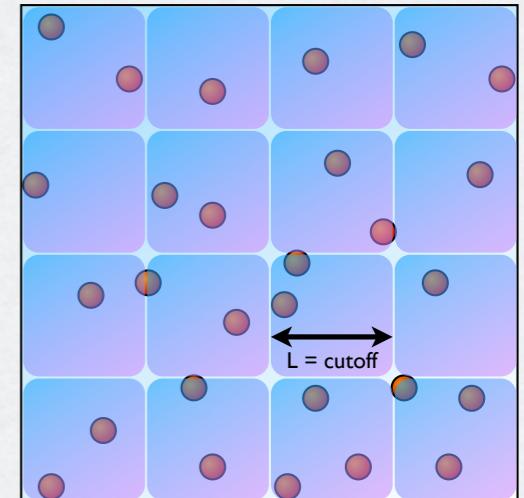
- Group particles according to the “address” of their cell



# Cell List Construction

## HOW TO FIND THE PARTICLE NEAREST NEIGHBORS?

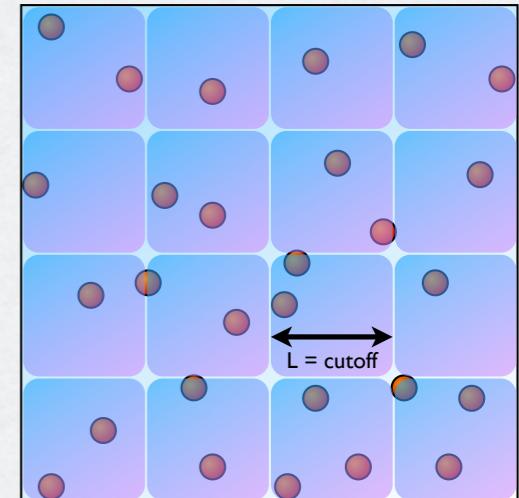
- Group particles according to the “address” of their cell
- Cell address for a single particle can be computed in  $O(1)$



# Cell List Construction

## HOW TO FIND THE PARTICLE NEAREST NEIGHBORS?

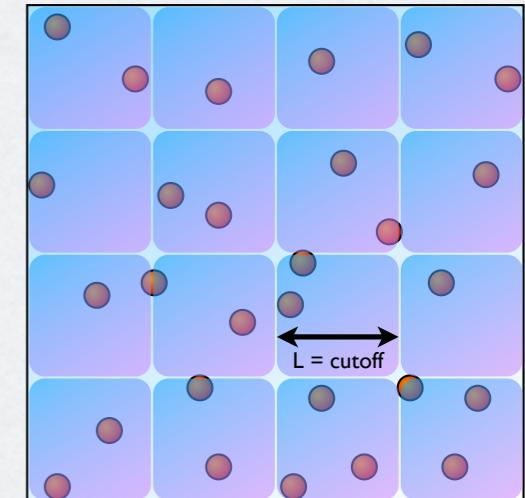
- Group particles according to the “address” of their cell
- Cell address for a single particle can be computed in  $O(1)$
- Insertion of a particle in its associated cell in  $O(1)$



# Cell List Construction

## HOW TO FIND THE PARTICLE NEAREST NEIGHBORS?

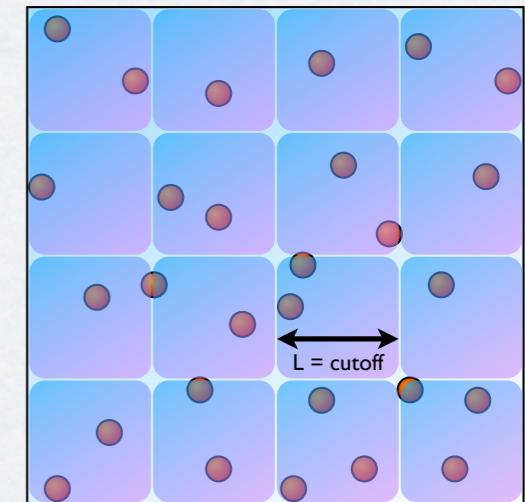
- Group particles according to the “address” of their cell
  - Cell address for a single particle can be computed in  $O(1)$
  - Insertion of a particle in its associated cell in  $O(1)$
- Cost for the creation of cell lists for  $N$  particles is in  $O(N)$



# Cell List Construction

## HOW TO FIND THE PARTICLE NEAREST NEIGHBORS?

- Group particles according to the “address” of their cell
  - Cell address for a single particle can be computed in  $O(1)$
  - Insertion of a particle in its associated cell in  $O(1)$
- Cost for the creation of cell lists for  $N$  particles is in  $O(N)$

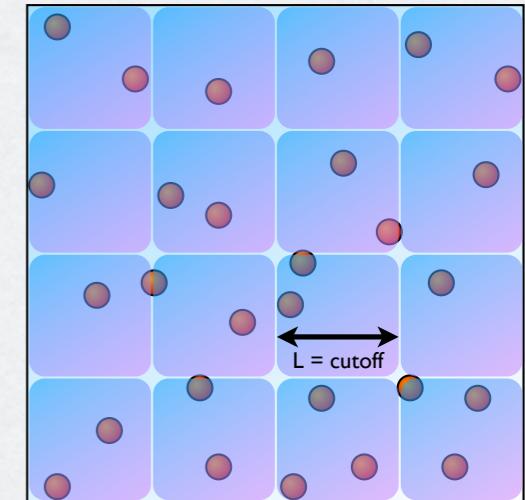


**STL-based implementation with  $M^3$  cells (3D, domain  $[0,1]^3$ )**

# Cell List Construction

## HOW TO FIND THE PARTICLE NEAREST NEIGHBORS?

- Group particles according to the “address” of their cell
  - Cell address for a single particle can be computed in  $O(1)$
  - Insertion of a particle in its associated cell in  $O(1)$
- Cost for the creation of cell lists for  $N$  particles is in  $O(N)$



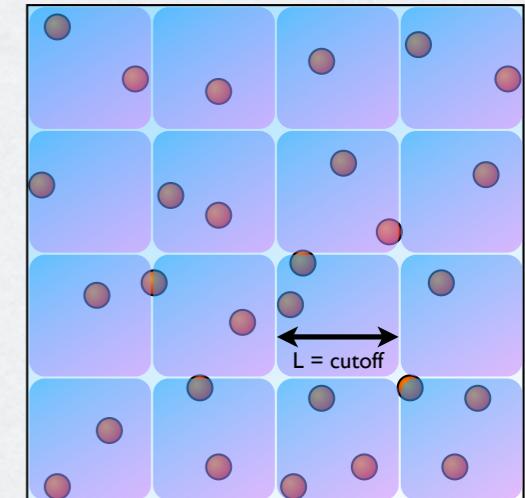
## STL-based implementation with $M^3$ cells (3D, domain $[0,1]^3$ )

- `std::vector<ParticleID*> cells(M*M*M)` contains a vector for each cell

# Cell List Construction

## HOW TO FIND THE PARTICLE NEAREST NEIGHBORS?

- Group particles according to the “address” of their cell
  - Cell address for a single particle can be computed in  $O(1)$
  - Insertion of a particle in its associated cell in  $O(1)$
- Cost for the creation of cell lists for  $N$  particles is in  $O(N)$



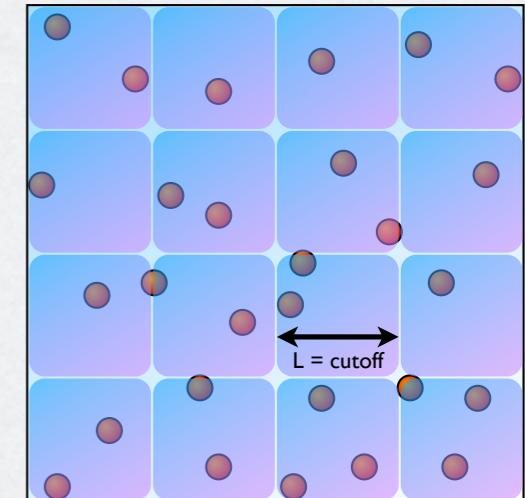
## STL-based implementation with $M^3$ cells (3D, domain $[0,1]^3$ )

- `std::vector<ParticleID*> cells(M*M*M)` contains a vector for each cell
- cell address for particle at  $(x,y,z)$ :  $(ix,iy,iz) = (\text{floor}(x/L)), \text{floor}(y/L), \text{floor}(z/L))$ ,  $L = 1/M$

# Cell List Construction

## HOW TO FIND THE PARTICLE NEAREST NEIGHBORS?

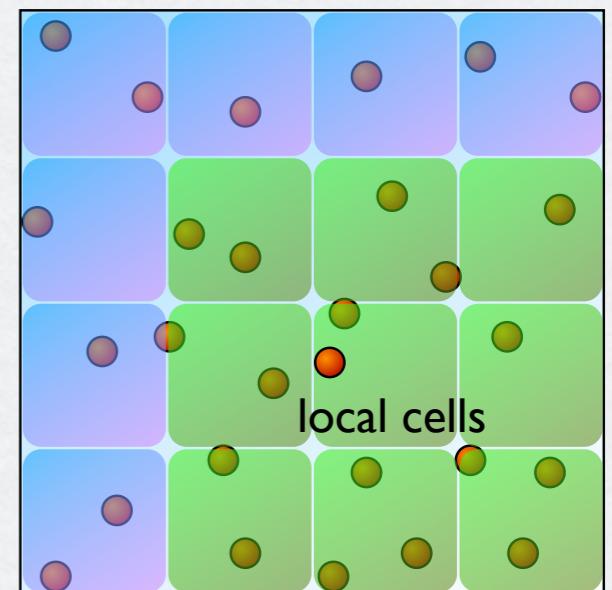
- Group particles according to the “address” of their cell
  - Cell address for a single particle can be computed in  $O(1)$
  - Insertion of a particle in its associated cell in  $O(1)$
- Cost for the creation of cell lists for  $N$  particles is in  $O(N)$



## STL-based implementation with $M^3$ cells (3D, domain $[0,1]^3$ )

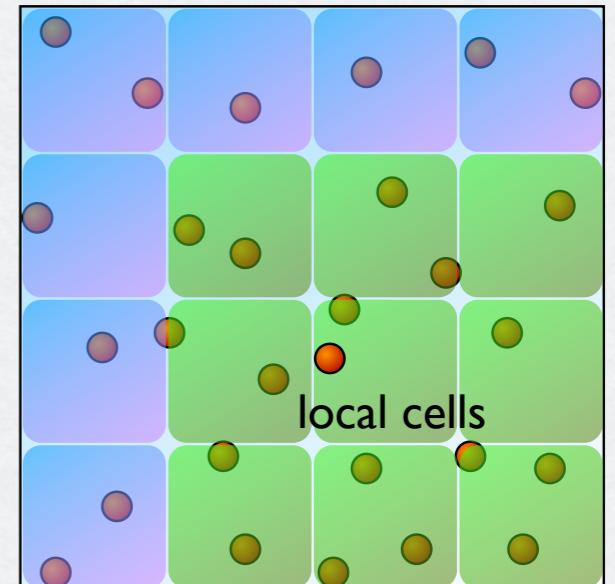
- `std::vector<ParticleID*> cells(M*M*M)` contains a vector for each cell
- cell address for particle at  $(x,y,z)$ :  $(ix,iy,iz) = (\text{floor}(x/L)), \text{floor}(y/L), \text{floor}(z/L))$ ,  $L = 1/M$ 
  - 1D address into cells:  $\text{idx} = ix + iy*M + iz*M*M = ix + M*(iy + M*iz)$

# Cell List Traversal



# Cell List Traversal

- Given an initialized  
`std::vector<ParticleID*> cells (M*M*M)`

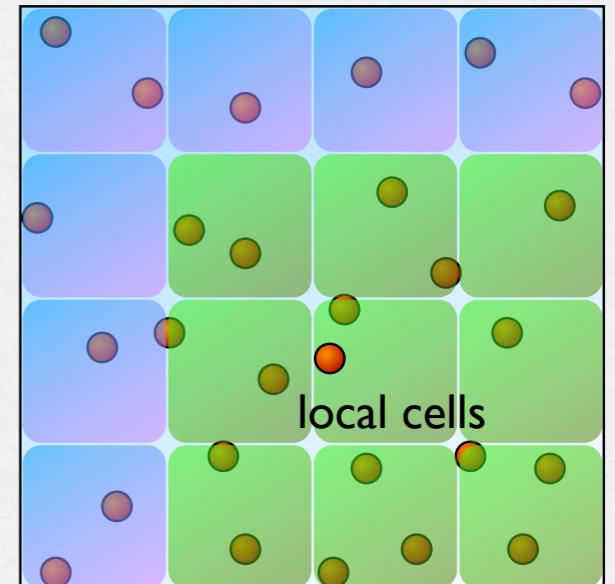


# Cell List Traversal

- Given an initialized  
`std::vector<ParticleID*> cells (M*M*M)`

- Algorithm:**

Loop over all `cells[i]`



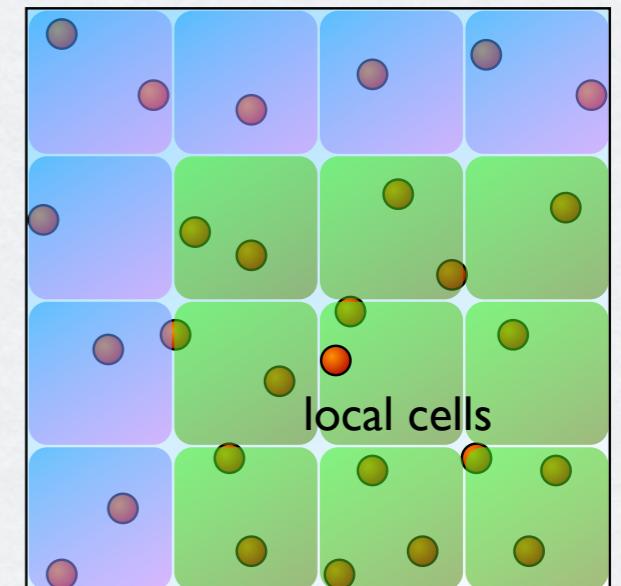
# Cell List Traversal

- Given an initialized  
`std::vector<ParticleID*> cells (M*M*M)`

- Algorithm:**

- Loop over all `cells[i]`

- Loop over all neighbors `cells[j]` (incl. `cells[i]` itself)



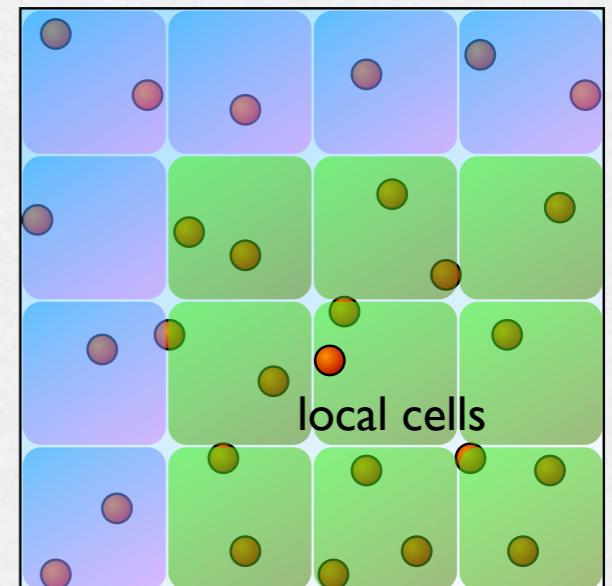
# Cell List Traversal

- Given an initialized  
`std::vector<ParticleID*> cells (M*M*M)`

- Algorithm:**

- Loop over all `cells[i]`

- Loop over all neighbors `cells[j]` (incl. `cells[i]` itself)
    - Loop over all particles `p[k]` in `cells[i]`



# Cell List Traversal

- Given an initialized  
`std::vector<ParticleID*> cells (M*M*M)`

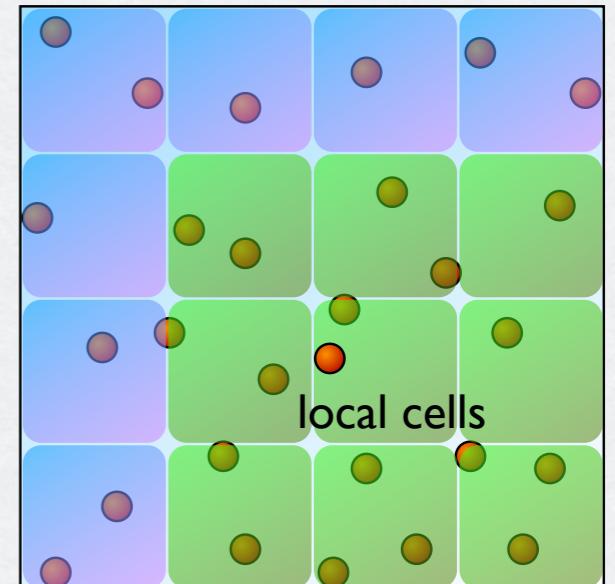
- Algorithm:**

- Loop over all `cells[i]`

- Loop over all neighbors `cells[j]` (incl. `cells[i]` itself)

- Loop over all particles `p[k]` in `cells[i]`

- Compute interaction of `p[k]` with all `p[l]` in `cells[j]` and update force of `p[k]`



# Cell List Traversal

- Given an initialized  
`std::vector<ParticleID*> cells (M*M*M)`

- Algorithm:**

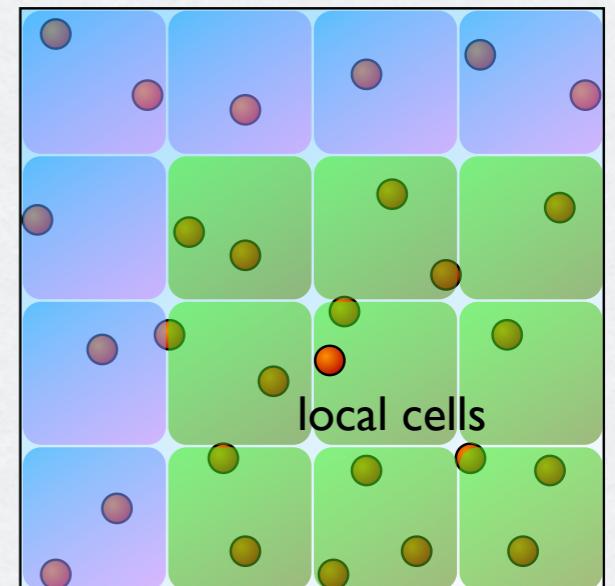
- Loop over all `cells[i]`

- Loop over all neighbors `cells[j]` (incl. `cells[i]` itself)

- Loop over all particles `p[k]` in `cells[i]`

- Compute interaction of `p[k]` with all `p[l]` in `cells[j]` and update force of `p[k]`

- Careful to skip interaction with own particle (i.e. `p[k] == p[l]`)



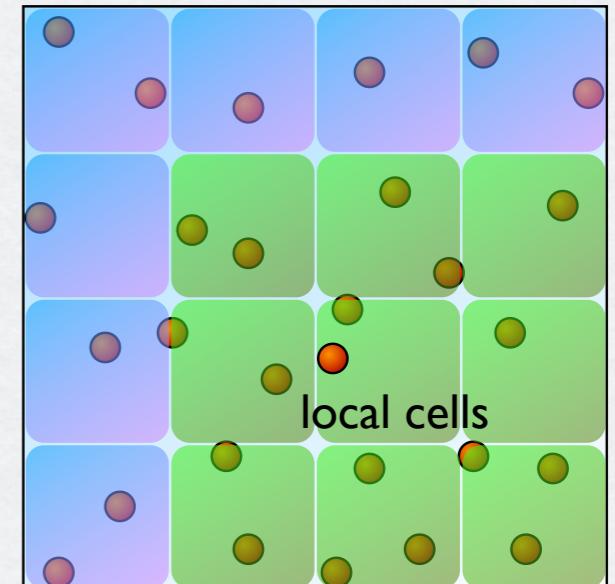
# Cell List Traversal

- Given an initialized  
`std::vector<ParticleID*> cells (M*M*M)`

- Algorithm:**

- Loop over all `cells[i]`

- Loop over all neighbors `cells[j]` (incl. `cells[i]` itself)
      - Loop over all particles `p[k]` in `cells[i]`
        - Compute interaction of `p[k]` with all `p[l]` in `cells[j]` and update force of `p[k]`
          - Careful to skip interaction with own particle (i.e. `p[k] == p[l]`)
          - Always check the cutoff distance: Cell list guarantees that all particles within the cutoff are in neighboring cells, not the other way around!

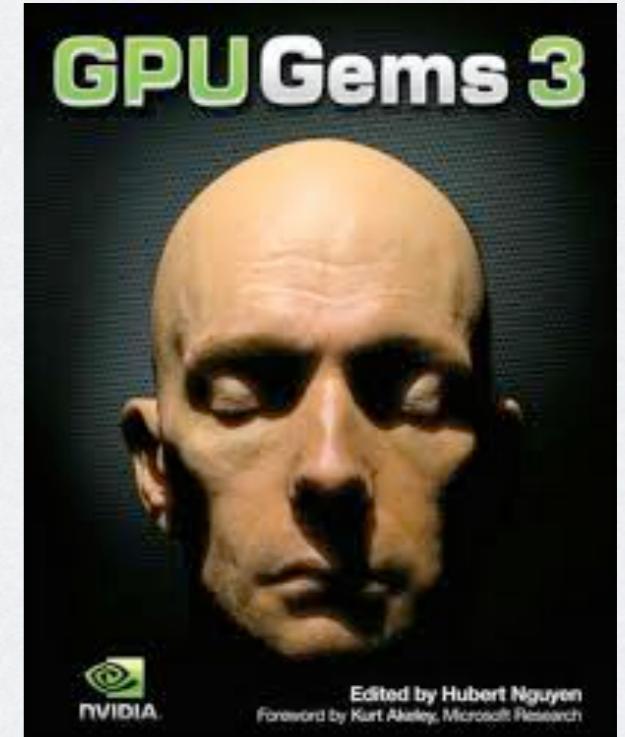


# Using Cell List

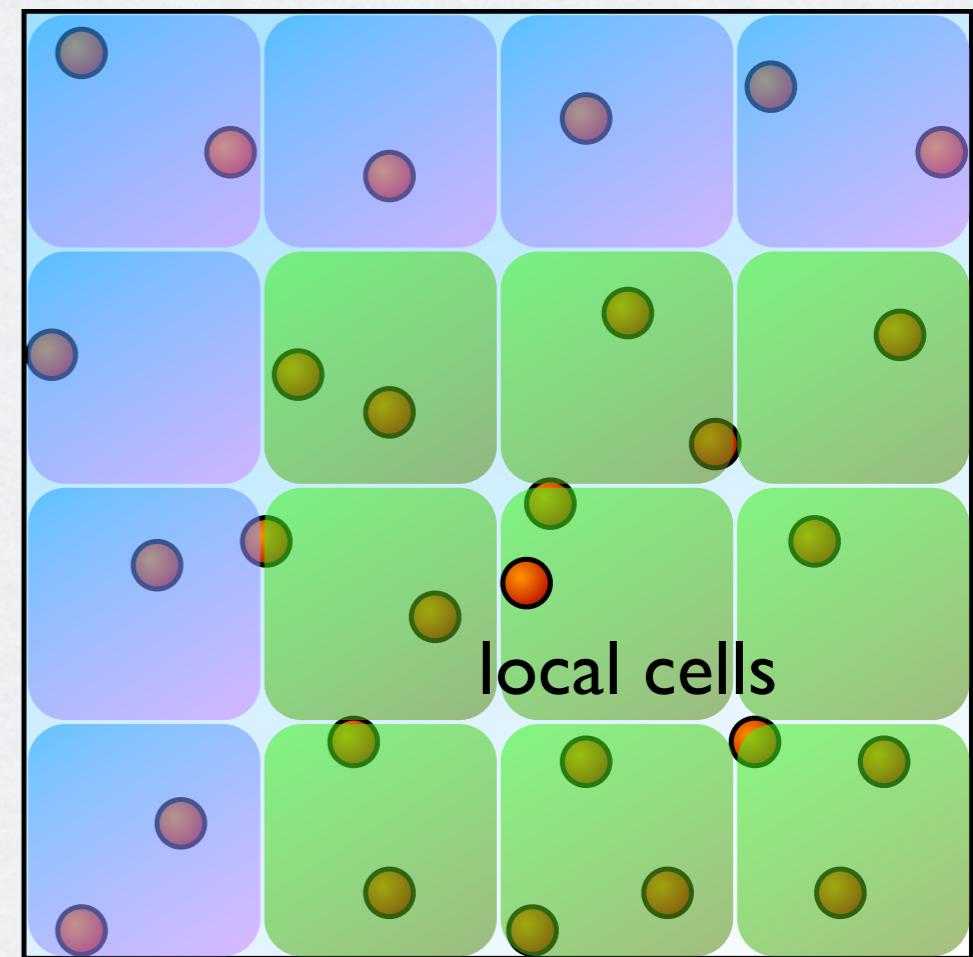
- Cell list is created at each time-step
- Cell list creation is  $\mathcal{O}(N)$ , it does not mean cheap!
- Particle position coherence can be exploited
  - Keep previous cell list
  - Refresh cell list:
    - Look for particles that moved into another cell
    - Erase the particle from the old cell
    - Insert the particle in the new cell

# Cell List Implementation

- Information taken from GPU Gems 3
  - ▶ GPU Gems 3 is a collection of state-of-the-art GPU programming examples
  - ▶ available online: [http://developer.nvidia.com/GPUGems3/gpugems3\\_ch32.html](http://developer.nvidia.com/GPUGems3/gpugems3_ch32.html)
- Broad-phase collision detection with CUDA
  - ▶ called “**Spatial Subdivision**” algorithm
- Analogous to cell-lists in MD

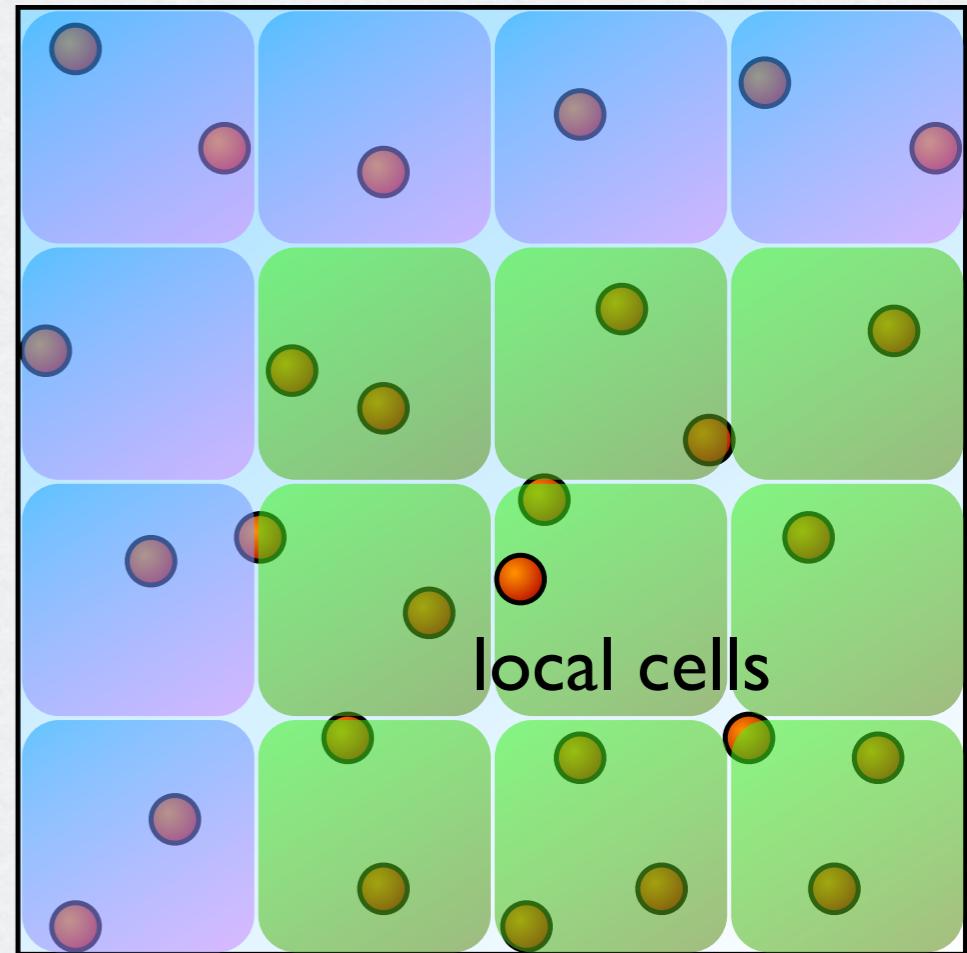


# Sequential Implementation



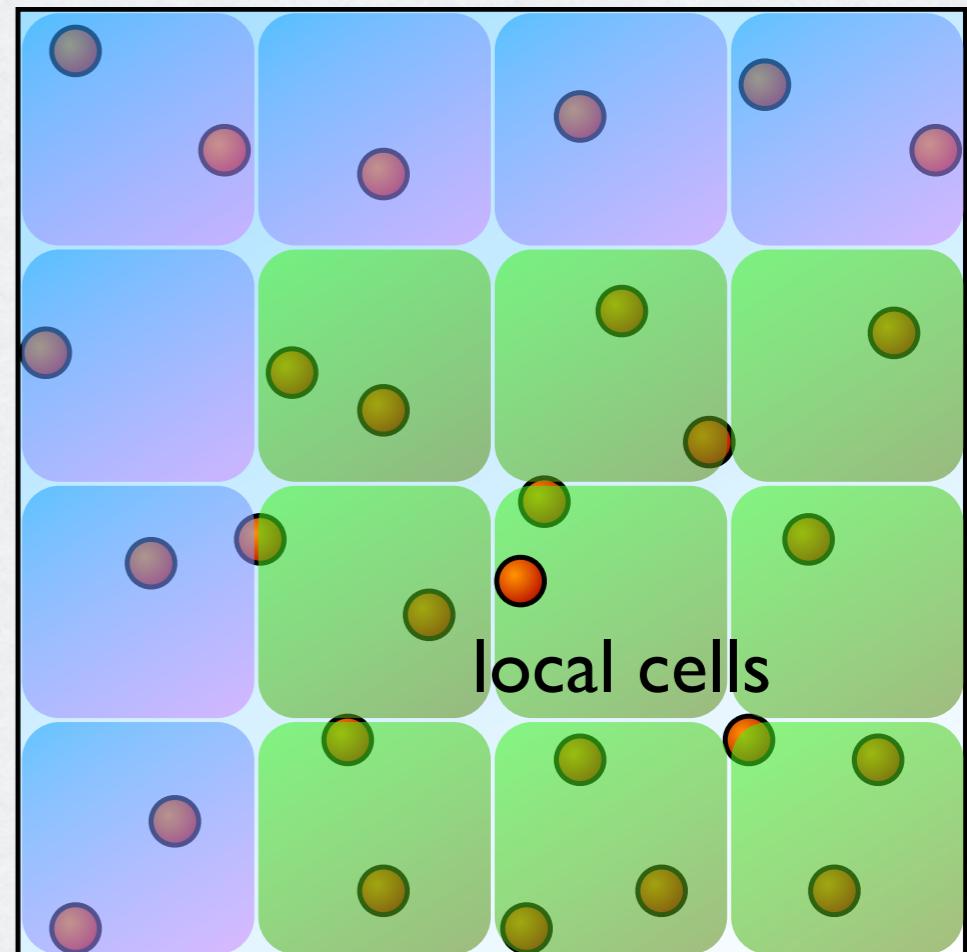
# Sequential Implementation

1. For each particle find the ID of its container cell



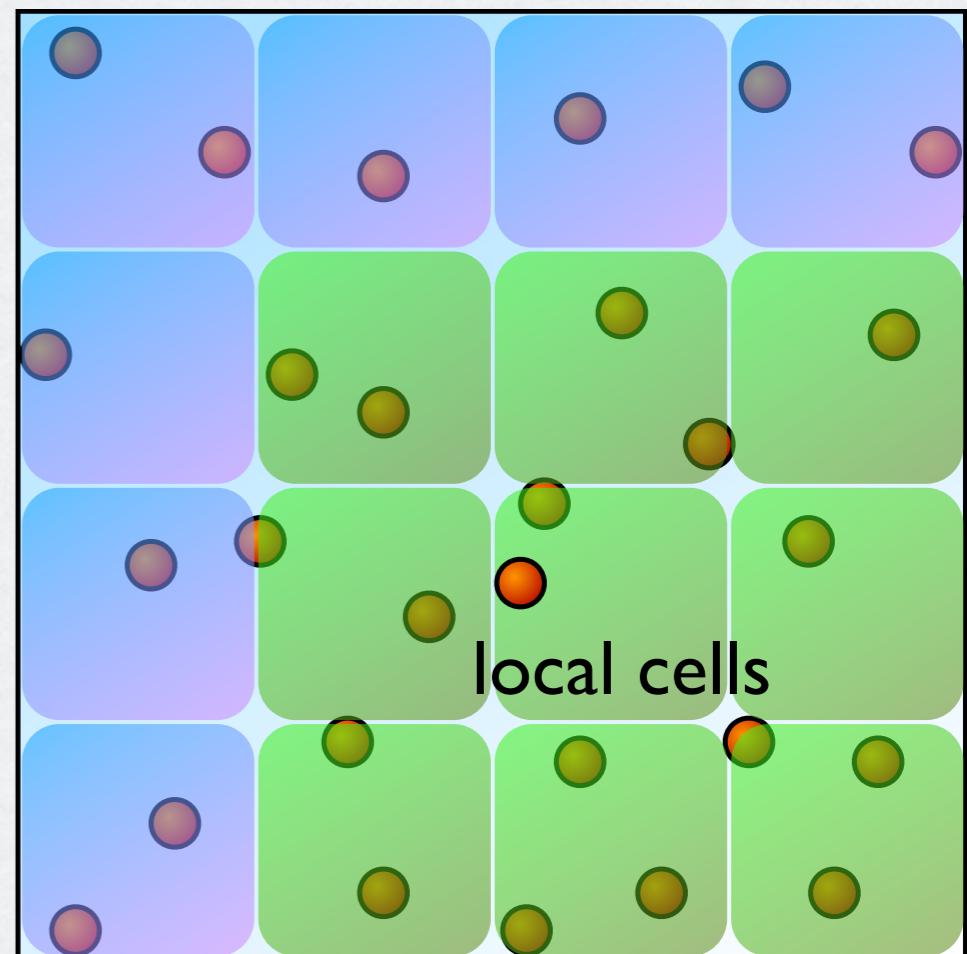
# Sequential Implementation

1. For each particle find the ID of its container cell
2. Form list of particle IDs and corresponding cell IDs



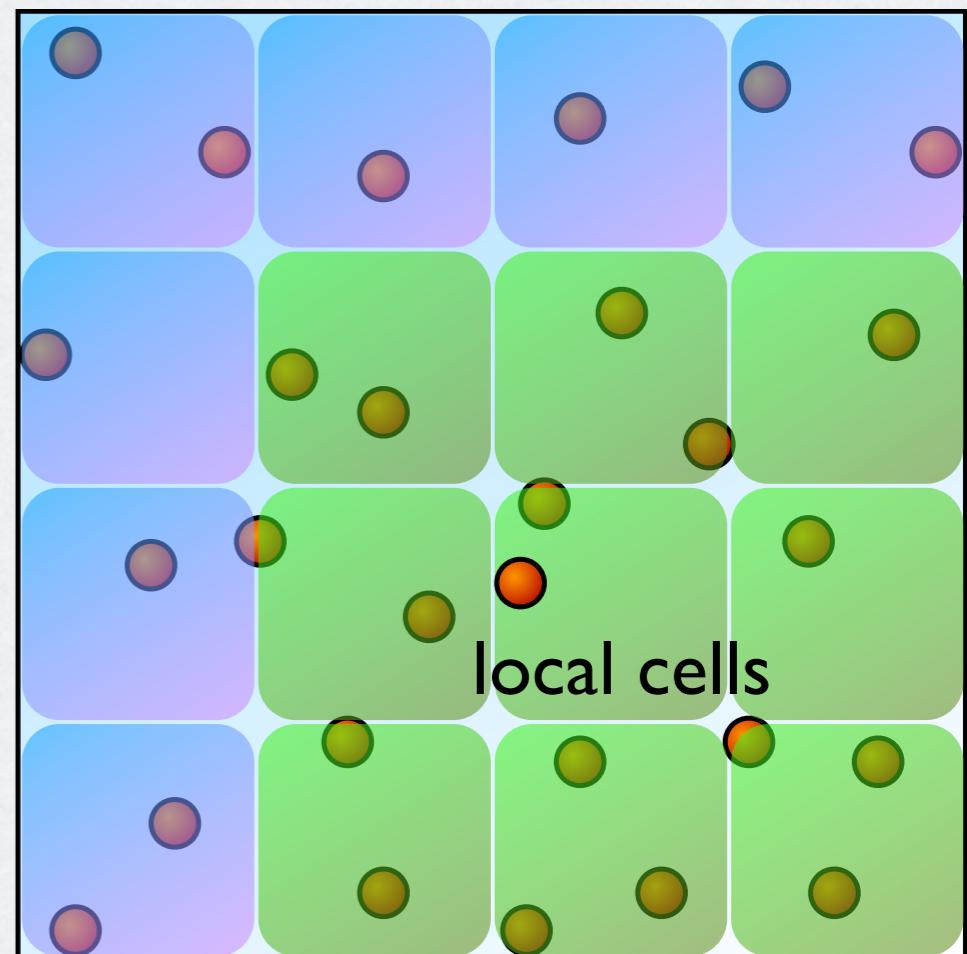
# Sequential Implementation

1. For each particle find the ID of its container cell
2. Form list of particle IDs and corresponding cell IDs
3. Sort the particle list by the cell ID



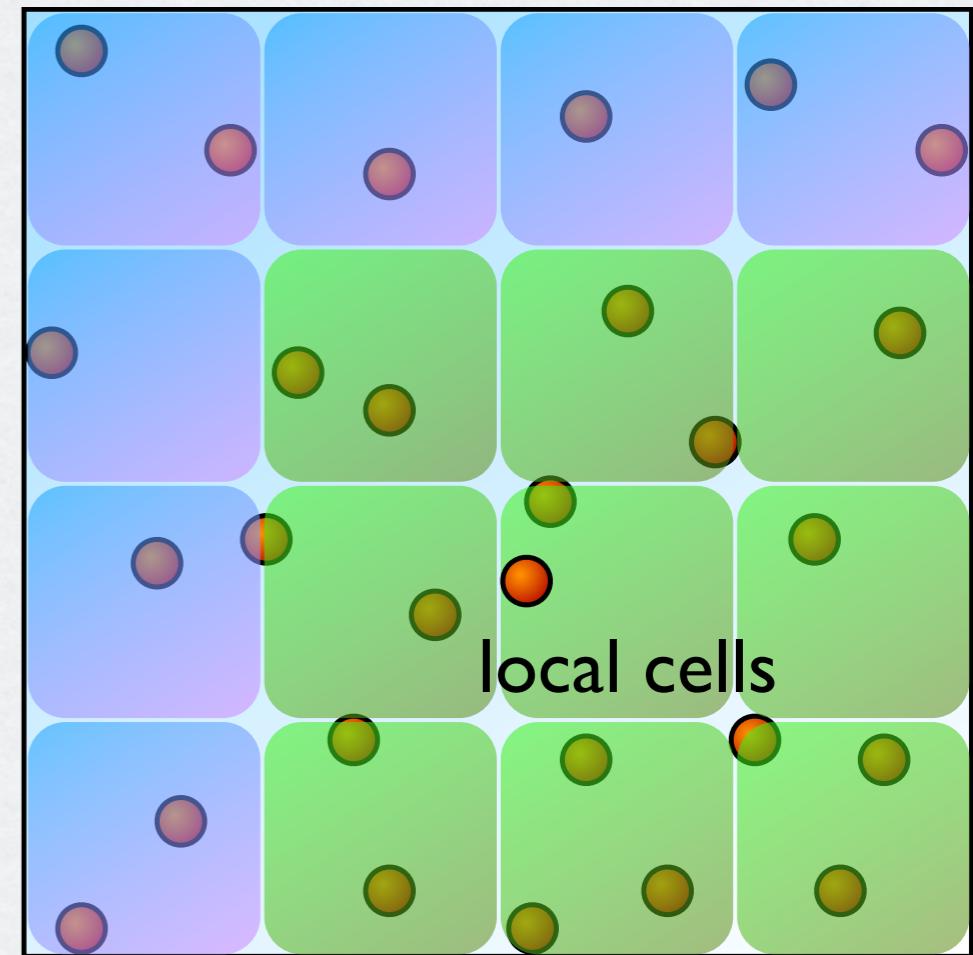
# Sequential Implementation

1. For each particle find the ID of its container cell
2. Form list of particle IDs and corresponding cell IDs
3. Sort the particle list by the cell ID
4. Process the cell list according to the cell IDs



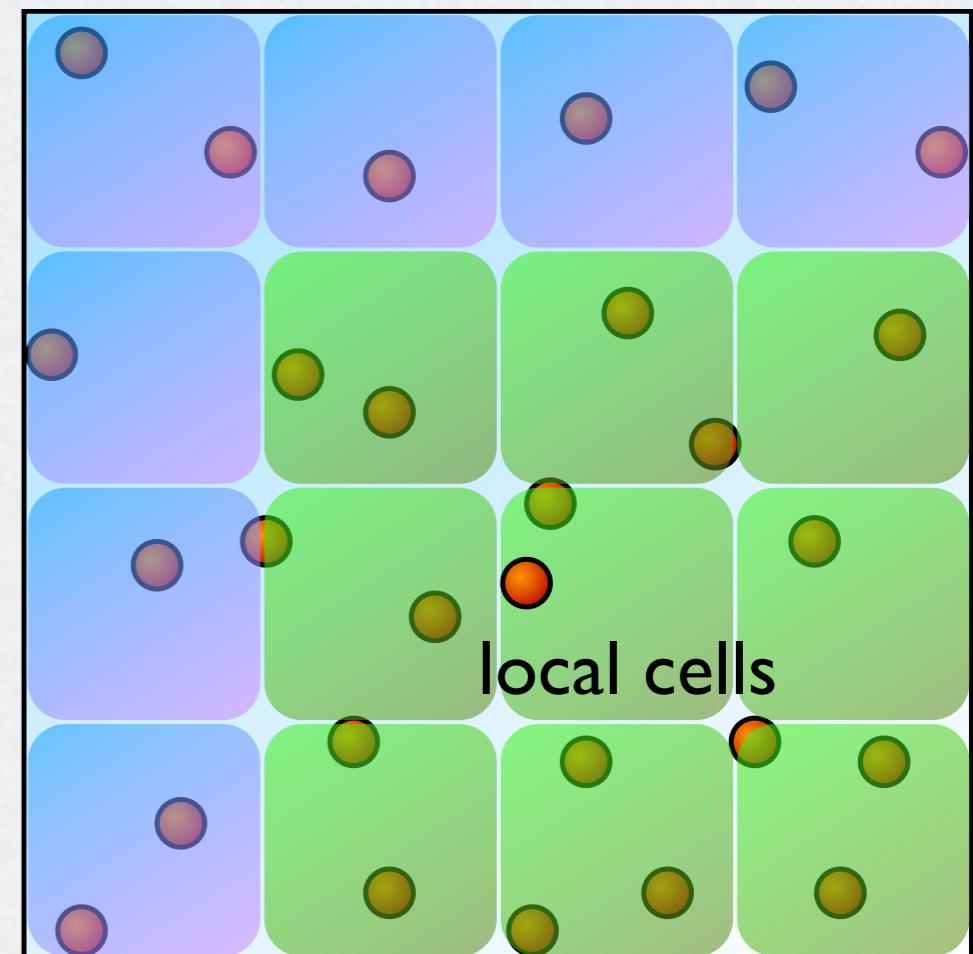
# Sequential Implementation

1. For each particle find the ID of its container cell
2. Form list of particle IDs and corresponding cell IDs
3. Sort the particle list by the cell ID
4. Process the cell list according to the cell IDs
5. Perform interactions between all particles with the same cell ID and corresponding neighboring cells



# Sequential Implementation

1. For each particle find the ID of its container cell
2. Form list of particle IDs and corresponding cell IDs
3. Sort the particle list by the cell ID
4. Process the cell list according to the cell IDs
5. Perform interactions between all particles with the same cell ID and corresponding neighboring cells

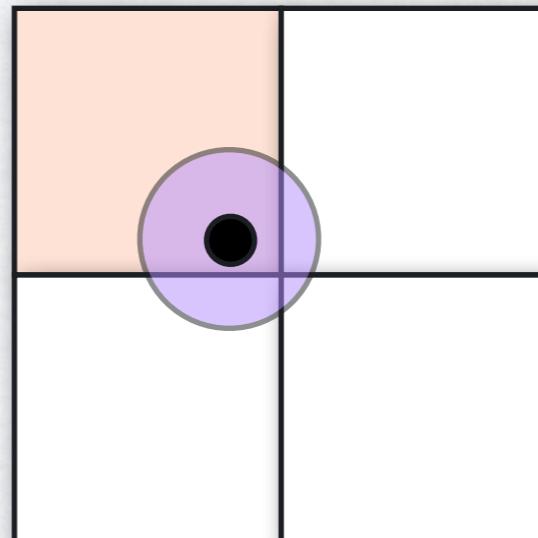


→ How to parallelize ?

# Data Structures

For each particle

particle ID  
container cell ID



particle-ID	0	1	2	3	4	5	6	7	8	9	10
cell-ID	6	5	8	6	3	4	4	9	7	10	1

The array contains pairs of (particle-ID, cell-ID)

# Constructing the 1D Cell-ID Array

Container cell ID H (hash) for particle i  
with position  $(x_i, y_i, z_i)$  in  $]-1, 1[^3$

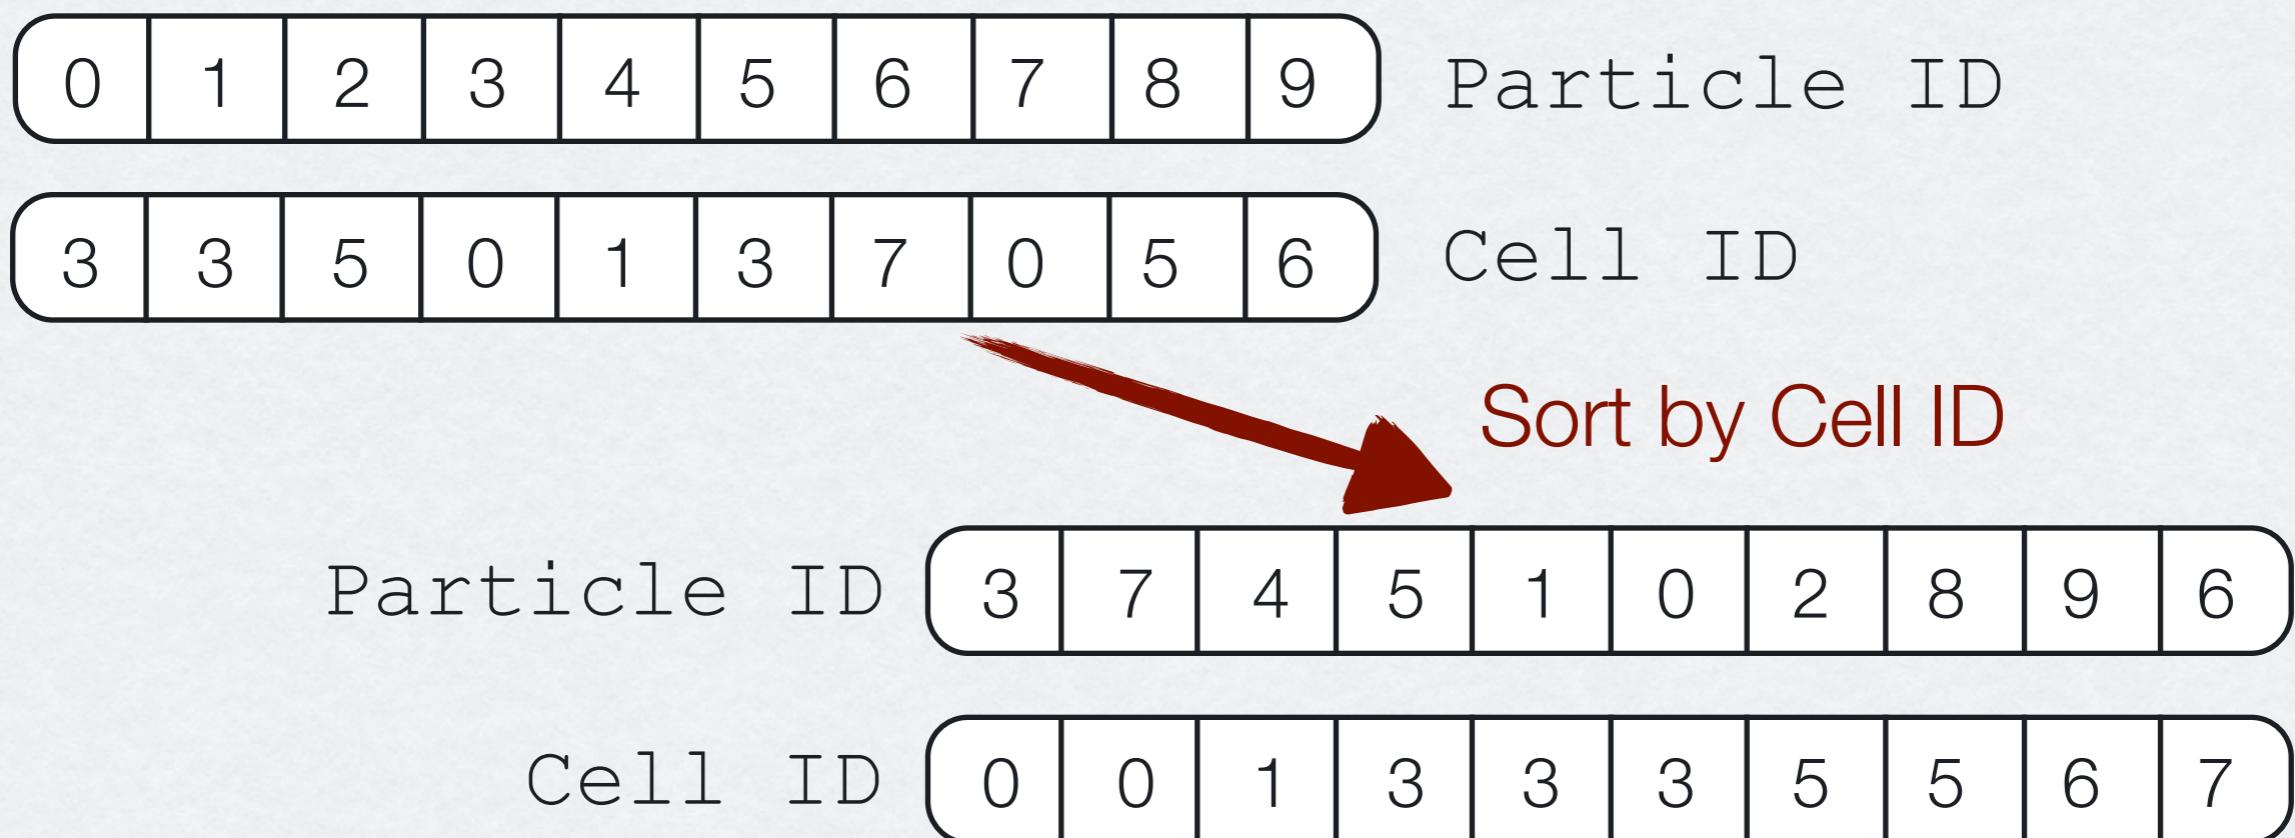
$$\text{cell ID} = \left\lfloor M \frac{x_i + 1}{2} \right\rfloor + M \left\lfloor M \frac{y_i + 1}{2} \right\rfloor + M^2 \left\lfloor M \frac{z_i + 1}{2} \right\rfloor$$

M: number of cells in a given direction



# Sorting the Cell ID Array

- Cell ID array is sorted by cell ID
- Particle ID array is also sorted by cell ID
- Cell IDs may appear multiple times
- Radix sort:  $\mathcal{O}(N)$  sorting



# Sorting with C++ - std::sort()

simple:

```
#include <algorithm>
#include <utility>
#include <vector>

// Custom comparator for pair,
// compares only by the first field
struct compare_first
{
    template <typename T>
    bool operator() (const T& a, const T& b)
    {
        return a.first < b.first;
    }
};

// Vector storing cellids as first member of pair
// and particleids as second member
std::vector<std::pair<unsigned int, unsigned int> > cellid_particleid(n_particles);

// Fill the vector with cellids
// x, y, z - arrays storing coordinates of particles
// get_cell_index() returns the index of enclosing cell
for (int i=0; i<n_particles; i++)
    cellid_particleid[i] = std::make_pair(get_cell_index(x[i], y[i], z[i]), i);

std::sort(cellid_particleid.begin(), cellid_particleid.end(), compare_first());
```

BUT: this sorting is  $\mathcal{O}(N \log N)$ !

# Creating the Interaction Cell List

# Creating the Interaction Cell List

- *Interaction cell (non-empty cell)*: contains 1 or more particles

# Creating the Interaction Cell List

- *Interaction cell* (*non-empty cell*): contains 1 or more particles
- *Interaction cell list*:
  - created from the sorted cell ID array
  - the sorted cell ID array is scanned for **changes of cell ID**
  - **changes** mark the end of one cell and the beginning of another.

# Creating the Interaction Cell List

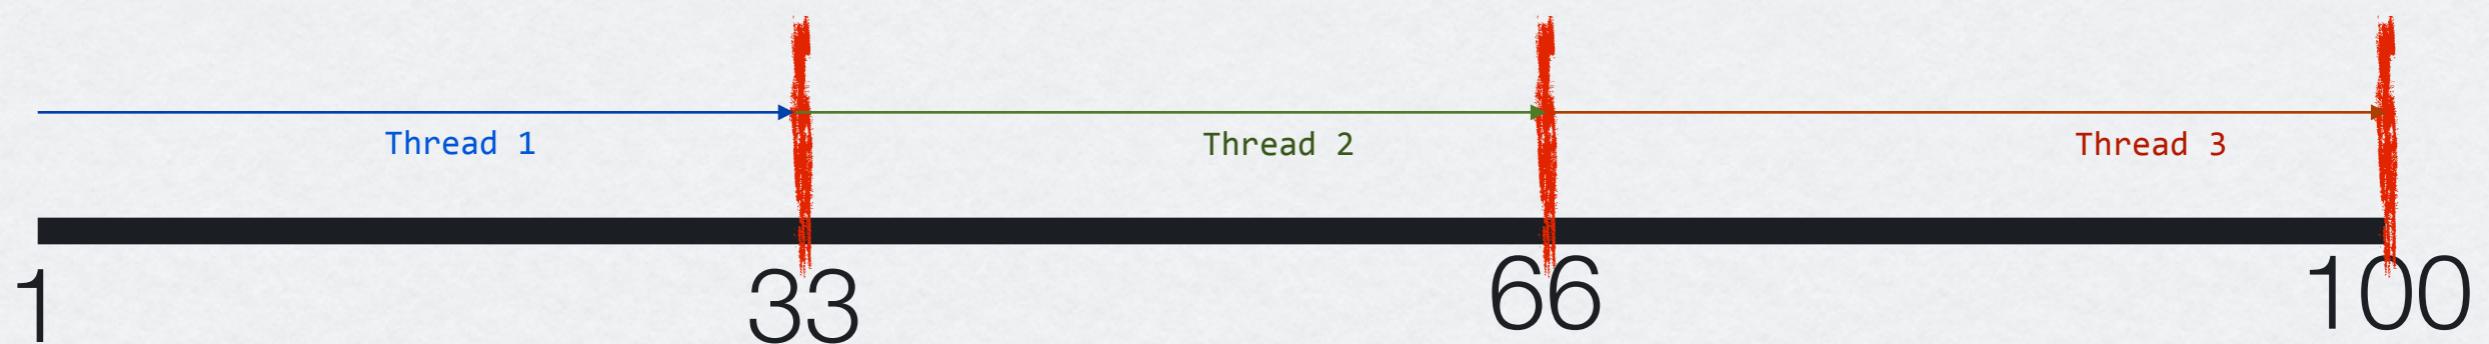
- *Interaction cell* (*non-empty cell*): contains 1 or more particles
- *Interaction cell list*:
  - created from the sorted cell ID array
  - the sorted cell ID array is scanned for **changes of cell ID**
  - **changes** mark the end of one cell and the beginning of another.
- *Parallelization*: each thread is assigned a (roughly equal) segment of the cell ID array to scan for changes

# Creating the Interaction Cell List



- **Parallelization:** threads assigned to portions of the 1D array.

# Creating the Interaction Cell List



- **Parallelization:** threads assigned to portions of the 1D array.
- **HOW TO:**
  - each thread scans past the end point of its pre-assigned segment and stops when it finds a change in cell-ID (except for the last thread)
  - each thread skips its first change as it has been found by a thread assigned to the preceding segment(except for the first thread)

# Creating the Interaction Cell List

- The sorted 1D cell ID array is scanned twice

# Creating the Interaction Cell List

- The sorted 1D cell ID array is scanned twice

- **First scan**

- count the number of particles within each interaction cell
- convert them into **offsets** in the 1D array ( through a parallel prefix sum in several passes)

# Creating the Interaction Cell List

- The sorted 1D cell ID array is scanned twice

- **First scan**

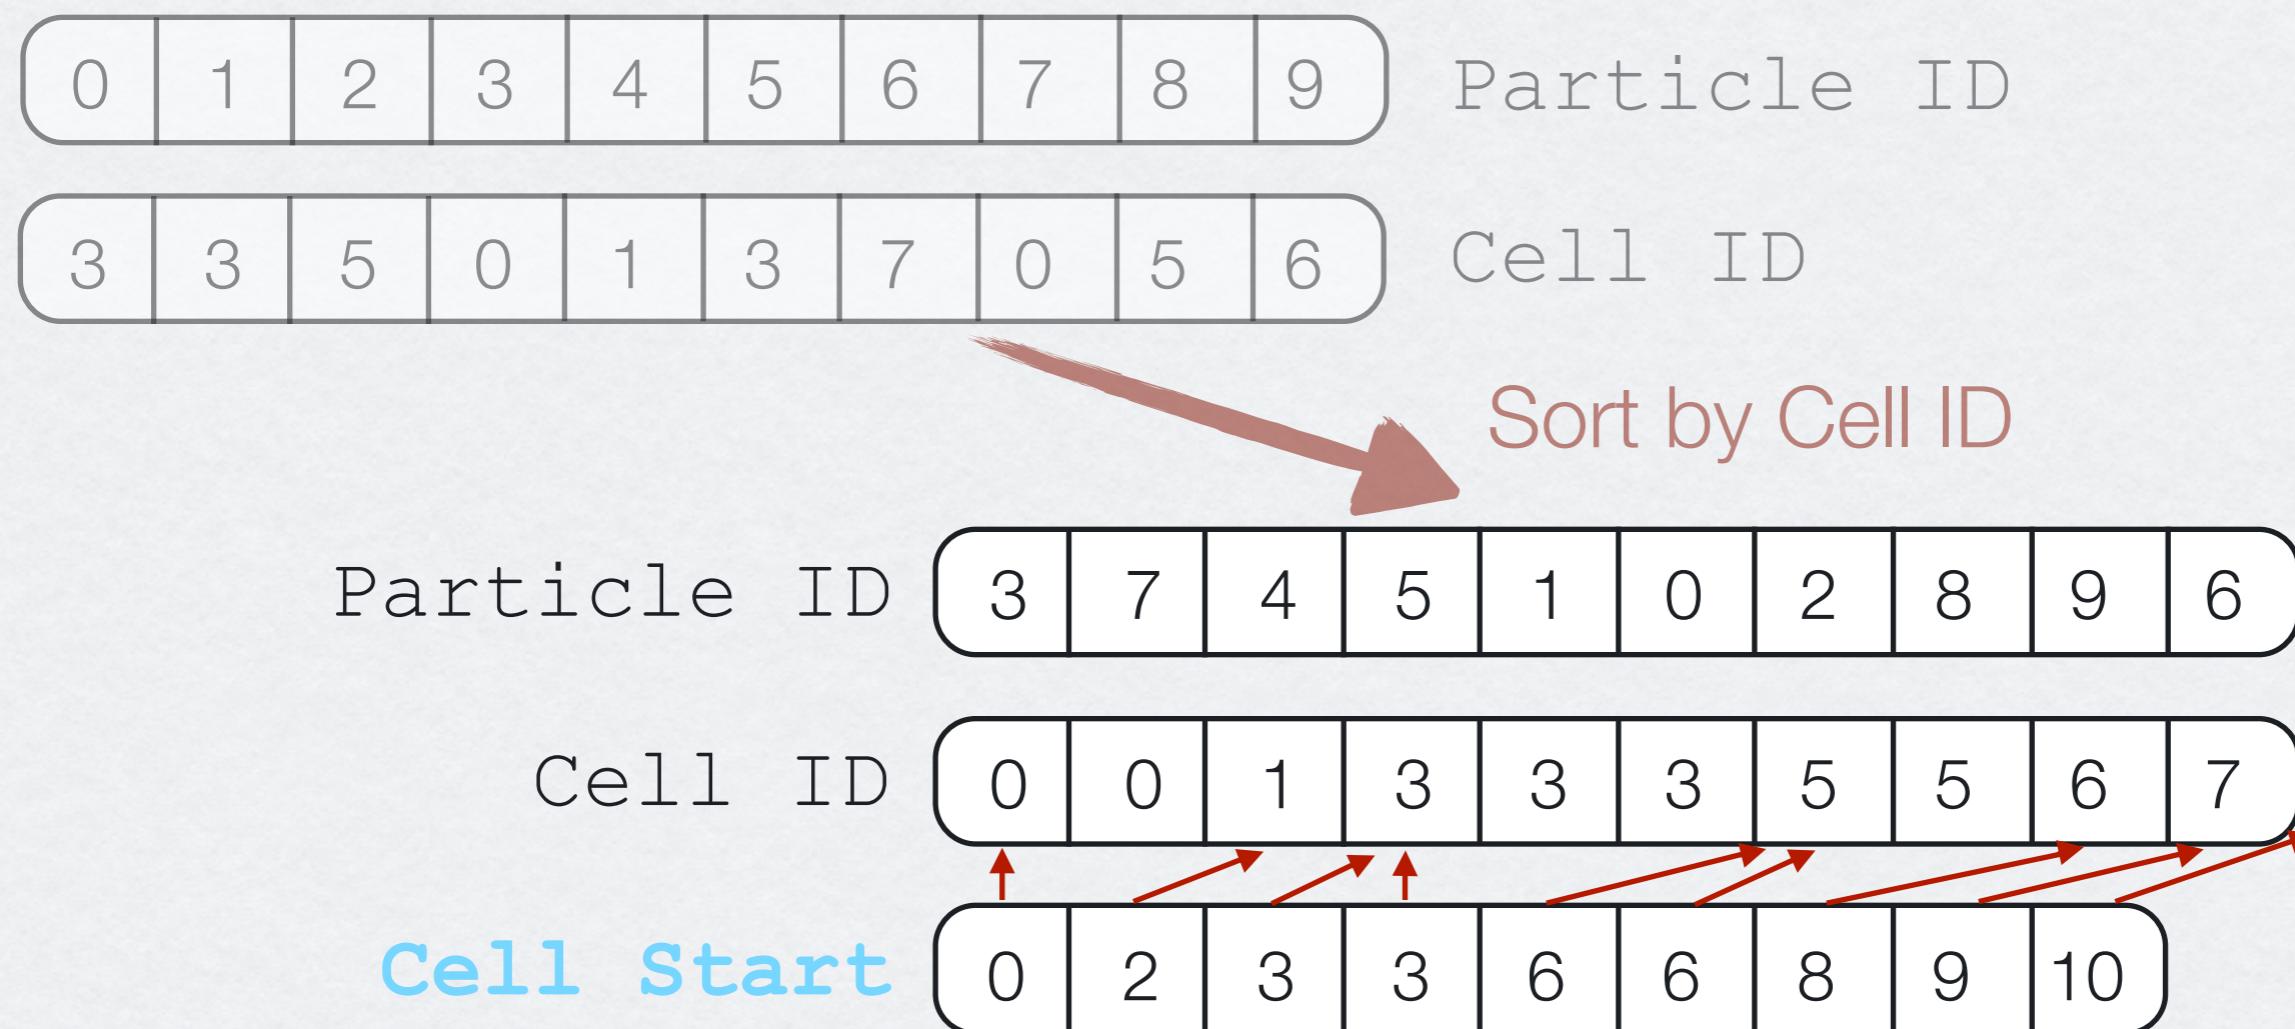
- count the number of particles within each interaction cell
- convert them into **offsets** in the 1D array (through a parallel prefix sum in several passes)

- **Second scan**

- create entries for each interaction cell in a **new array (“Cell Start”)**:
- **This array tells at which index in the 1D cell-ID array starts each cell**

# Sorting the Cell ID Array

- Cell Start array marks the first particle in Particle ID array that is contained by the cell
- The size of Cell Start array is the number of cells plus one
- The last element always marks the end of Particle ID array



# Cell Start Array

In C++ Cell Start calculation is straightforward

```
int current_cell = -1;
for (int i = 0; i < n_particles; i++)
{
    for (int j = current_cell; j < cellids[i]; j++)
    {
        start[j+1] = i;
    }
    current_cell = cellids[i];
}

for (int j = current_cell; j <= n_cells; j++)
start[j+1] = n_particles;
```

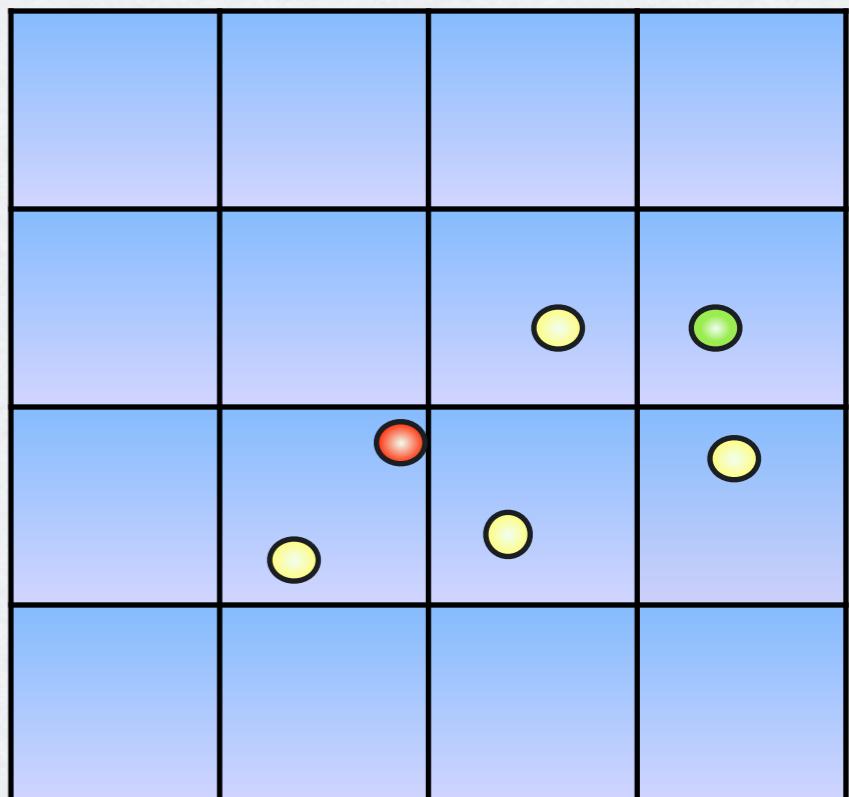
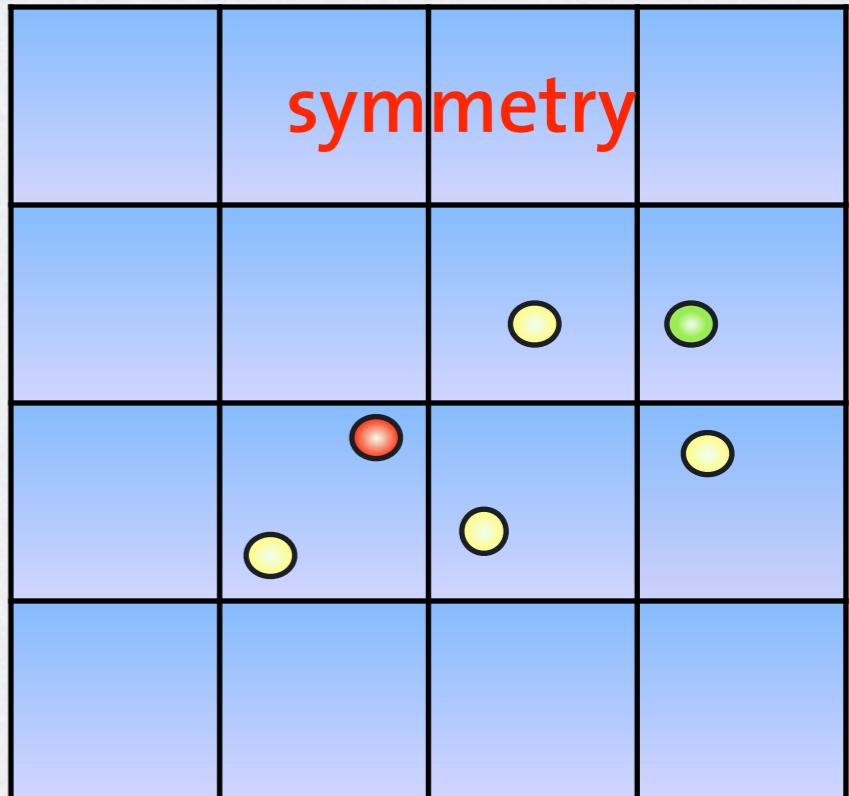
# Traversing the list

- After the creation of the list all the **particle neighbors are found** by:
  - **visiting all the cells that are adjacent to the cell containing the target particle**
- This make up a total of 9/27 (2D/3D) cells to look up
- Additional care has to be taken for the cells on the border of domain.  
Special case : Periodic Boundary Conditions.

# Parallelization over Particles

→ one thread processes one particle

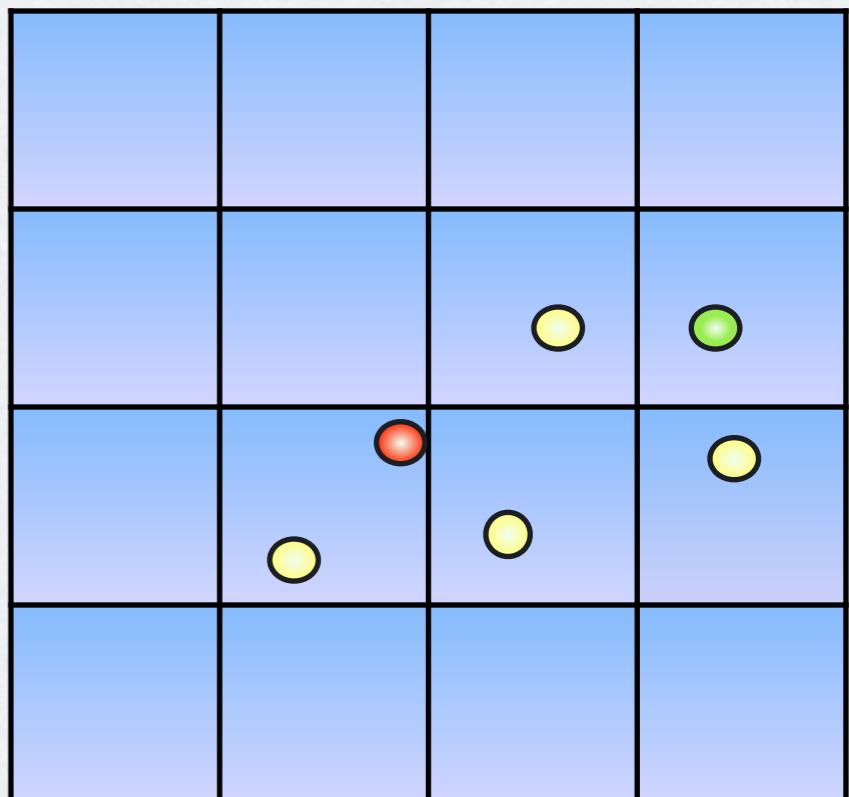
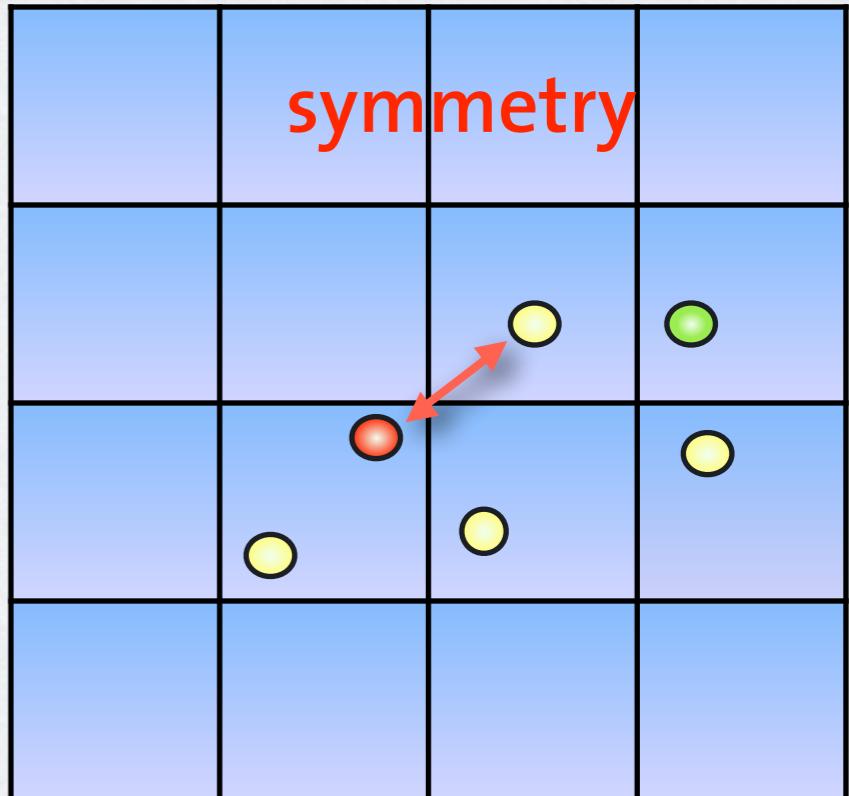
- particle of thread 1
- particle of thread 2
- other particles



# Parallelization over Particles

→ one thread processes one particle

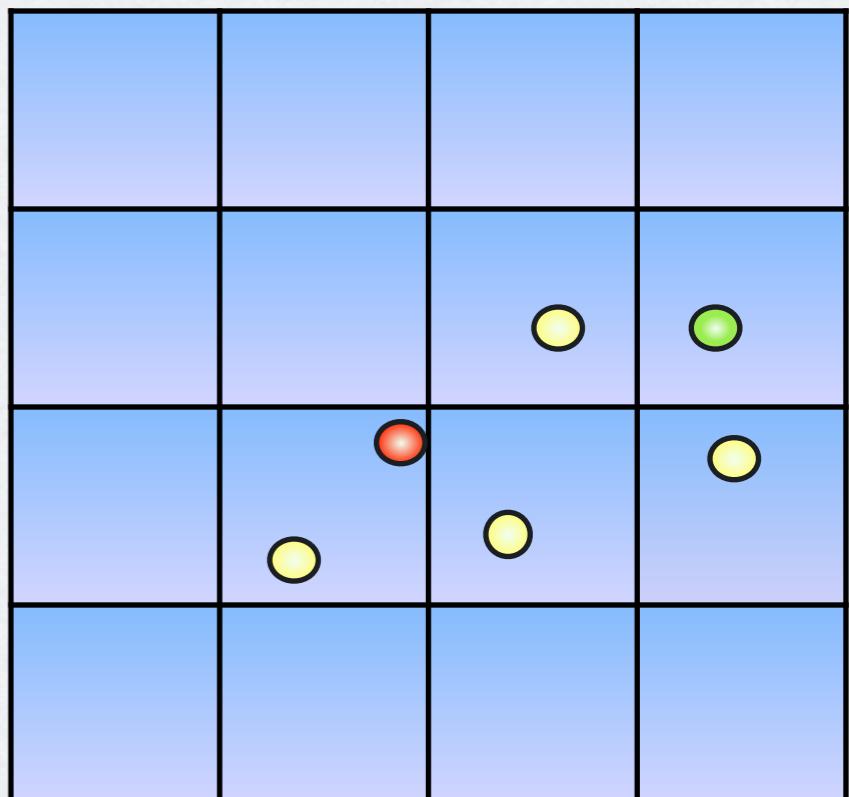
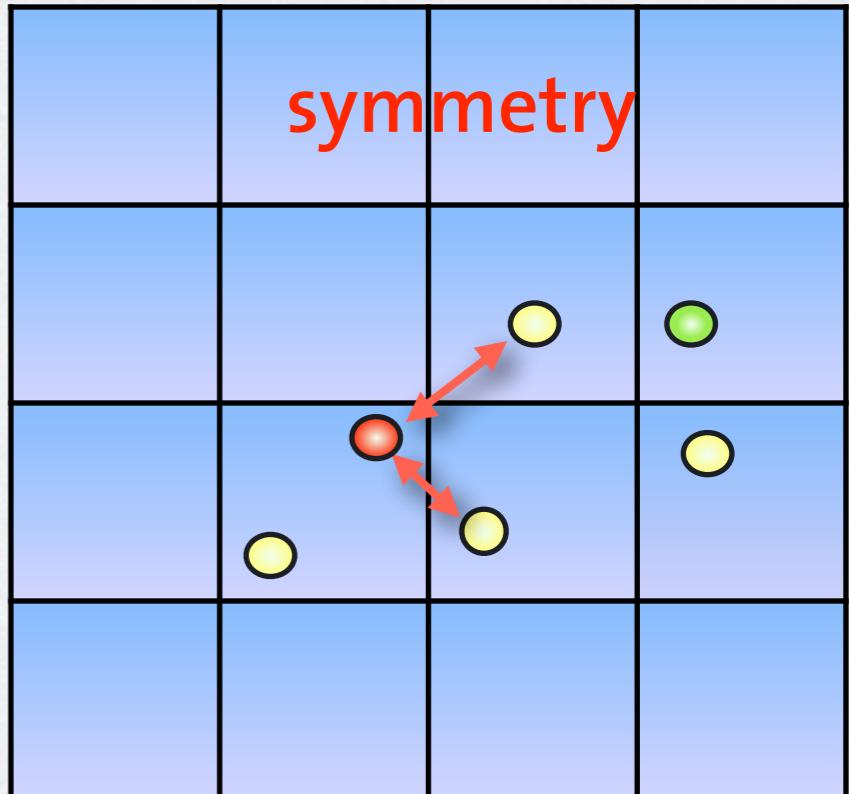
- particle of thread 1
- particle of thread 2
- other particles



# Parallelization over Particles

→ one thread processes one particle

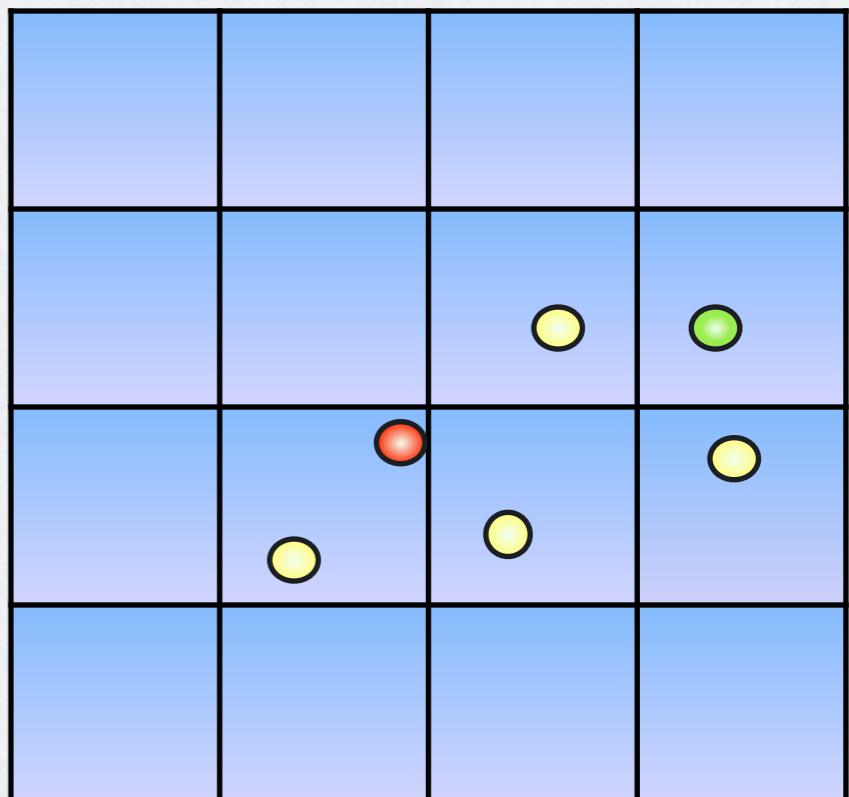
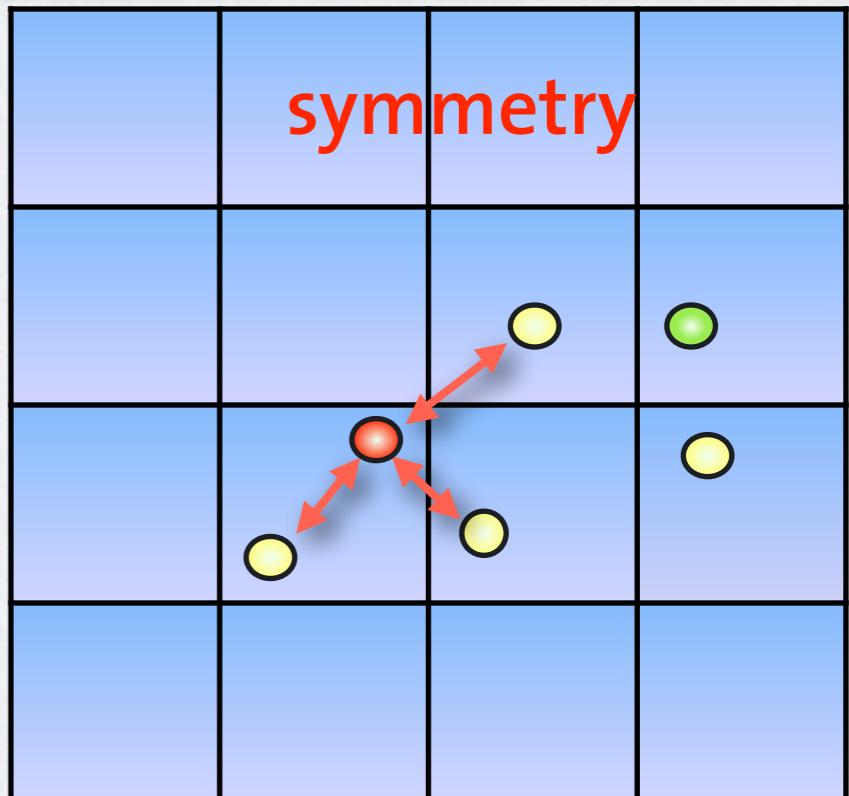
- particle of thread 1
- particle of thread 2
- other particles



# Parallelization over Particles

→ one thread processes one particle

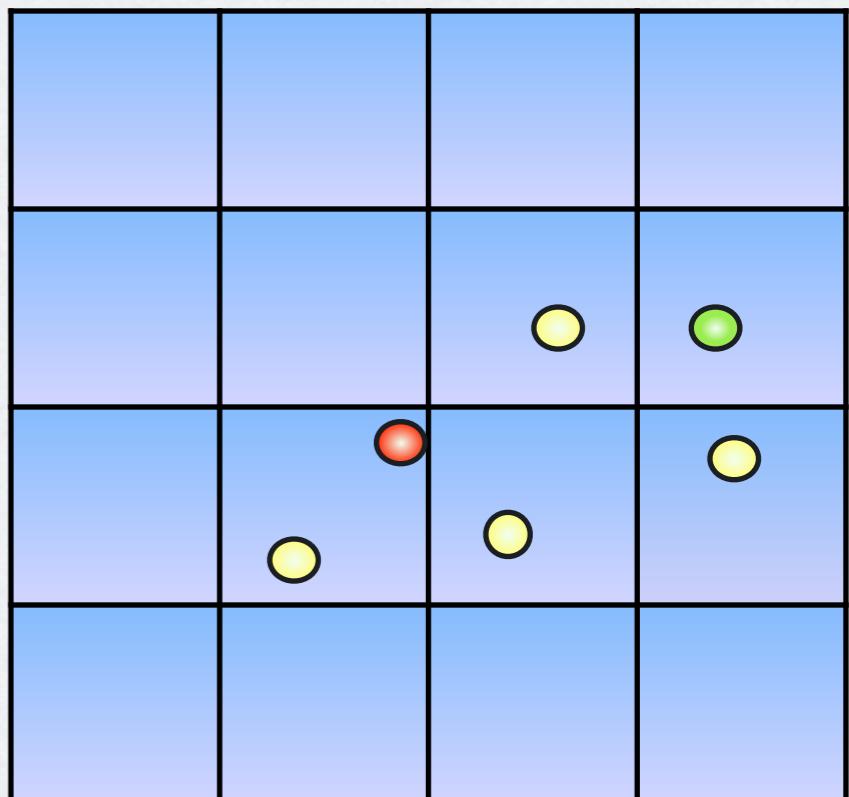
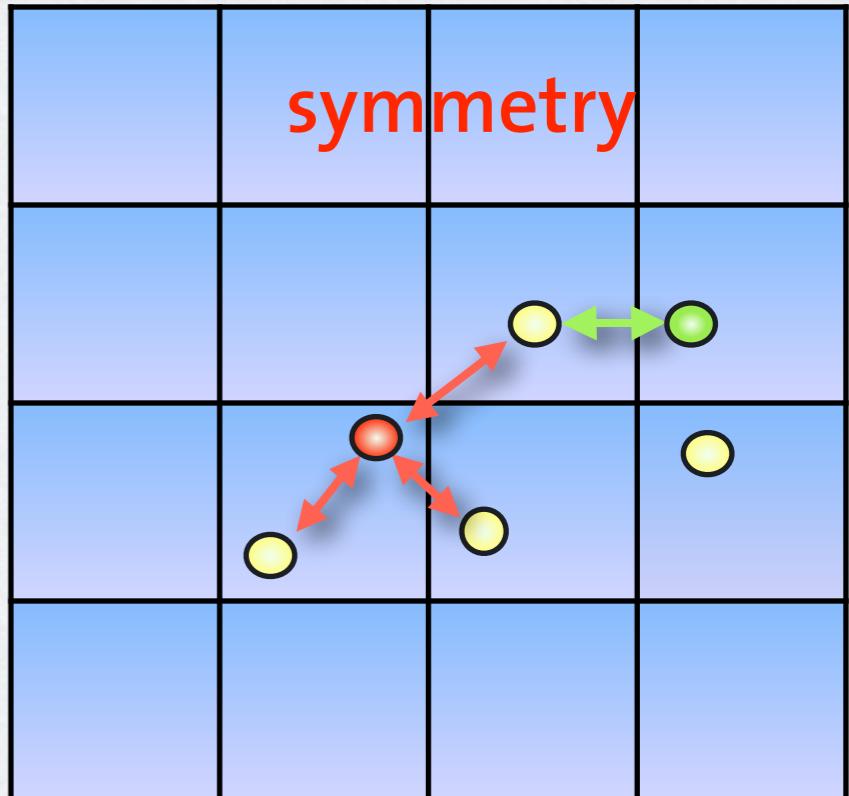
- particle of thread 1
- particle of thread 2
- other particles



# Parallelization over Particles

→ one thread processes one particle

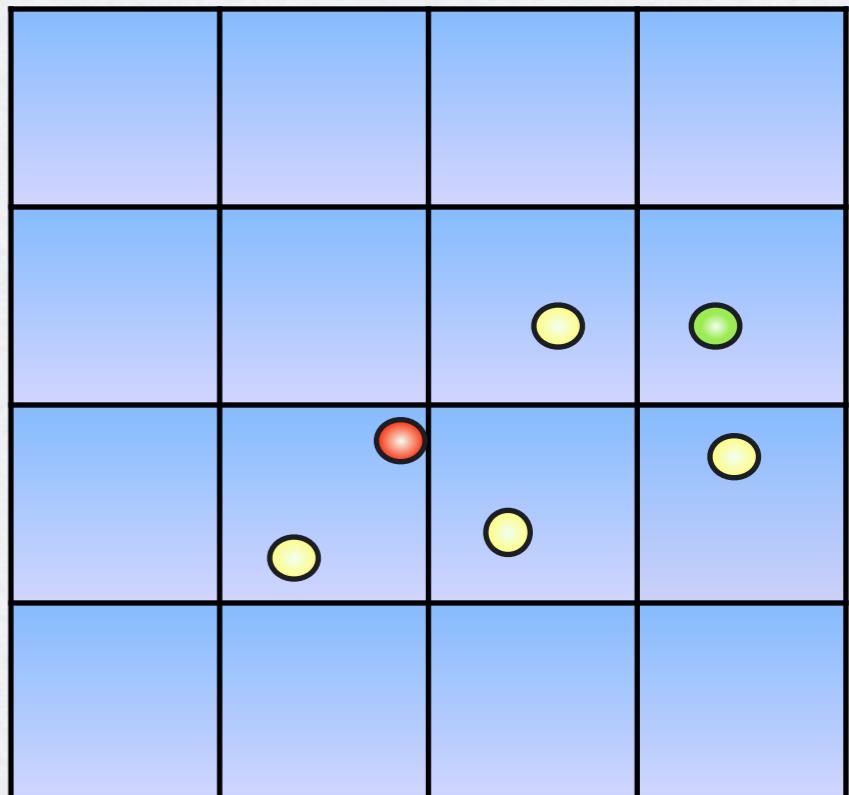
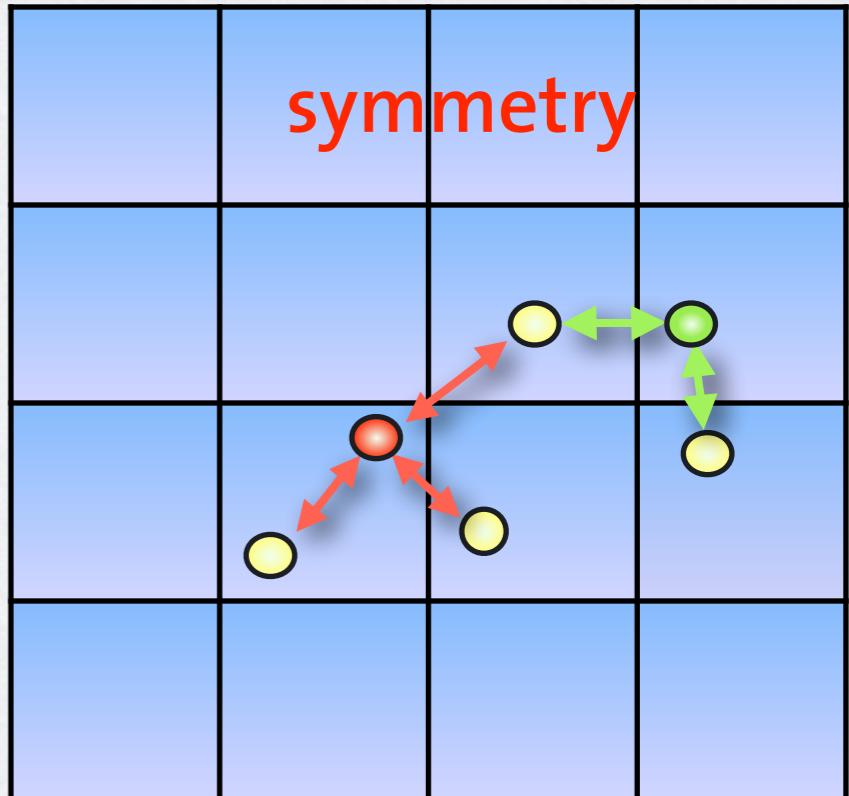
- particle of thread 1
- particle of thread 2
- other particles



# Parallelization over Particles

→ one thread processes one particle

- particle of thread 1
- particle of thread 2
- other particles



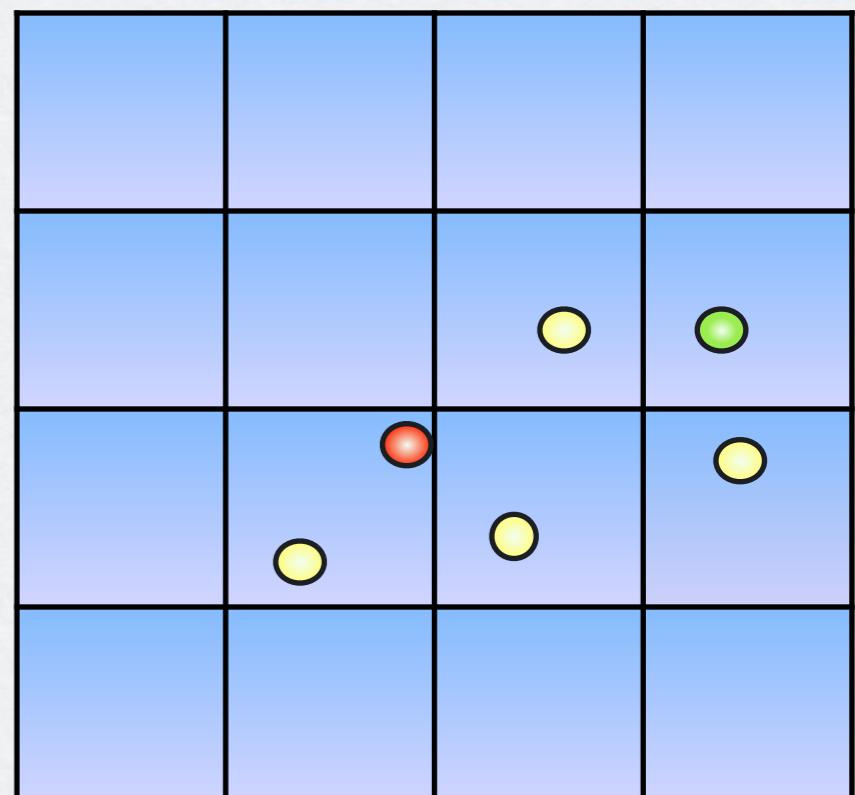
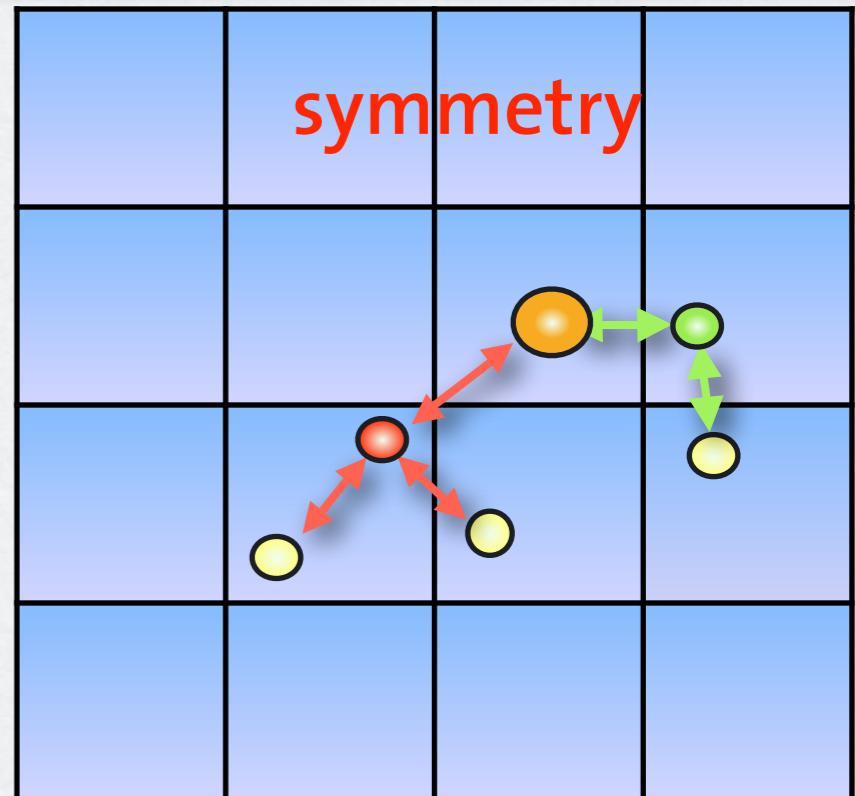
# Parallelization over Particles

→ one thread processes one particle

- particle of thread 1
- particle of thread 2
- other particles

## PROBLEMS

- race condition!



# Parallelization over Particles

→ one thread processes one particle

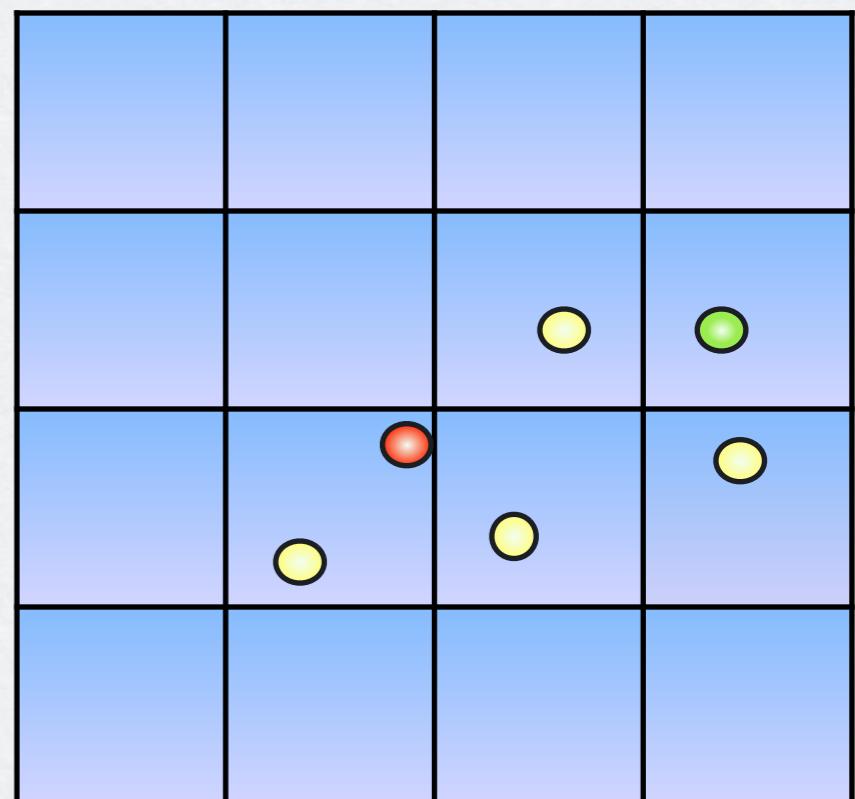
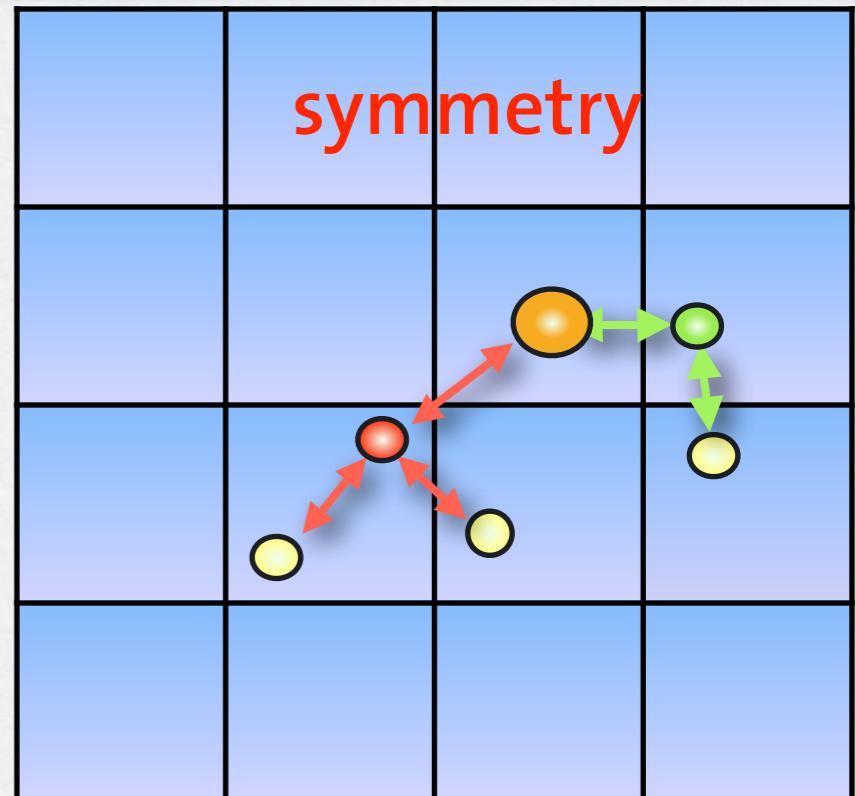
- particle of thread 1
- particle of thread 2
- other particles

## PROBLEMS

- race condition!

- load imbalance

(some threads do more work than others)



# Parallelization over Particles

→ one thread processes one particle

- particle of thread 1
- particle of thread 2
- other particles

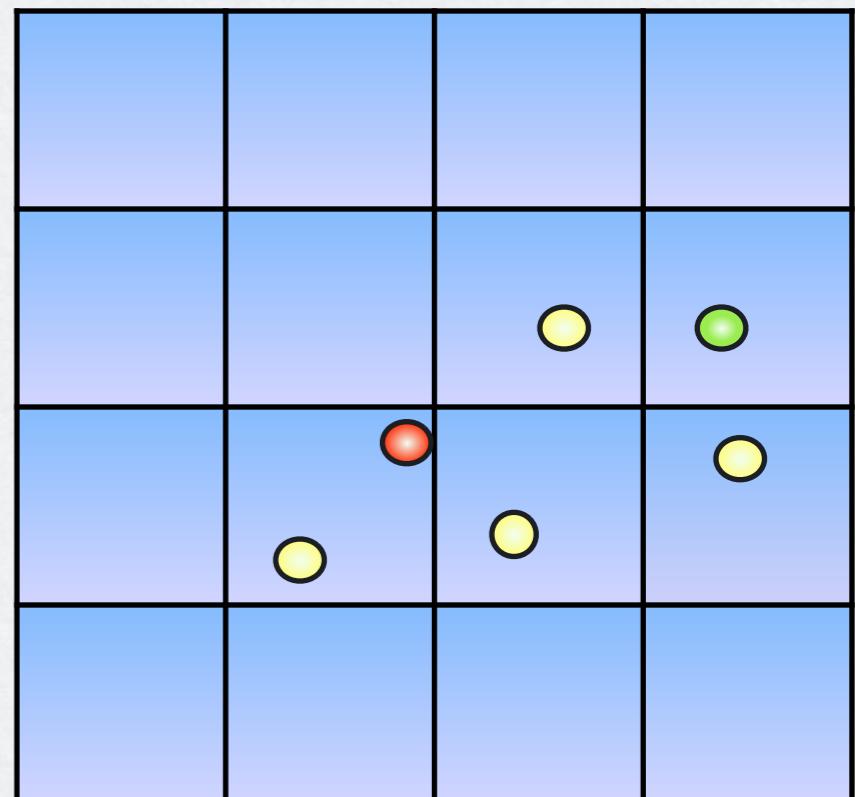
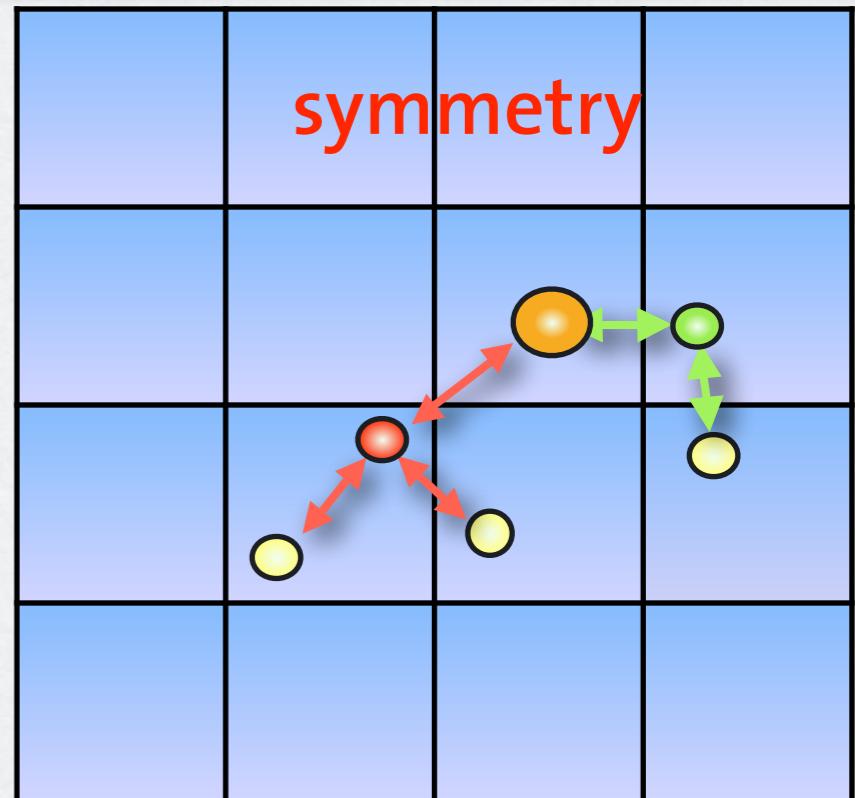
## PROBLEMS

- race condition!

- load imbalance

(some threads do more work than others)

## HOW TO FIX IT:



# Parallelization over Particles

→ one thread processes one particle

- particle of thread 1
- particle of thread 2
- other particles

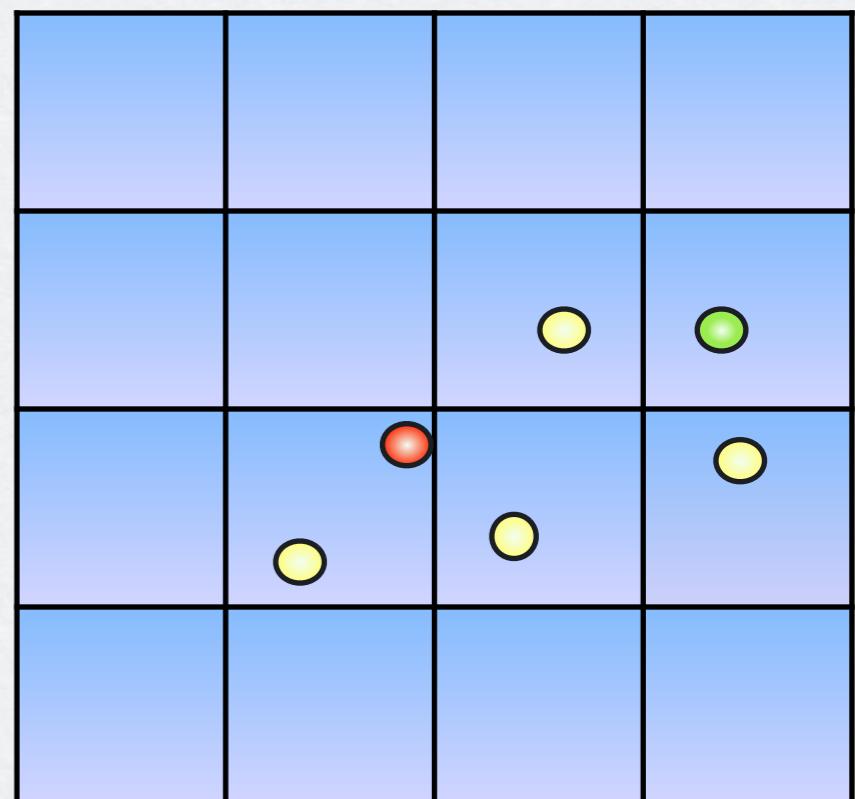
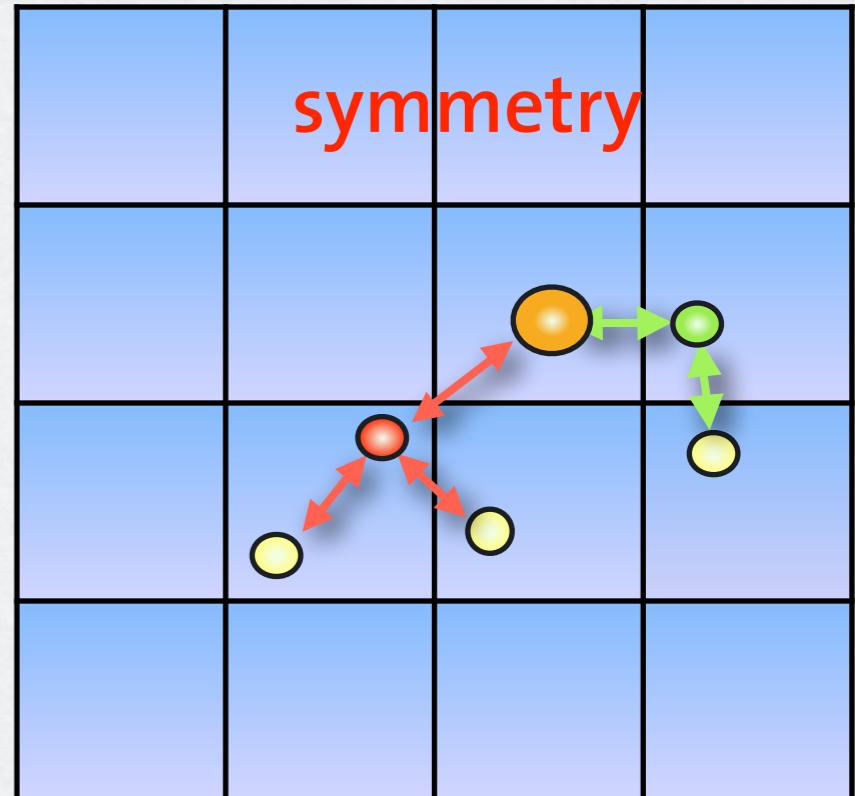
## PROBLEMS

- race condition!
- load imbalance  
(some threads do more work than others)

## HOW TO FIX IT:

- remove symmetry to avoid race conditions :  
each particle is updated separately

- Tasking



# Parallelization over Particles

→ one thread processes one particle

- particle of thread 1
- particle of thread 2
- other particles

## PROBLEMS

- race condition!

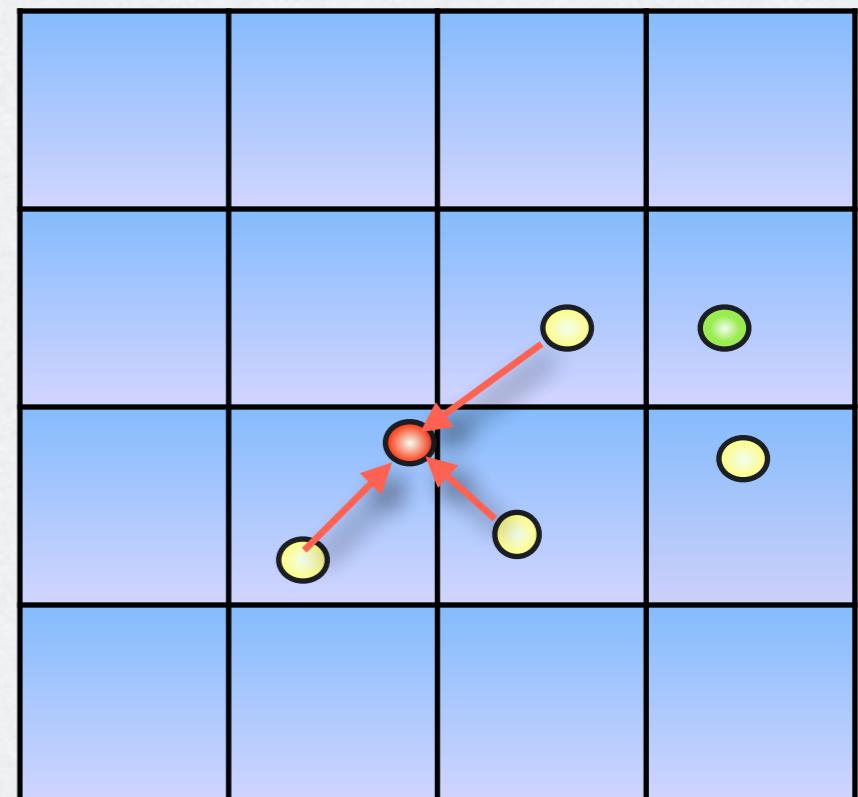
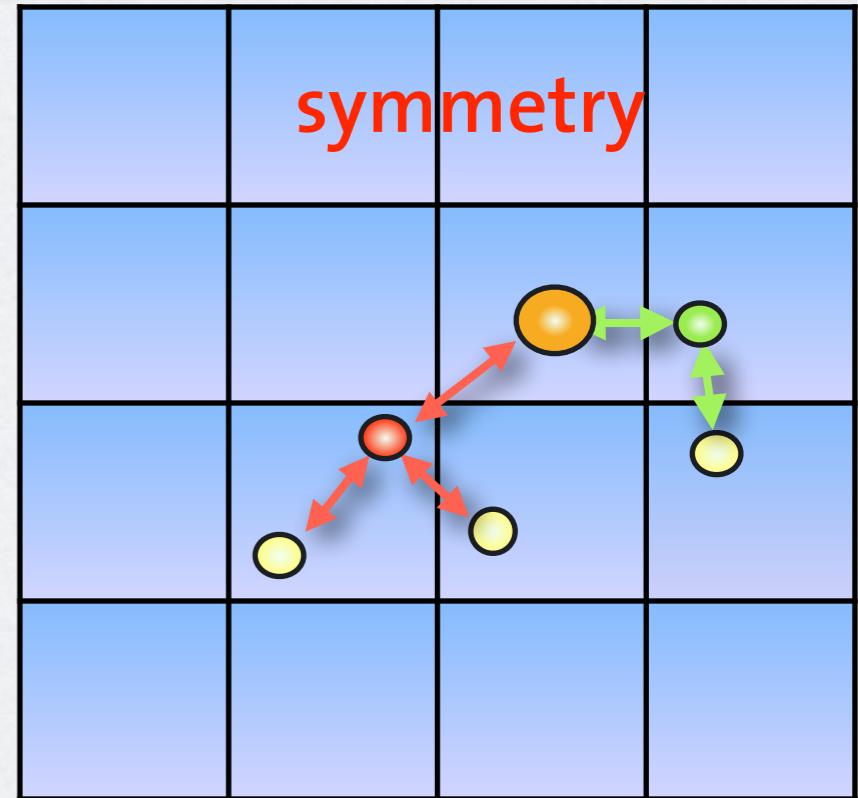
- load imbalance

(some threads do more work than others)

## HOW TO FIX IT:

remove symmetry to avoid race conditions :  
each particle is updated separately

Tasking



# Parallelization over Particles

→ one thread processes one particle

- particle of thread 1
- particle of thread 2
- other particles

## PROBLEMS

- race condition!

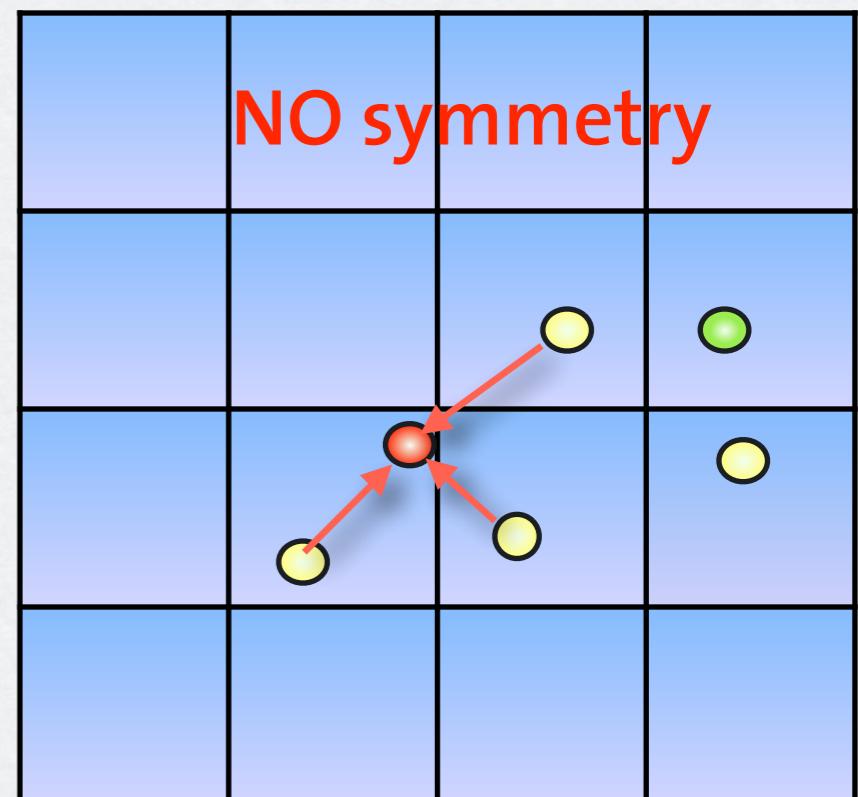
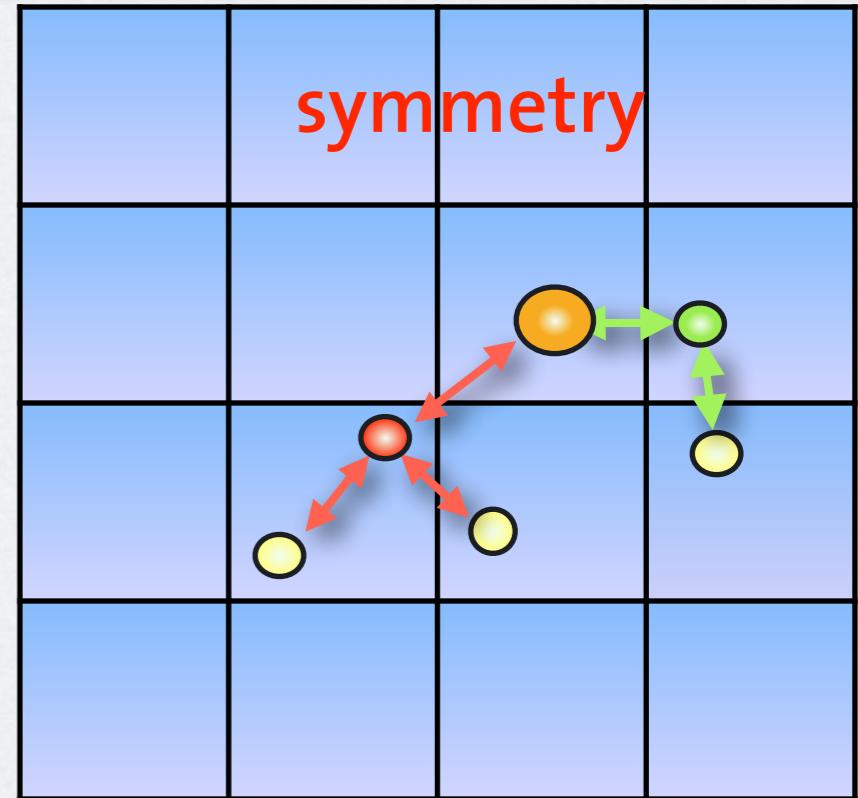
- load imbalance

(some threads do more work than others)

## HOW TO FIX IT:

remove symmetry to avoid race conditions :  
each particle is updated separately

Tasking



# Parallelization over Particles

→ one thread processes one particle

- particle of thread 1
- particle of thread 2
- other particles

## PROBLEMS

- race condition!

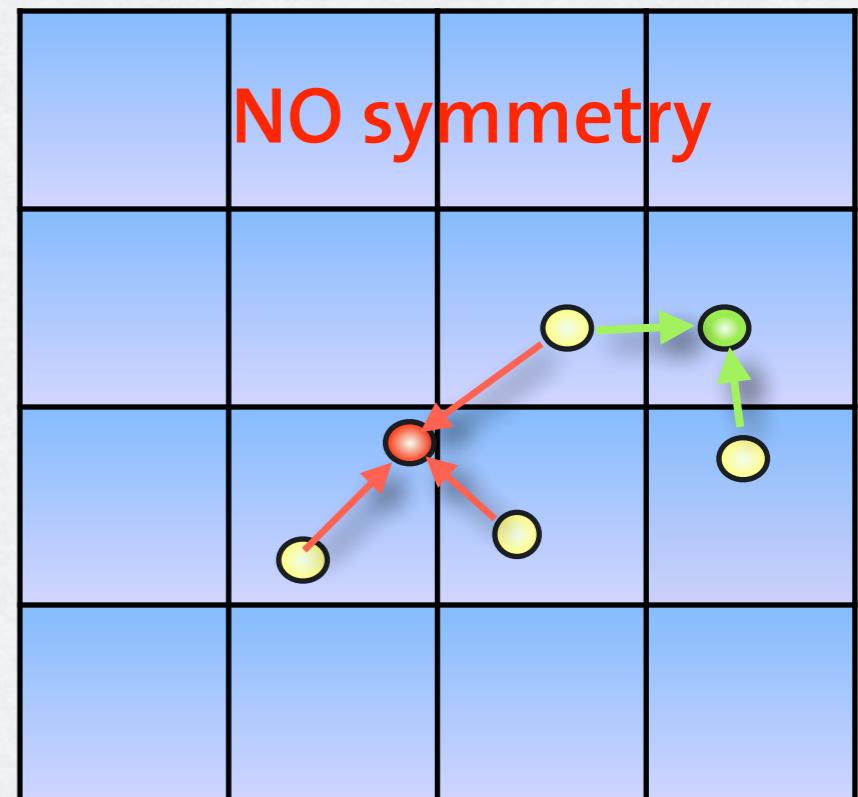
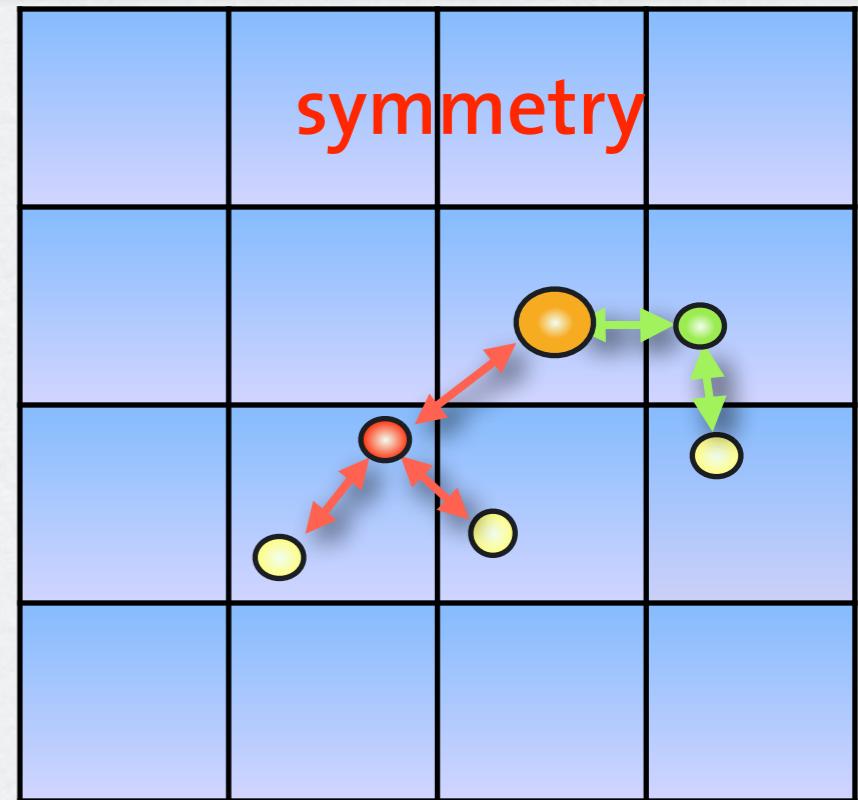
- load imbalance

(some threads do more work than others)

## HOW TO FIX IT:

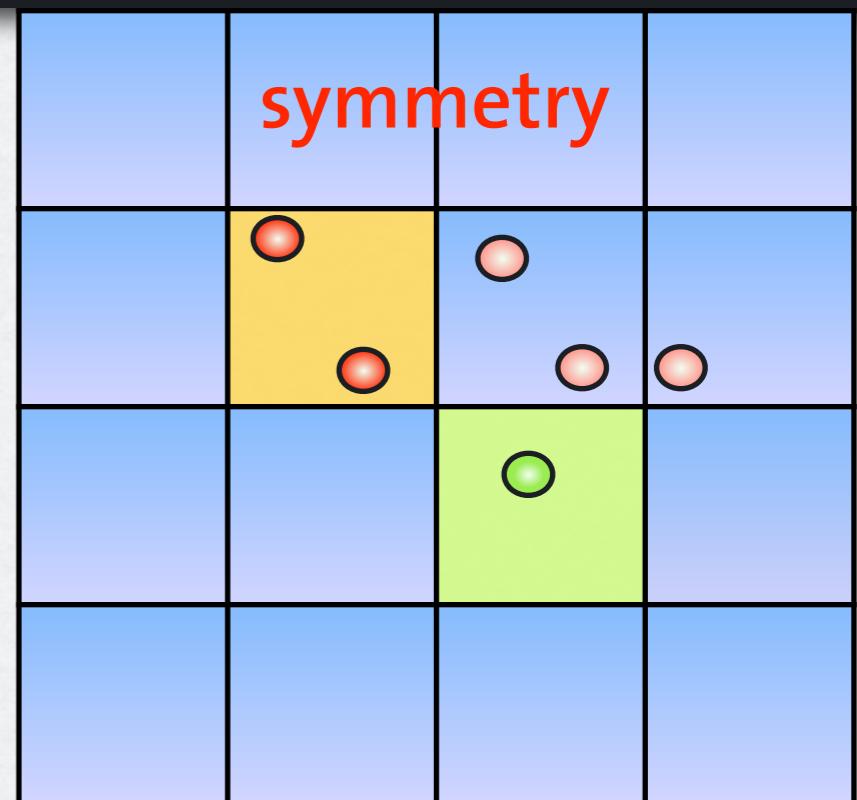
- remove symmetry to avoid race conditions :  
each particle is updated separately

- Tasking



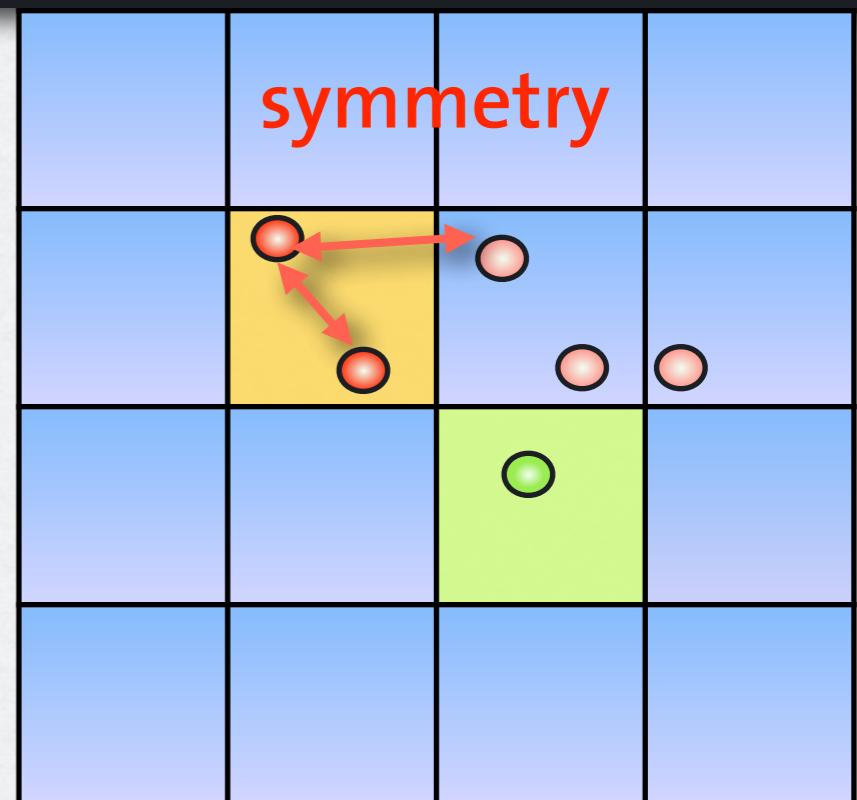
# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells



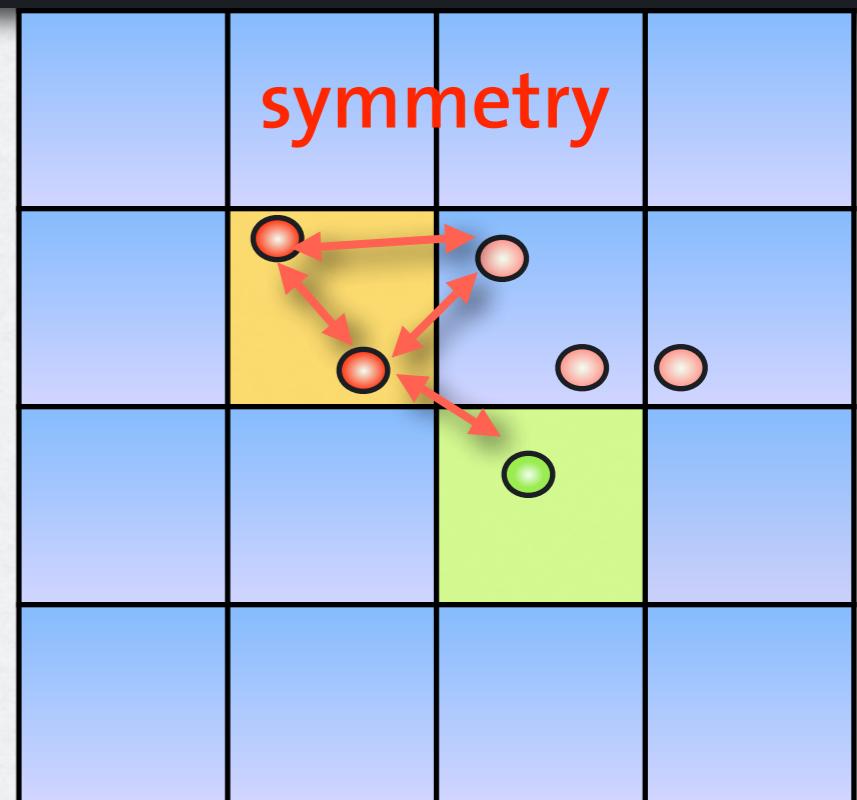
# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells



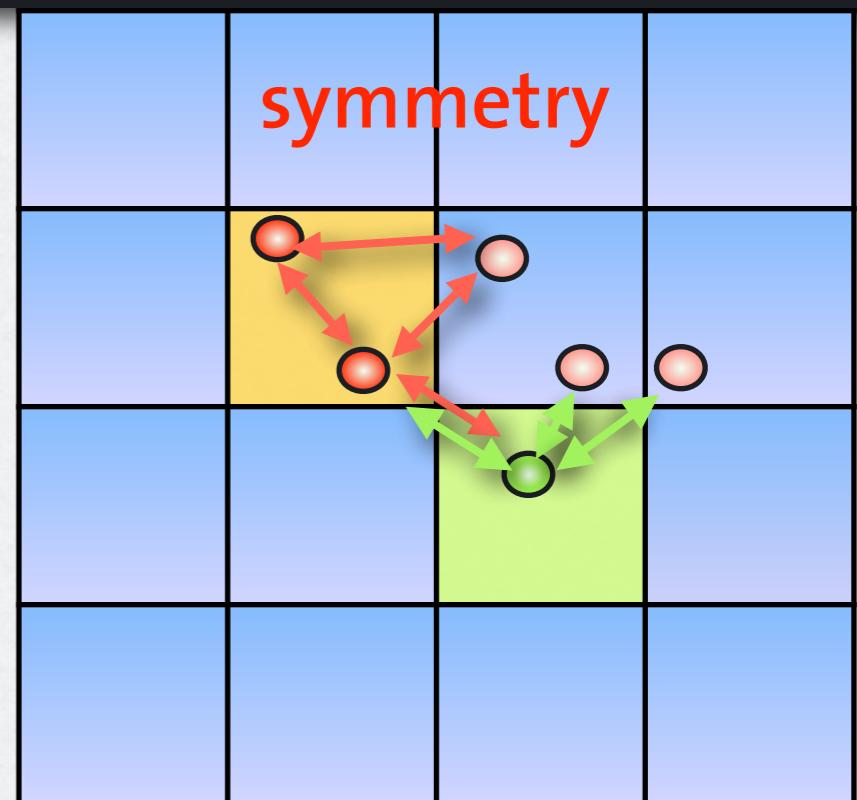
# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells



# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

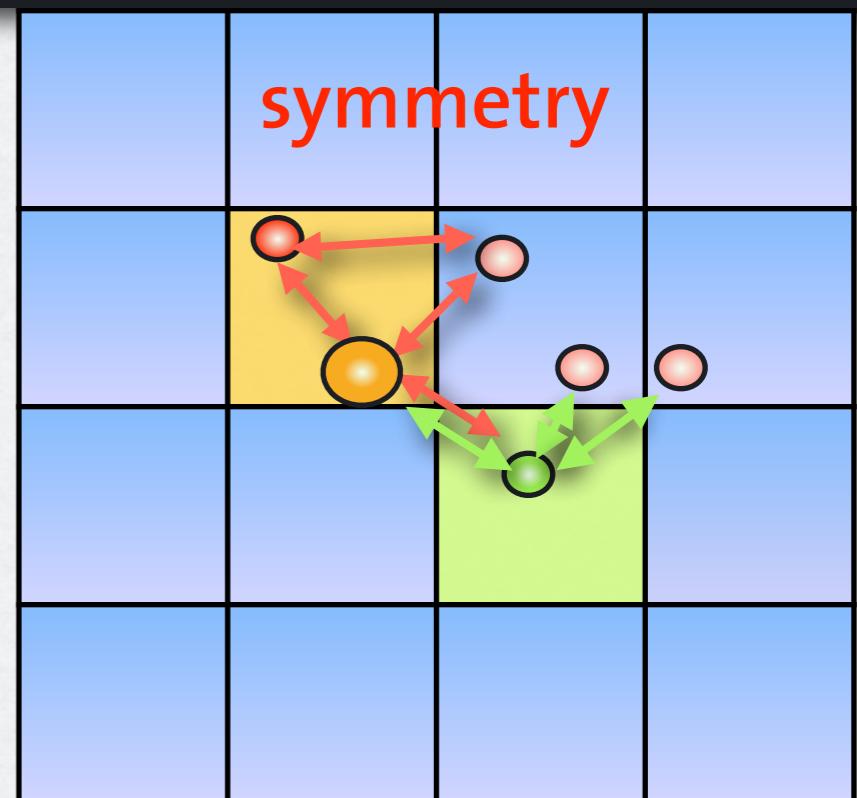


# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

## POTENTIAL PROBLEMS

- race condition!

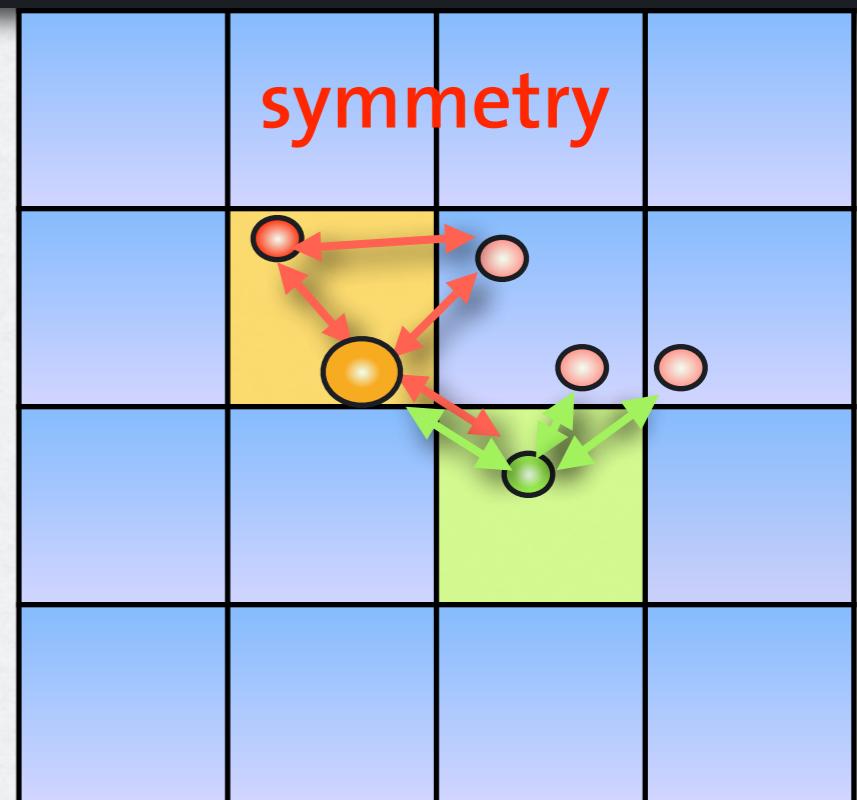


# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

## POTENTIAL PROBLEMS

- race condition!
- load imbalance (e.g. if some cell is empty)



# Parallelization over Cells

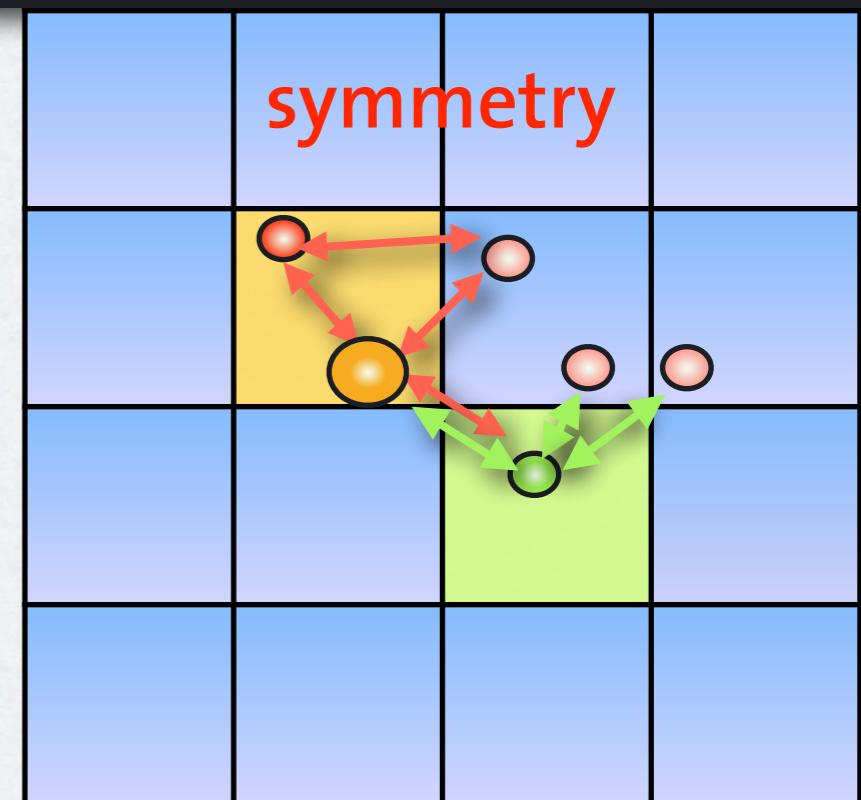
- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

## POTENTIAL PROBLEMS

- race condition!
- load imbalance (e.g. if some cell is empty)

## How to reduce load imbalance:

- skip empty cells
- TASKING : handling each cell is a task



# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

## POTENTIAL PROBLEMS

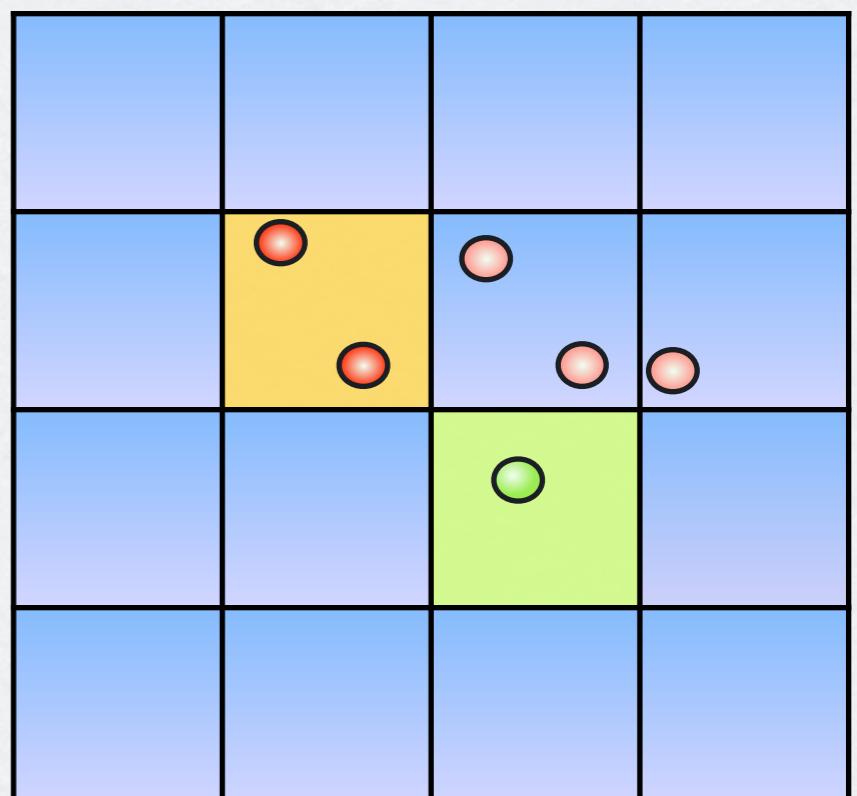
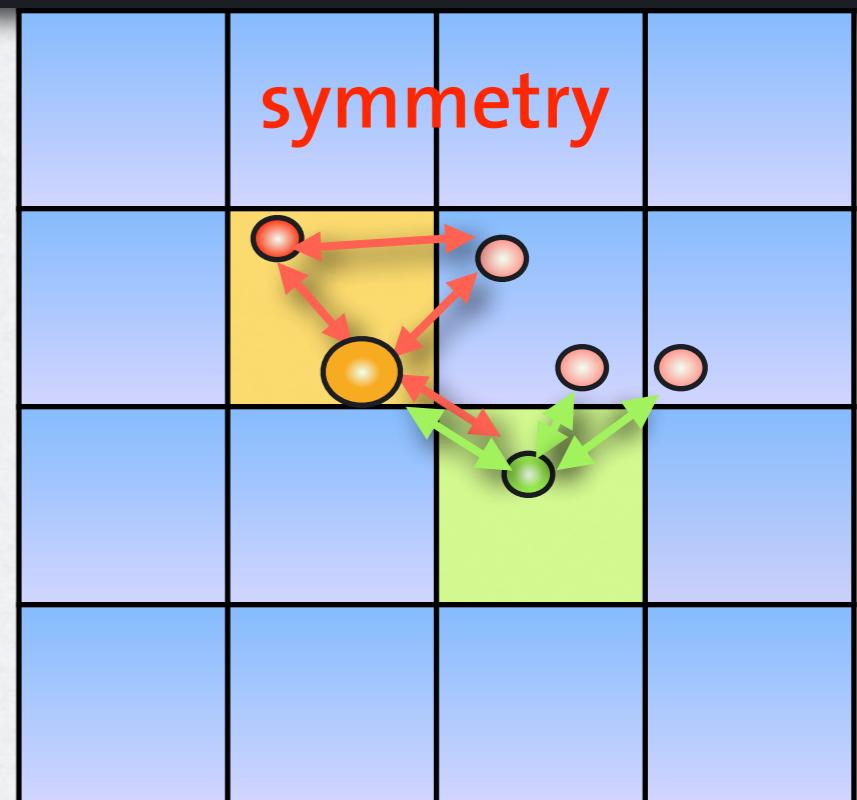
- race condition!
- load imbalance (e.g. if some cell is empty)

## How to reduce load imbalance:

- skip empty cells
- TASKING : handling each cell is a task

## How to fix race condition:

- each particle updates just itself



# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

## POTENTIAL PROBLEMS

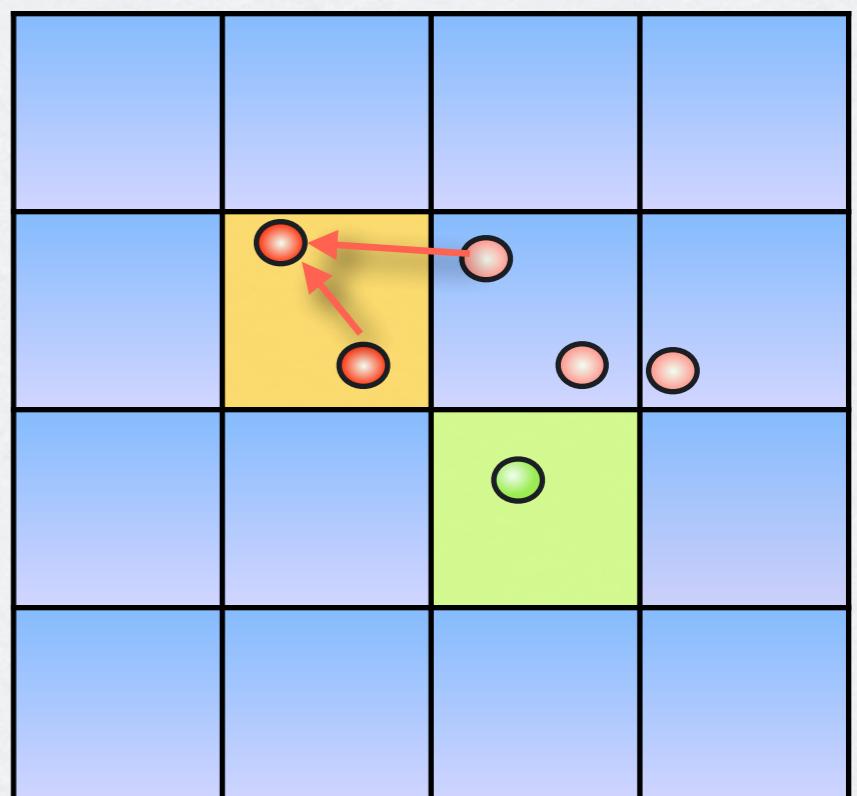
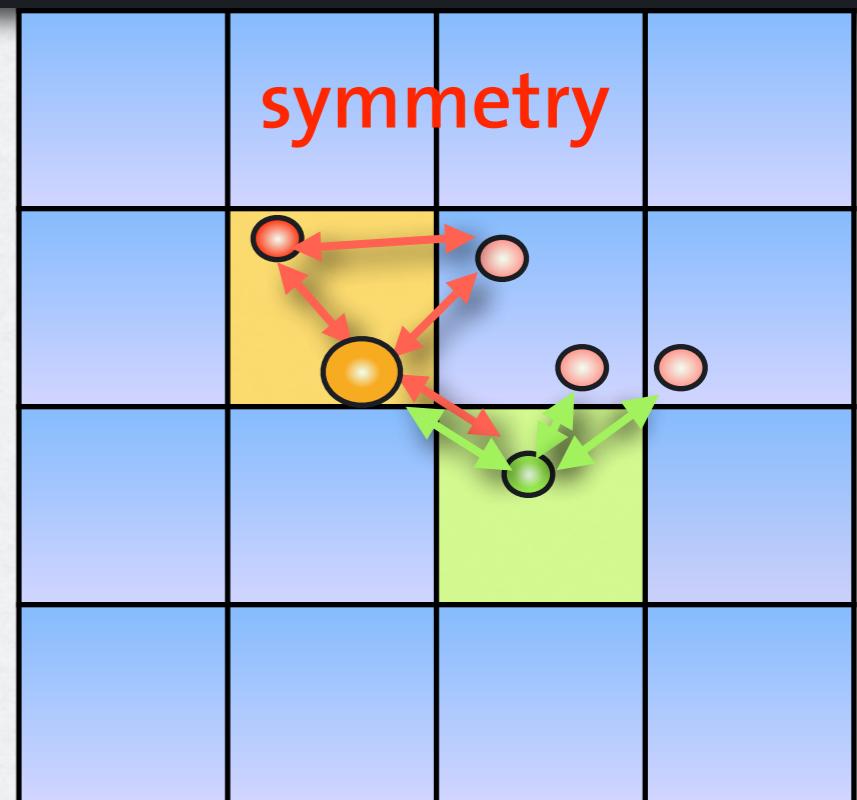
- race condition!
- load imbalance (e.g. if some cell is empty)

## How to reduce load imbalance:

- skip empty cells
- TASKING : handling each cell is a task

## How to fix race condition:

- each particle updates just itself



# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

## POTENTIAL PROBLEMS

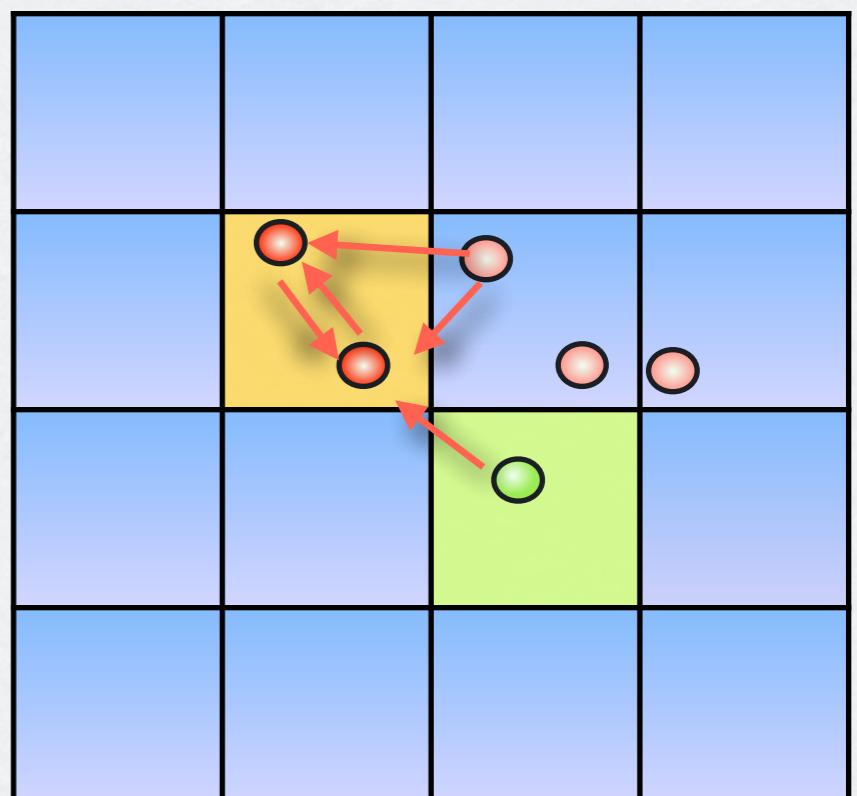
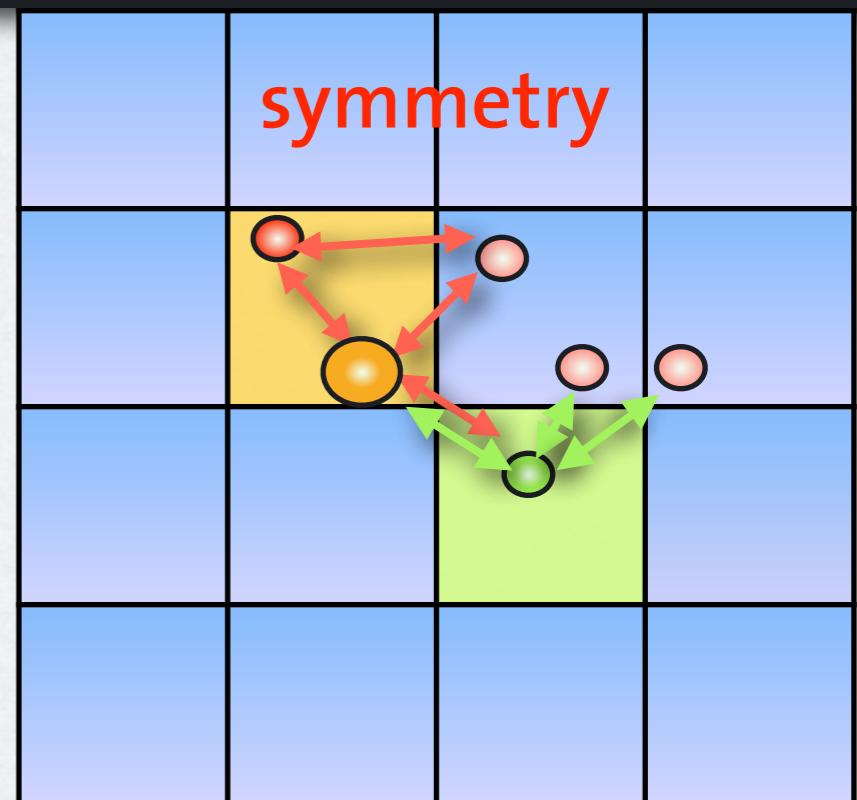
- race condition!
- load imbalance (e.g. if some cell is empty)

## How to reduce load imbalance:

- skip empty cells
- TASKING : handling each cell is a task

## How to fix race condition:

- each particle updates just itself



# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

## POTENTIAL PROBLEMS

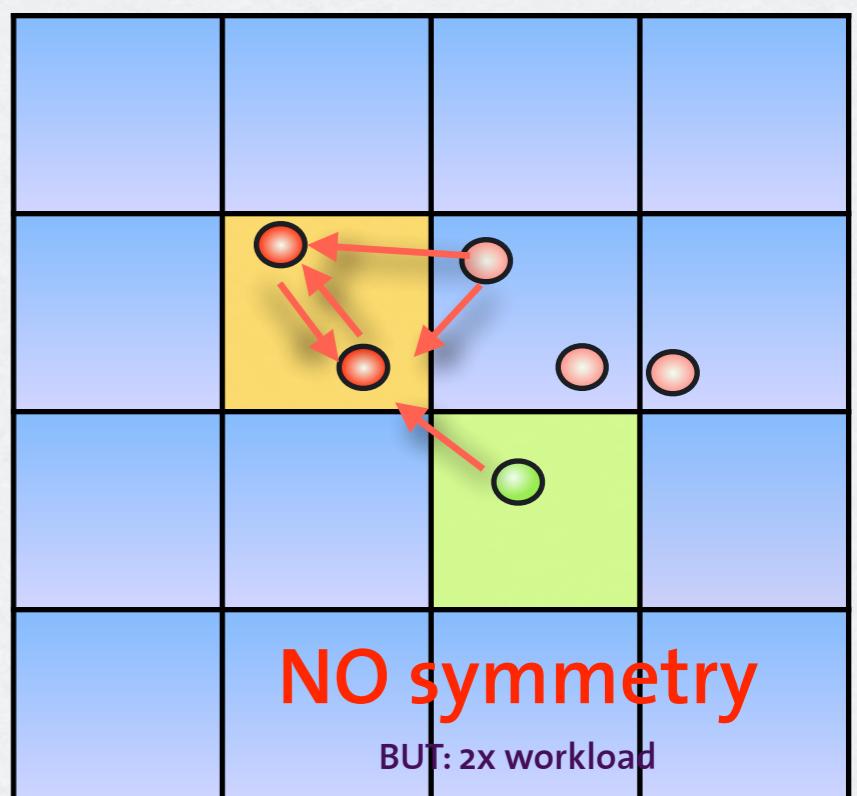
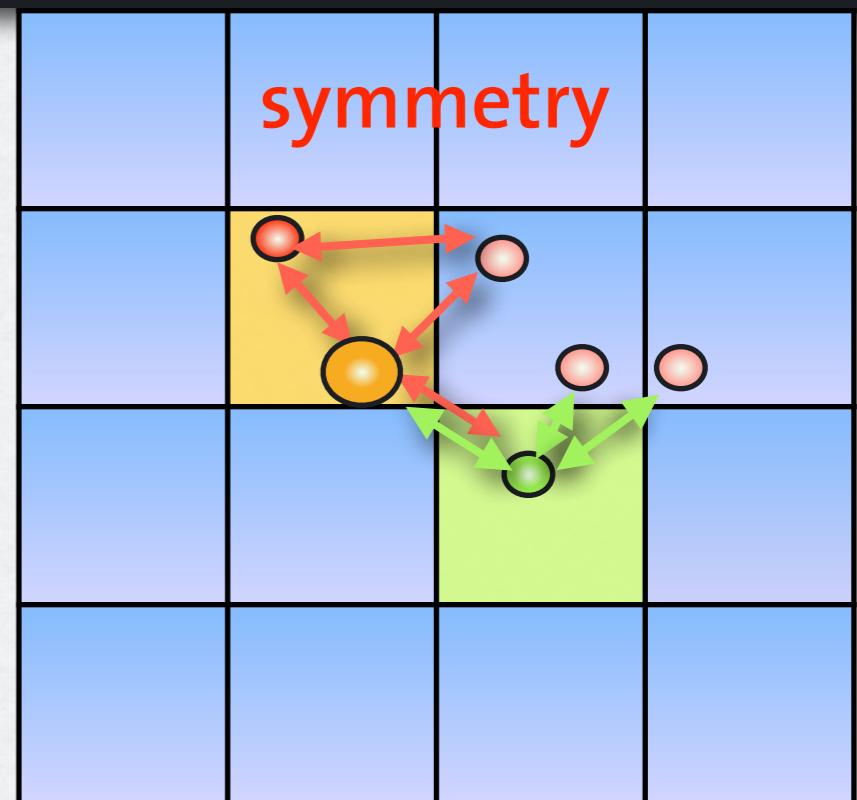
- race condition!
- load imbalance (e.g. if some cell is empty)

## How to reduce load imbalance:

- skip empty cells
- TASKING : handling each cell is a task

## How to fix race condition:

- each particle updates just itself



# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

## POTENTIAL PROBLEMS

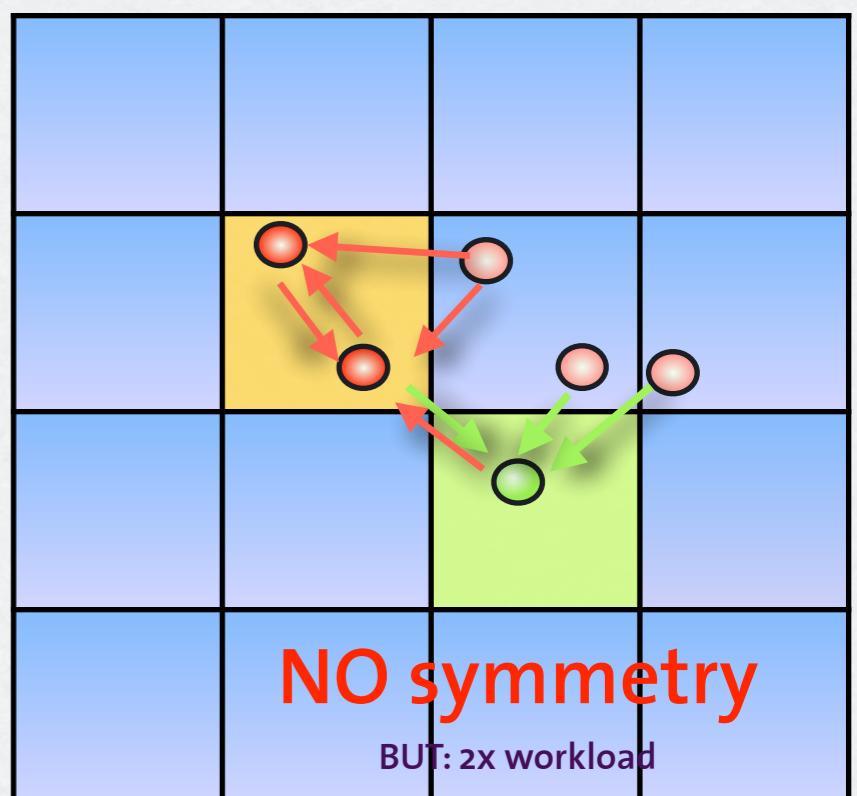
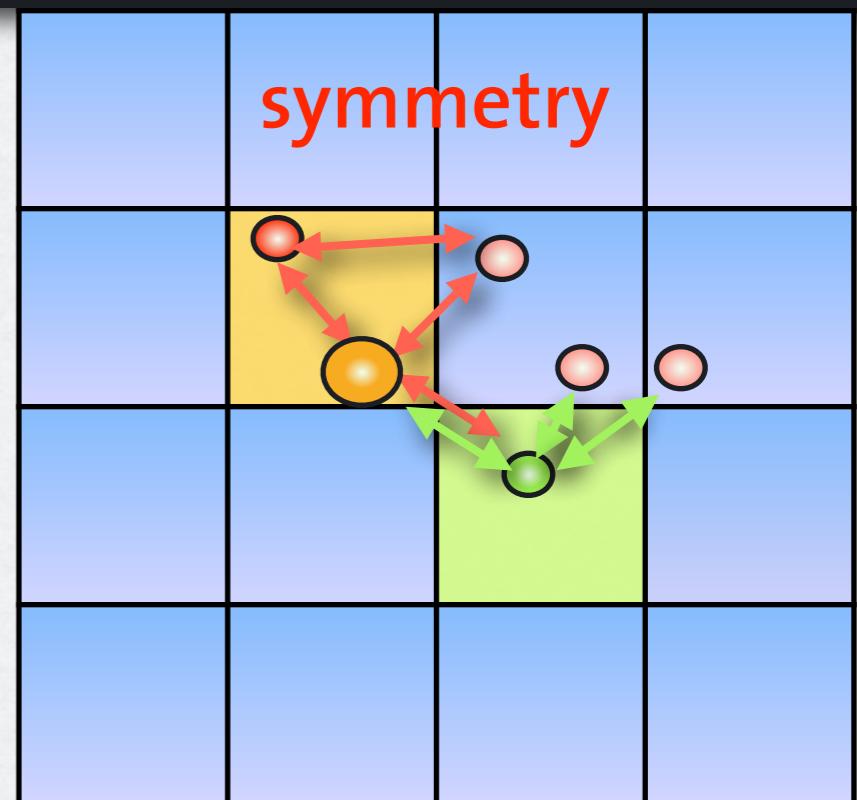
- race condition!
- load imbalance (e.g. if some cell is empty)

## How to reduce load imbalance:

- skip empty cells
- TASKING : handling each cell is a task

## How to fix race condition:

- each particle updates just itself

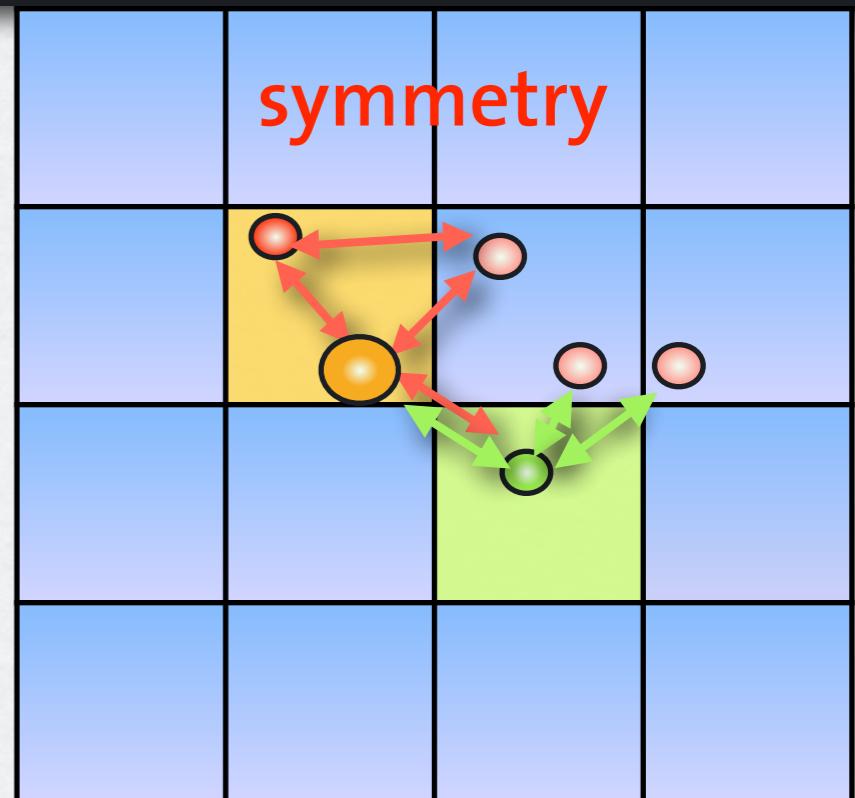


# Parallelization over Cells

- one thread processes one cell
  - particles in cell of thread 1
  - particles in cell of thread 2
  - particles in other cells

## POTENTIAL PROBLEMS

- race condition!
- load imbalance (e.g. if some cell is empty)

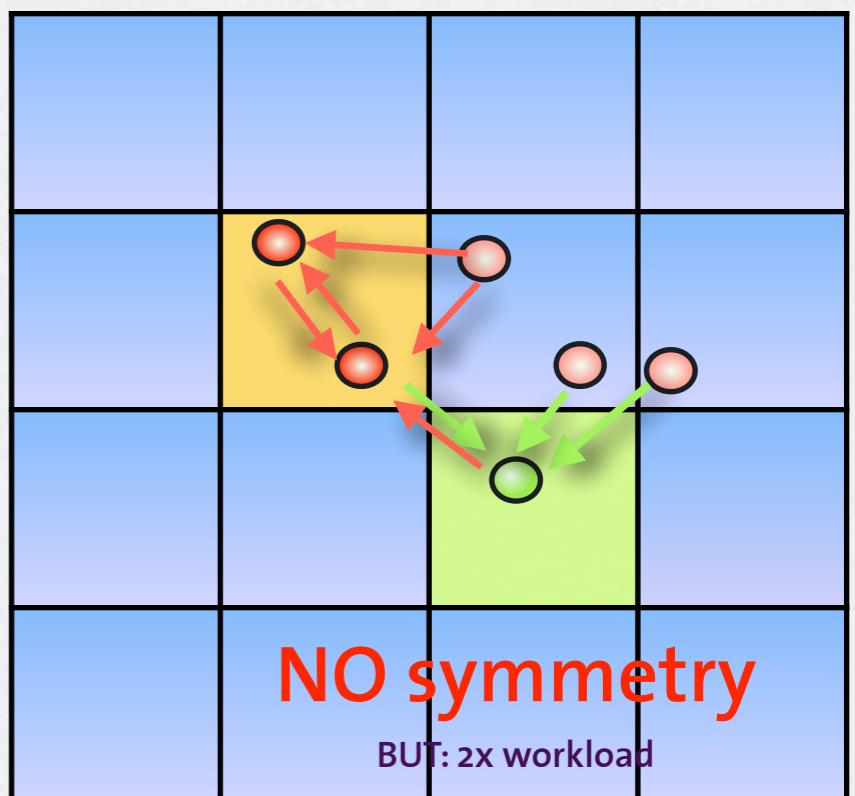


## How to reduce load imbalance:

- skip empty cells
- TASKING : handling each cell is a task

## How to fix race condition:

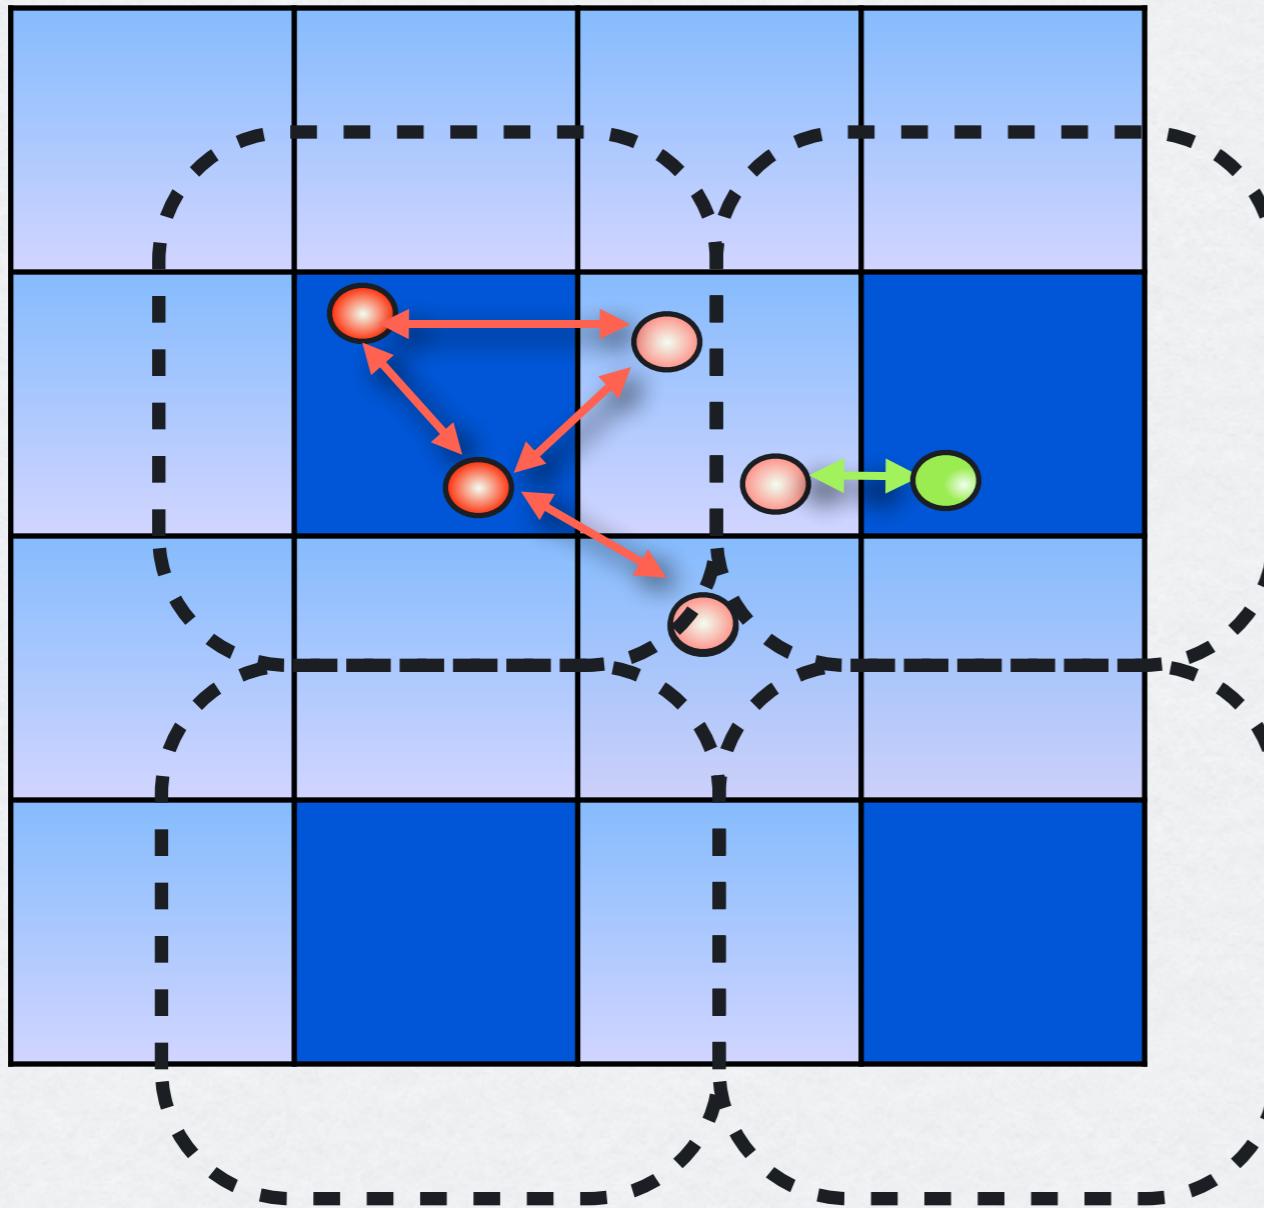
- each particle updates just itself
- multipass (next slide)



# Divide and Conquer - Multipass

Define Cell size  $\geq$  interaction *diameter* of particles

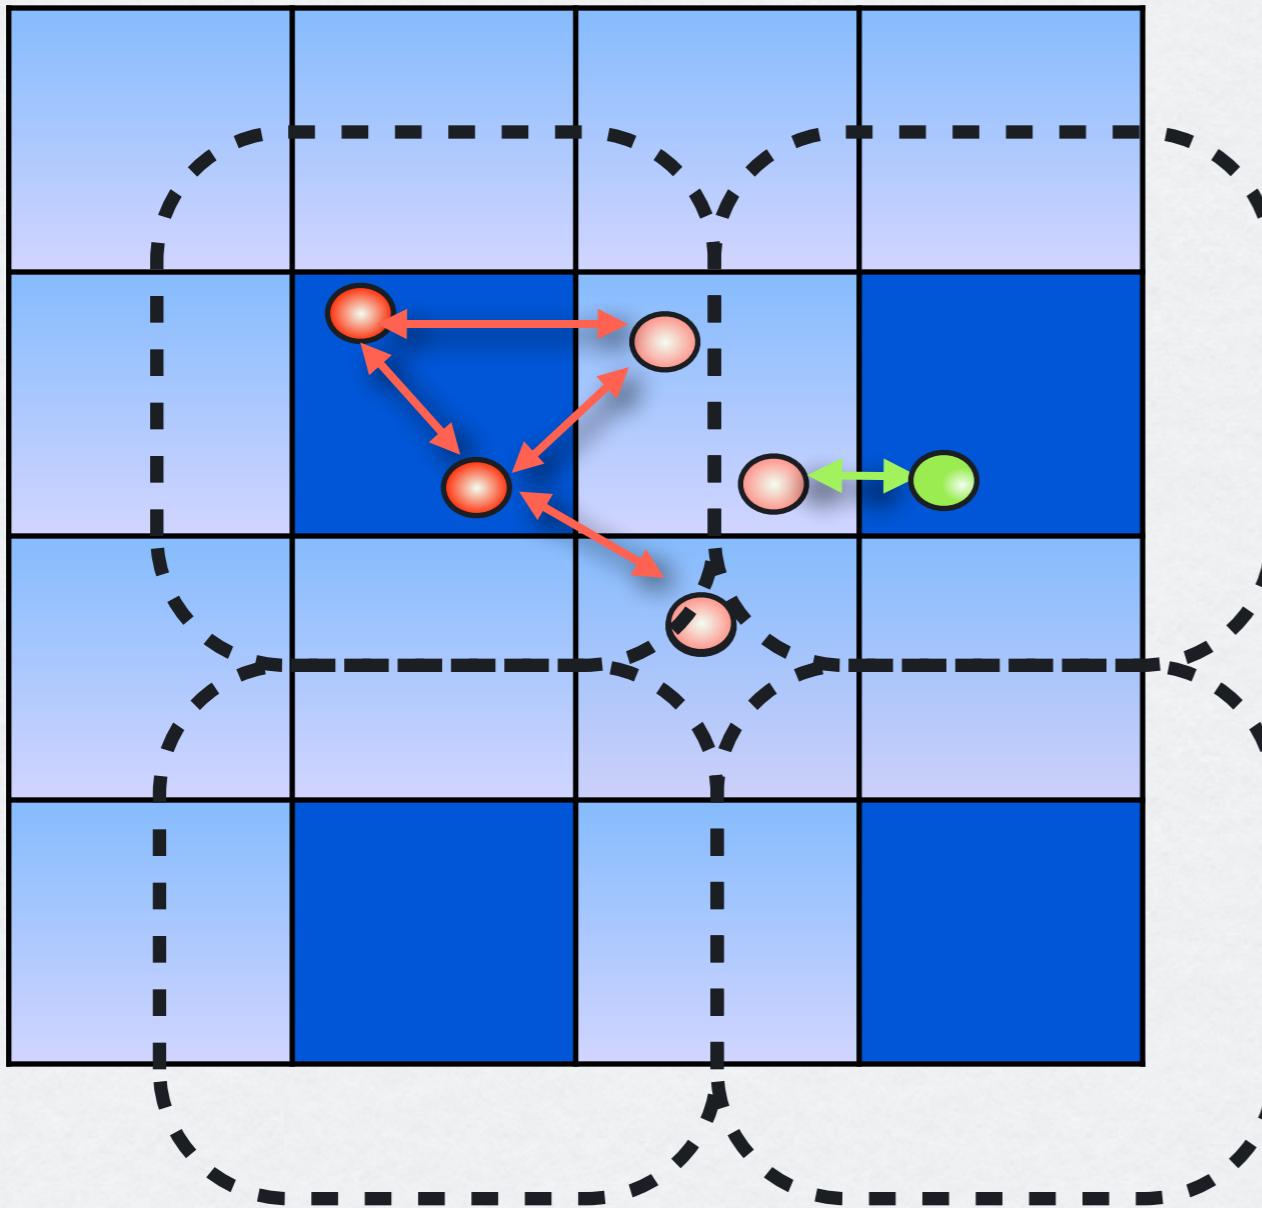
“Pass” : Process non-neighboring cells simultaneously



# Divide and Conquer - Multipass

Define Cell size  $\geq$  interaction *diameter* of particles

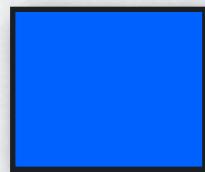
“Pass” : Process non-neighboring cells simultaneously



NO race condition

# Divide and Conquer - Multipass

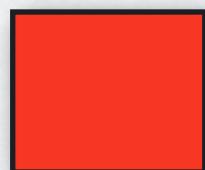
→ in 2D there must be 4 passes to cover all possible cells



cells of pass 1



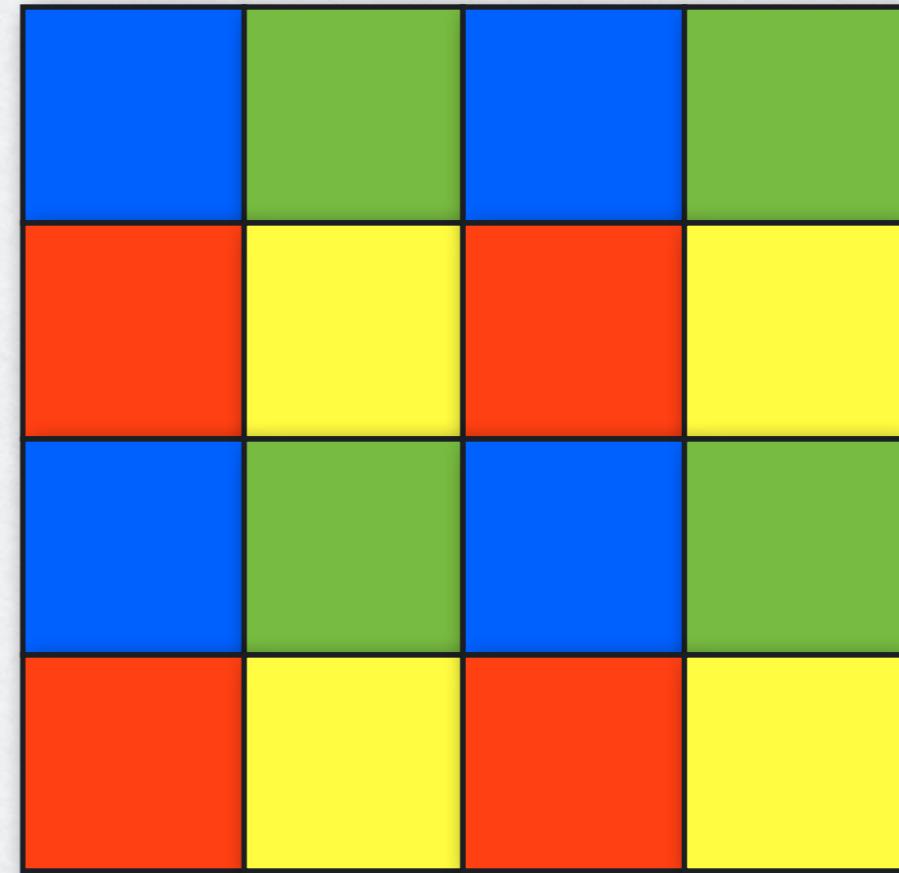
cells of pass 2



cells of pass 3



cells of pass 4



cells of the same color can be processed at the same time

→ in 3D: 8 passes are needed

# Issues with Cell Lists

- Curse of dimensionality:
  - 2D - 35% of the particles checked are within the cutoff
  - 3D - 16% of the particles checked are within the cutoff
- Building cell lists every timestep might be expensive