

# Project 1

## N-Body on GPU

High Performance Computing for Science and Engineering

October 1, 2013

**CSElab**

Computational Science & Engineering Laboratory

<http://www.cse-lab.ethz.ch>

**ETH**

Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

# Administrative Notes

- Please contact one of the TAs if you need anything
  - Please avoid writing to Anna (our secretary)
  - Please avoid writing to Prof. Koumoutsakos
- Hand-in to your TA
  - If you have a special request, please ask by the end of the day
  - Please check again the list tomorrow evening for changes
- Exam
  - Tuesday, 17.12.2013, 14:00 - 17:00
  - Computer rooms

# The Brutus Cluster

High Performance Computing for Science and Engineering

October 1, 2013

**CSElab**

Computational Science & Engineering Laboratory

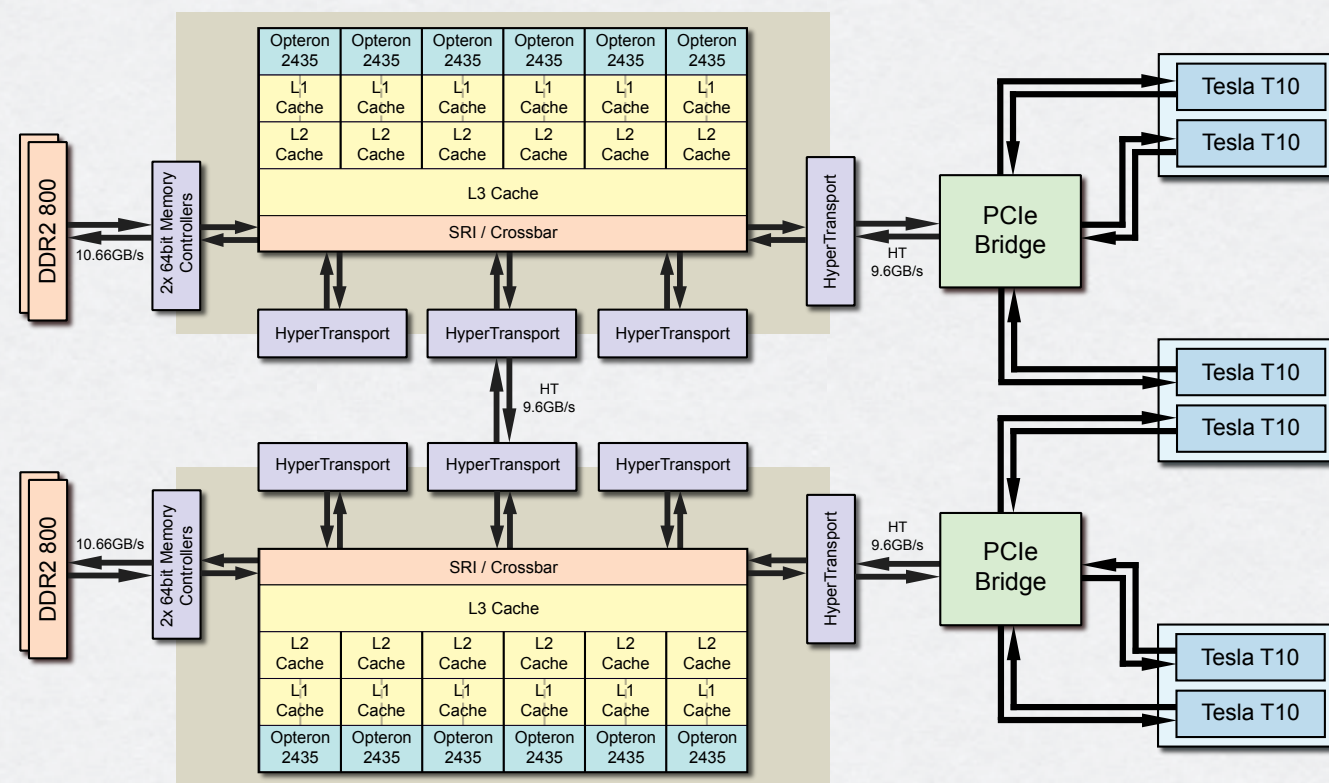

<http://www.cse-lab.ethz.ch>

**ETH**

Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

# Anatomy of a Cluster

## Brutus

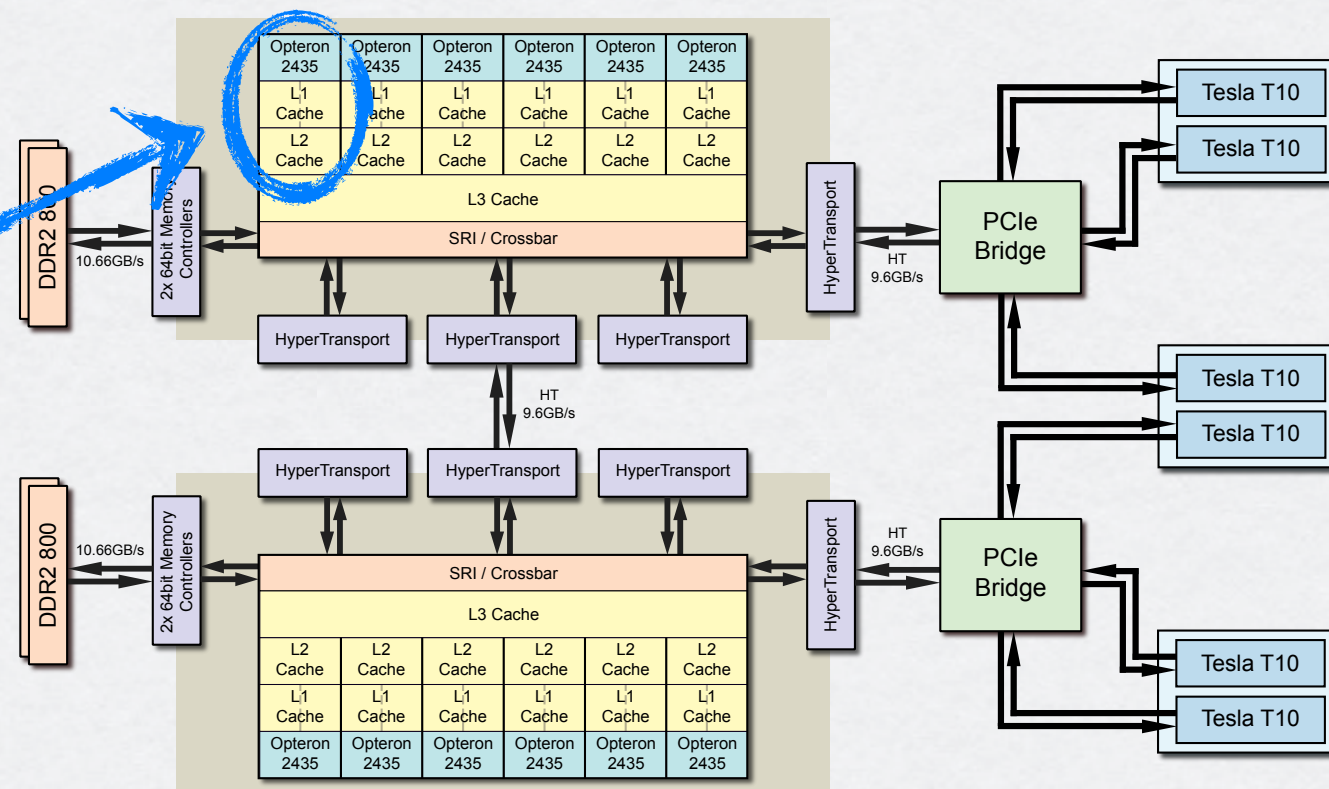



### Details

- Effective bandwidth with 12 cores: 20GB/s (STREAM Benchmark)

# Anatomy of a Cluster

# Brutus



## Details

- Effective bandwidth with 12 cores: 20GB/s (STREAM Benchmark)


CPU core

**C++**

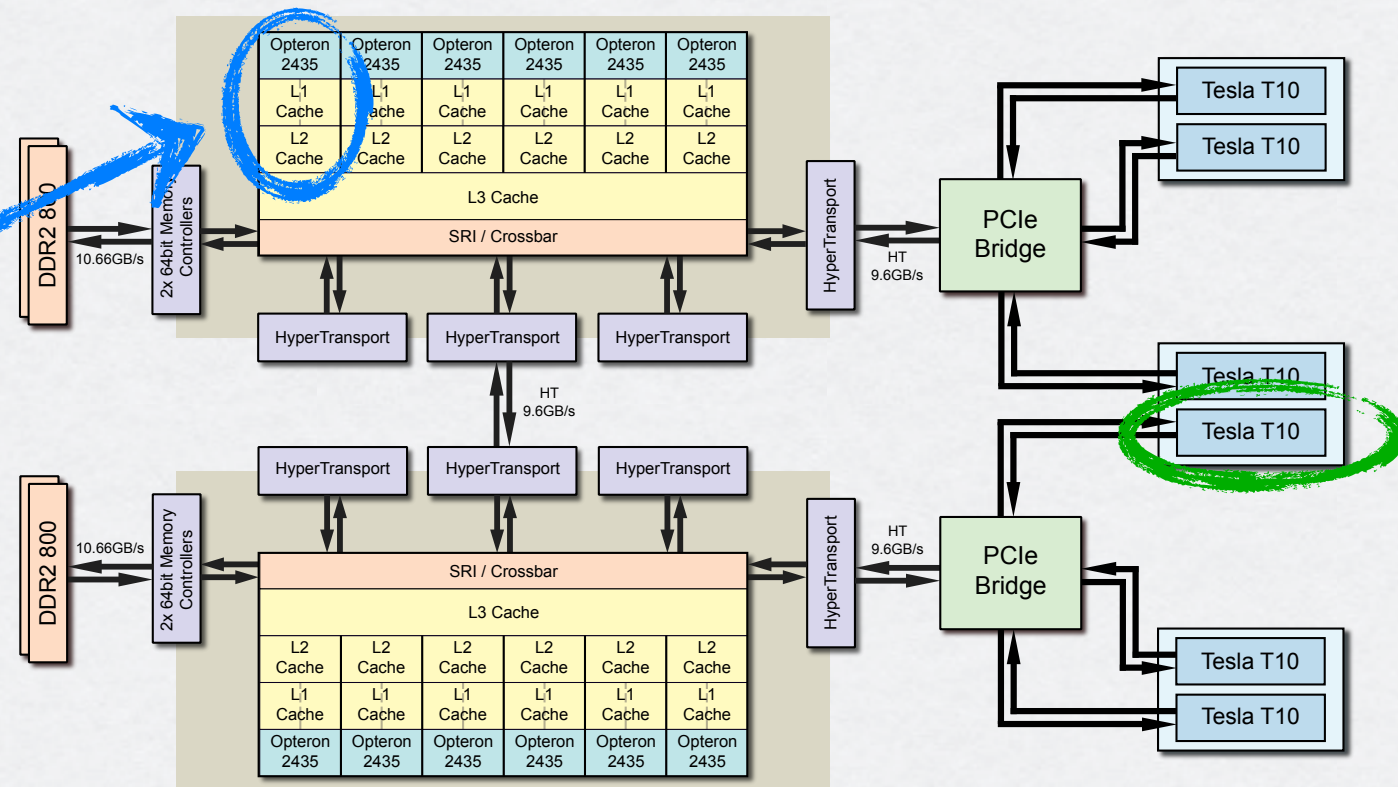


# Anatomy of a Cluster

## Brutus



CPU core  
C++




GPUs  
CUDA

Details  
- Effective bandwidth with 12 cores: 20GB/s (STREAM Benchmark)

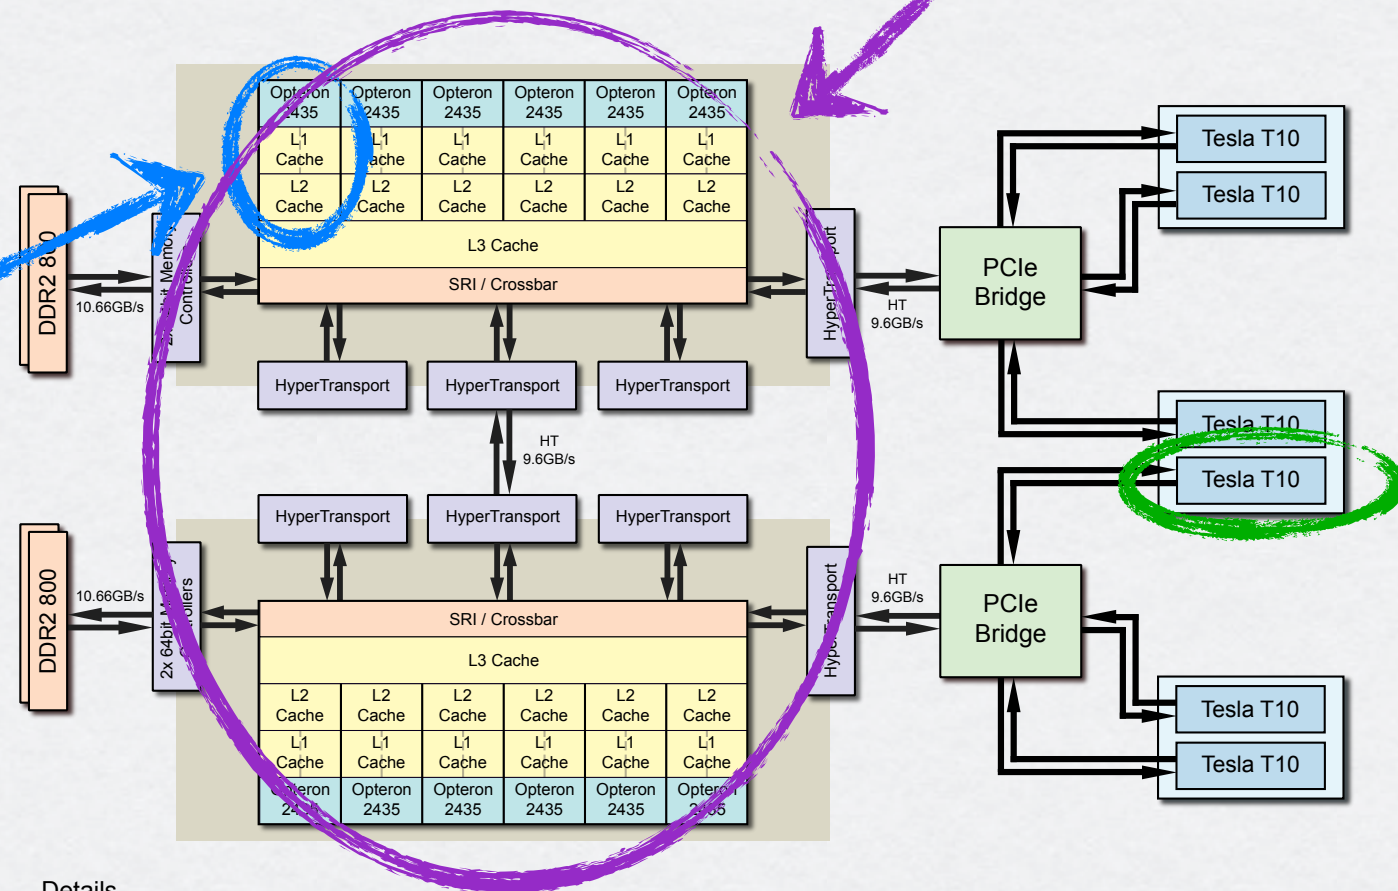
# Anatomy of a Cluster

## Brutus



Node: multiple processors  
Shared Memory  
OpenMP

CPU core  
C++




GPUs  
CUDA

Details  
- Effective bandwidth with 12 cores: 20GB/s (STREAM Benchmark)



# Anatomy of a Cluster

## Brutus



Cluster: network of nodes

Distributed memory

**MPI**

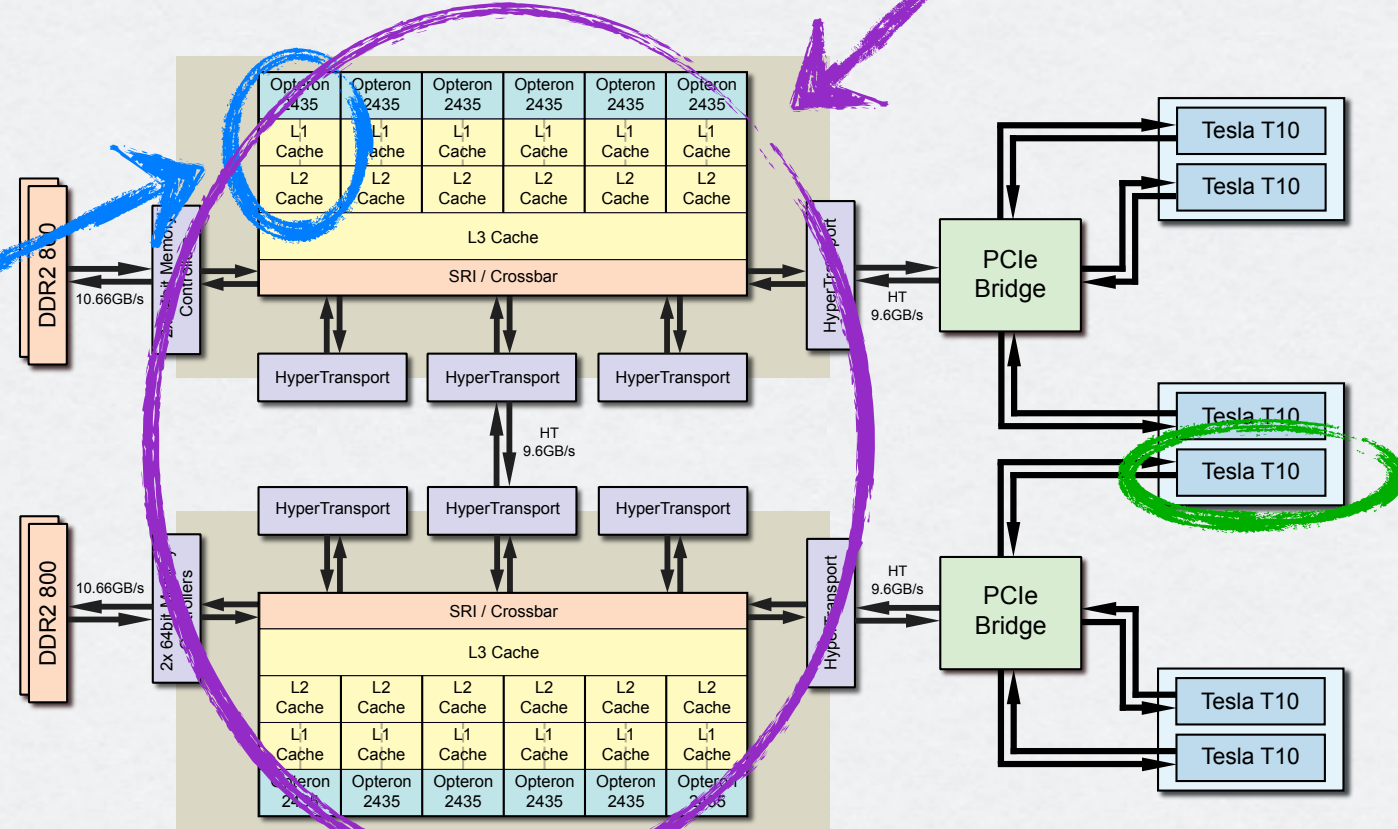
Node: multiple processors

Shared Memory

**OpenMP**

CPU core

**C++**



GPUs  
**CUDA**

Details

- Effective bandwidth with 12 cores: 20GB/s (STREAM Benchmark)



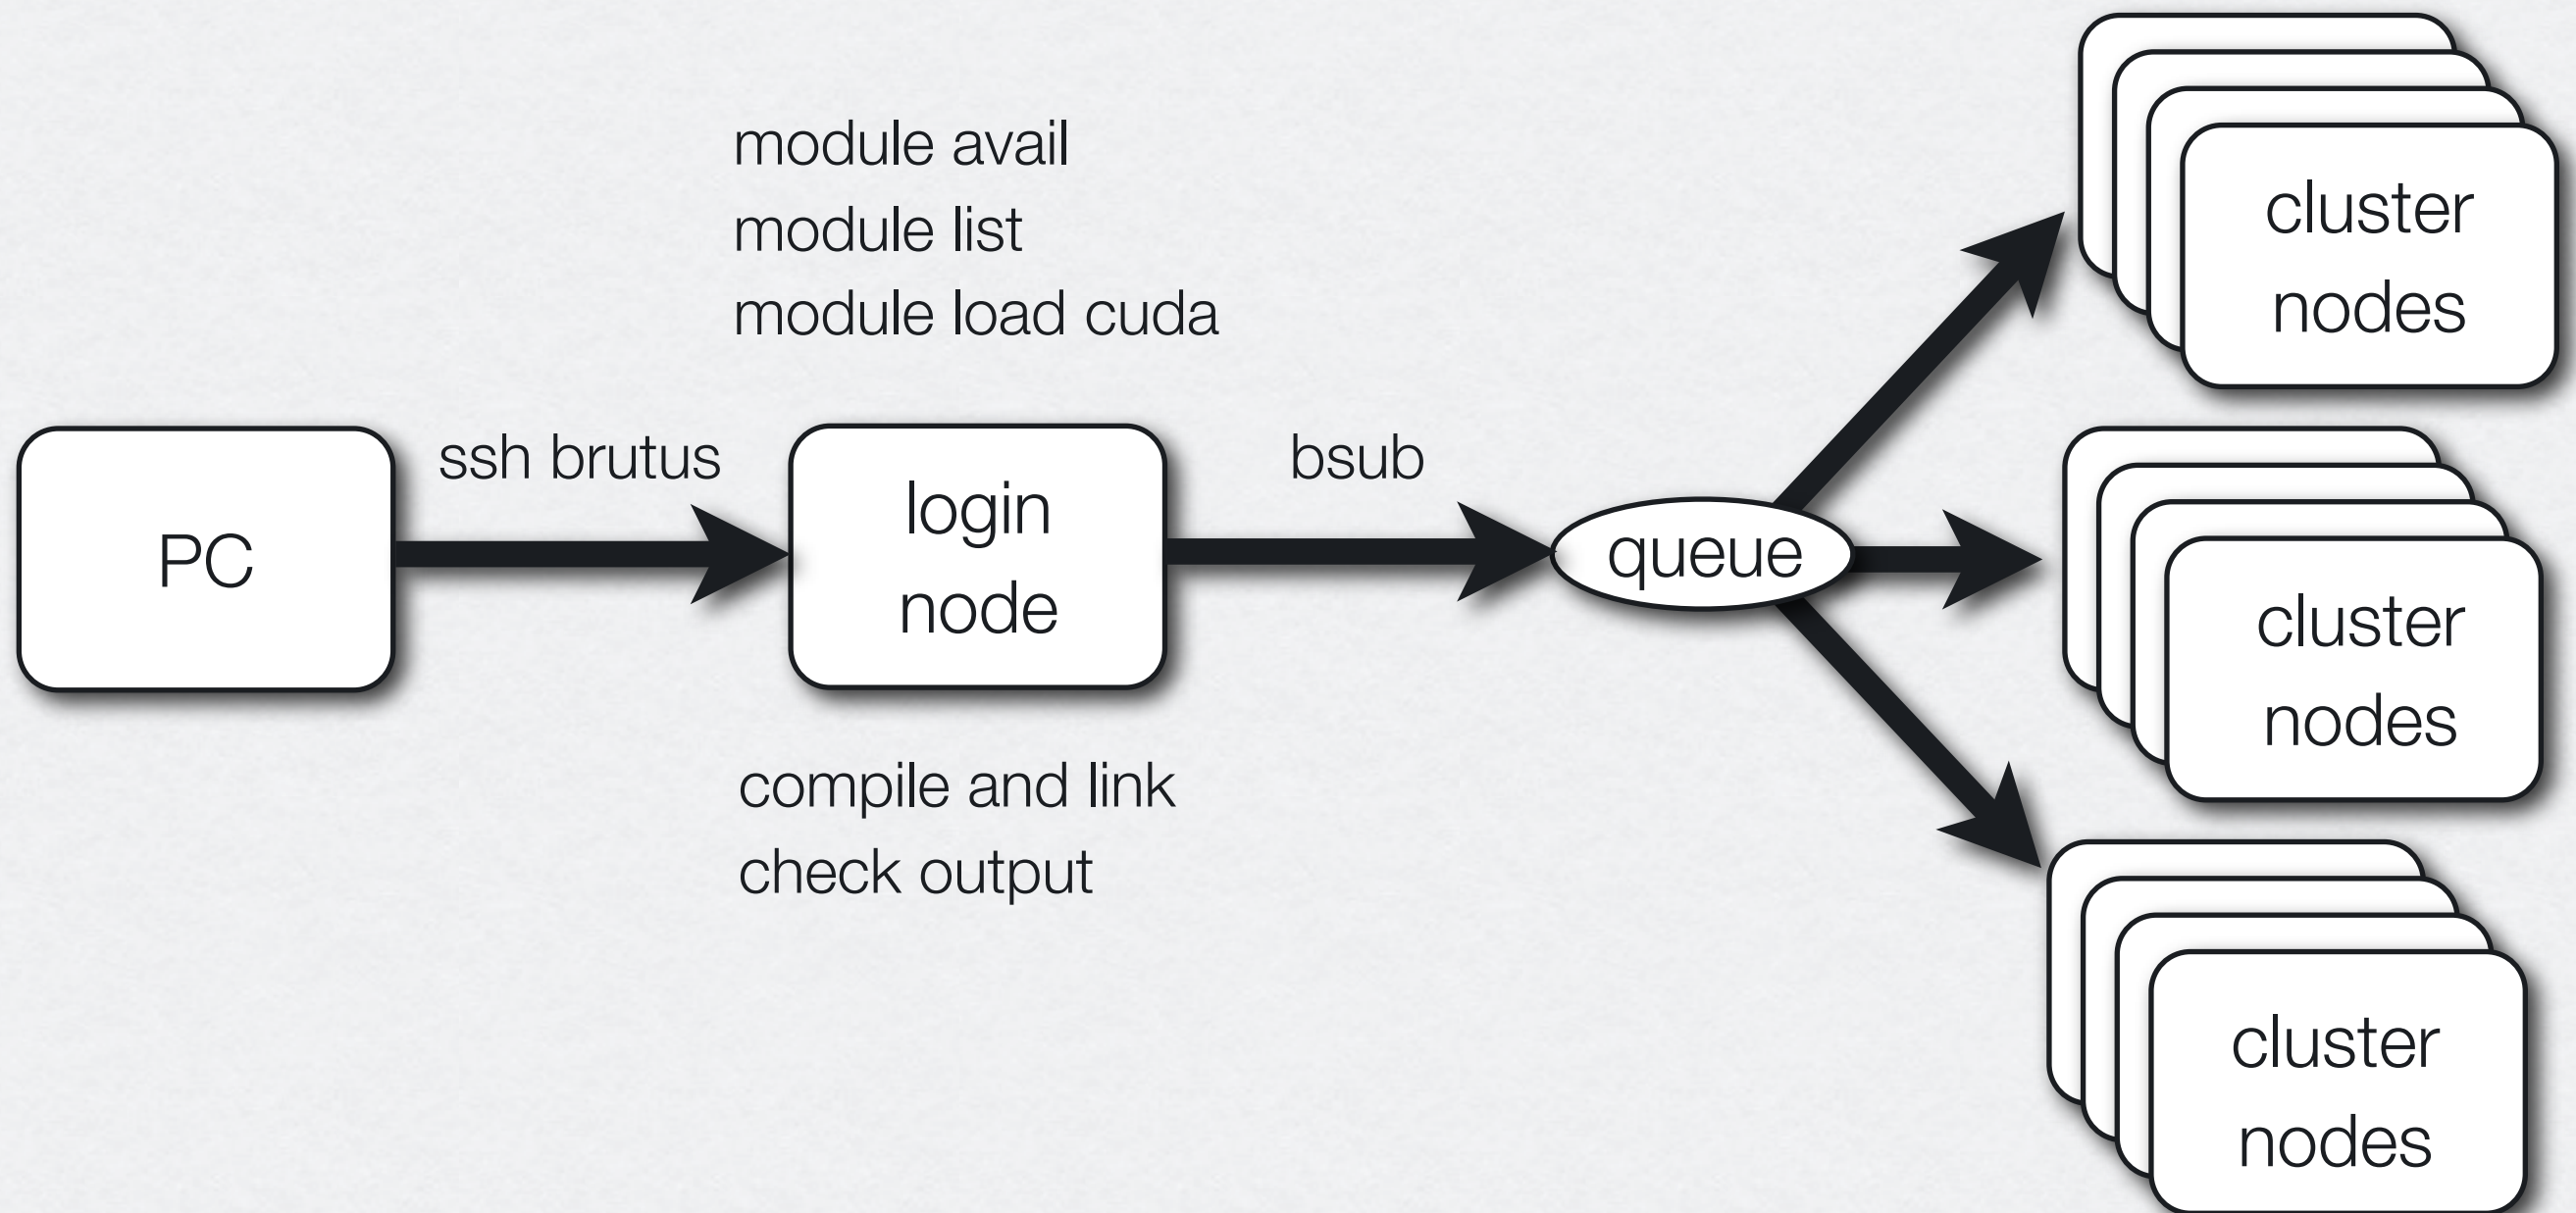
# Brutus

- High performance cluster of ETH Zurich
- Composed of different kinds of compute nodes
  - 36x Nvidia Tesla C2050 (Fermi Architecture)
  - 120 nodes with 48 cores each
  - Many others (check [brutuswiki.ethz.ch](http://brutuswiki.ethz.ch))

# Accessing and Using Brutus

Request account (instructions on [brutuswiki.ethz.ch](http://brutuswiki.ethz.ch))

- put “HPCSE I” as “project”



# Account Request

The first thing you must do is to request for an account

- As STUDENTS attending a parallel computing class
- Brutus wiki / Contact us / Brutus account request
- [https://www1.ethz.ch/id/services/list/comp\\_zentral/cluster/brutus\\_acc\\_req\\_pre\\_EN](https://www1.ethz.ch/id/services/list/comp_zentral/cluster/brutus_acc_req_pre_EN)



# Basic steps

1. Connect to a login node of Brutus
2. Copy or edit your program files
3. Compile your program
4. Submit a job / run your program on compute nodes
5. Check your job (status and output)

# 1. Connect

- `ssh -Y <username>@brutus.ethz.ch`
  - `-Y`: Enables trusted X11 forwarding
  - Access to one of the Brutus login nodes

# 2. Develop

- Copy your files to Brutus, e.g.
  - `scp code.tar.gz <username>@brutus.ethz.ch:code.tar.gz`
- Use a text editor to write/modify your code
  - vi, emacs, nano, nedit



# 3. Compile

- You will need the appropriate programming tools and libraries to compile your code
  - By default, only the GNU compiler (gcc-4.4.7) is available
- Just load the environment module you need
- Examples
  - `module load gcc/4.6` (newer version of gcc)
  - `module load cuda` (for GPU codes)
  - `module list` (shows loaded modules)
  - `module avail` (what is available)
  - `module unload cuda` (unloads a module)

# 3. Compile

- Compile your code and produce the executable
- Examples:
  - `g++ cputest.cpp -o CPUexec`
  - `nvcc cudatest.cu -o CUDAexec`

# 4. Submit your job

- The login nodes are used only for development
- The program must run on a compute node
- To do that, you must use the bsub command
- CPU code: `bsub -W 00:10 -n 1 ./CPUexec`
- GPU job: `bsub -W 00:10 -n 1 -R gpu ./CUDAexec`
- You can submit script files too: `bsub -n 1 < myscript`



# 5. Check your job

- Some useful commands
  - bqueues: displays information about queues
  - bjobs: displays information about jobs (bjobs -l -u <username>)
  - bkill <jobID>: kills a job
- Output files
  - lsf.o<jobID>: created in your working directory when the job finishes
    - includes information about your job (statistics, etc.) and the messages the program prints (standard output)

# The N-Body Problem

High Performance Computing for Science and Engineering

October 1, 2013

**CSElab**

Computational Science & Engineering Laboratory



<http://www.cse-lab.ethz.ch>

**ETH**

Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

# N-Body - What is it About

- N-body problem is a classical problem of predicting the individual motions of a group of objects interacting with each other under some physical law
- This can be celestial bodies under gravitational forces, molecules under Lennard-Jones potential and so on





# General formulation

- $n$  bodies are represented as  $n$  points with masses  $m_1, m_2, \dots, m_n$
- Positions:  $\mathbf{q}_1(t), \mathbf{q}_2(t), \dots, \mathbf{q}_n(t)$
- Initial conditions for positions  $\mathbf{q}_1(0), \mathbf{q}_2(0), \dots, \mathbf{q}_n(0)$  and velocities  $\dot{\mathbf{q}}_1(0), \dot{\mathbf{q}}_2(0), \dots, \dot{\mathbf{q}}_n(0)$  are given
- Motion or evolution of  $\mathbf{q}_j$  is given by a 2<sup>nd</sup>-order ODE:


$$m_j \ddot{\mathbf{q}}_j = \sum_{\substack{k=1 \\ k \neq j}}^n \mathbf{F}(\mathbf{q}_k, \mathbf{q}_j)$$

- By introducing the variables  $\mathbf{x}_j = \mathbf{q}_j$  and  $\mathbf{v}_j = \dot{\mathbf{q}}_j$  we can convert the 2<sup>nd</sup>-order ODE to a system of first order ODEs:

$$\begin{pmatrix} \dot{\mathbf{x}}_j \\ m_j \dot{\mathbf{v}}_j \end{pmatrix} = \begin{pmatrix} \mathbf{v}_j \\ \sum_{k \neq j} \mathbf{F}(\mathbf{x}_k, \mathbf{x}_j) \end{pmatrix}$$

# Discretization

- The equations have  $6n$  degrees of freedom in 3D:

$$\begin{aligned}\mathbf{x}_j^i &:= \mathbf{q}_j(t_i) \\ \mathbf{v}_j^i &:= \dot{\mathbf{q}}_j(t_i)\end{aligned}$$

$$t_i = i \cdot \delta t$$


- Right-hand side is evaluated as follows:

$$\mathbf{a}_j^i := \ddot{\mathbf{q}}_j(t_i) = \sum_{\substack{k=1 \\ k \neq j}}^n \mathbf{F}(\mathbf{q}_k, \mathbf{q}_j)$$

- Time integration with Velocity-Verlet scheme:

$$\begin{aligned}\mathbf{x}_j^{i+1} &:= \mathbf{x}_j^i + \delta t \mathbf{v}_j^i + \frac{1}{2} \delta t^2 \mathbf{a}_j^i \\ \mathbf{v}_j^{i+1} &:= \mathbf{v}_j^i + \frac{\delta t}{2} (\mathbf{a}_j^{i+1} + \mathbf{a}_j^i)\end{aligned}$$

# Gravitation





# Gravitation

- Gravitational force depends on the distance  $\mathbf{r}$  between two point masses  $m$  and  $M$  as

$$\mathbf{F}(\mathbf{r}) = G \frac{mM\mathbf{r}}{|\mathbf{r}|^3}$$

- So the resulting expression for accelerations in gravitational N-body is as follows:

$$m_j \ddot{\mathbf{q}}_j = G \sum_{\substack{k=1 \\ k \neq j}}^n \frac{m_j m_k (\mathbf{q}_k - \mathbf{q}_j)}{|\mathbf{q}_k - \mathbf{q}_j|^3}$$

# Molecular Dynamics

- Force depends on the distance  $\mathbf{r}$  between two point molecules as

$$\mathbf{F}(r) = 24 \varepsilon \frac{\mathbf{r}}{r^2} \left( \frac{2\sigma^{12}}{r^{12}} - \frac{\sigma^6}{r^6} \right)$$

- So the resulting expression for accelerations in molecular N-body is as follows:


$$m_j \ddot{\mathbf{q}}_j = -24 \varepsilon \sum_{\substack{k=1 \\ k \neq j}}^n \frac{\mathbf{q}_k - \mathbf{q}_j}{|\mathbf{q}_k - \mathbf{q}_j|^2} \left( \frac{2\sigma^{12}}{|\mathbf{q}_k - \mathbf{q}_j|^{12}} - \frac{\sigma^6}{|\mathbf{q}_k - \mathbf{q}_j|^6} \right)$$

# Molecular Dynamics

This is just  
one guy

- The force is given by Lennard-Jones potential:

$$\mathbf{F}(\mathbf{r}) = -\text{grad } U(r) = -\text{grad } 4\epsilon \left( \frac{\sigma^{12}}{r^{12}} - \frac{\sigma^6}{r^6} \right)$$



- The potential vanishes quickly with growth of  $r$ , so you can truncate it at a certain cut-off distance  $r_{cut}$  (usually  $2.5 \cdot \sigma$ )
- To avoid a jump discontinuity implying an infinite force the potential must be shifted:

$$U(r) = \begin{cases} U(r) - U(r_{cut}), & \text{for } r < r_{cut} \\ 0, & \text{for } r \geq r_{cut} \end{cases}$$



# Project

High Performance Computing for Science and Engineering

October 1, 2013

**CSElab**

Computational Science & Engineering Laboratory

<http://www.cse-lab.ethz.ch>

**ETH**

Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

# Project 1

1. A-priori performance analysis
2. N-Body problem on CPU
3. N-Body problem on GPU
4. Accelerating the N-body problem
5. CPU and GPU comparisons

# A-Priori Performance Analysis

- Why?
  - Estimate achievable performance
  - Understand the algorithm
  - Estimate hardware bottlenecks
- Initial step towards Roofline model



# A-Priori Performance Analysis

- Quantities that determine performance
  - Total amount  $T$  of memory transfers (read/writes)
  - Total number  $C$  of floating point operations (FLOP)
  - Total amount  $M$  of memory required to store the data
- Example kernel:  $K(\mathbf{x}, \alpha) \rightarrow x_i := \alpha x_i \mid i = 0, \dots, n - 1$ 
  - $T = 2n$
  - $C = n$
  - $M = n+1$
- State any assumption you make!

# N-Body on CPU

- Starting point for GPU code
  - Simplest base implementation
  - Understand the algorithm
  - Baseline to verify and debug GPU code
- Software
  - C/C++
  - Design to use different forces
  - Design to accomodate GPU code (classes, templates,...)
  - Structure for testing

# N-Body on CPU

- Code from scratch
  - Be careful and plan ahead!
  - [cppreference.com](http://cppreference.com) and [cplusplus.com](http://cplusplus.com) are your friends!
  - Object oriented design will facilitate
    - Use of different forces (gravitation, Lennard-Jones)
    - GPU implementation
- We provide:
  - `Timer.h` - class for timing
  - `ArgumentParser.h` - class for reading command line arguments



# Box-Muller

Random numbers from a normal distribution

- Required for initial conditions of velocities in MD
- Use `drand48()` to generate  $U_1, U_2$  (uniformly distributed)
- Box-Muller returns two uncorrelated normally distributed values

$$Z_0 = \sqrt{-2 \ln U_1} \cos(2\pi U_2)$$

$$Z_1 = \sqrt{-2 \ln U_1} \sin(2\pi U_2)$$

# N-Body on GPU

- CUDA N-Body Problem
  1. Port C++ code to GPU with CUDA
  2. Optimization
- Keep various optimizations and baseline
  - Easier debugging between optimization steps
  - Easier measurements of performance changes
- CUDA
  - We will introduce it next week
  - CUDA SDK is full of examples

# Algorithm Optimization

- Lennard-Jones potential
  - Short range interactions
  - ➔ Cell lists: reduce complexity from  $O(N^2)$  to  $O(N)$
- Cell lists
  - Will be introduced in two weeks



# CPU vs GPU

- Performance comparison
  - “GPUs are 100x faster than CPUs” - **INCOMPLETE INFO**
  - CPU model, #threads, compiler and version, memory
  - GPU model, grid structure, memory
- Accuracy comparison
  - CPUs and GPUs might have different standards for numerics
  - Non-associativity in parallel computing causes different results
  - Useful to detect bugs as well

# Report

- Hand-in
  - Code
  - PDF with results and comments
  - Optional: movie of the simulation (VMD, OpenGL,...)
- Rule: feedback is proportional to the effort in reporting
  - Even if something does not work, tell us about it!