

Exercise 1 - Feedback

High Performance Computing for Science and Engineering

October 10, 2014

Reporting Performance

- Given a diffusion solver using finite differences:
 - A takes 5s per timestep
 - B takes 1.8s per timestep
- Which one would you choose?

Reporting Performance

- Given a diffusion solver using finite differences:
 - A takes 5s per timestep with a single thread @ 233MHz
 - B takes 1.8s per timestep with multithreading @ 3GHz
- Which one would you choose?

Reporting Performance

- Context is everything!
 - What is the hardware setting?
 - CPU model, number of cores, memory bandwidth, GPU model, network,...
 - What about the software?
 - compiler, version, optimization flags
 - precision used: double/float/half?

Reporting Results

- Write the parameters used in the simulations:
 - it helps the reader reproduce the results
 - it helps you relearn the code when you get back at it
 - it helps debugging: not all values are equally good!

There is no bug free code

- Don't be afraid of showing wrong results!
- Discuss what is not working and show it to us!
 - It might help in looking for bugs
 - It might help you finding out what to try next or see if you forgot something

Exercise 3

High Performance Computing for Science and Engineering


October 10, 2014

Von Neumann Stability Analysis

Used to understand stability of finite difference schemes

Assumes a solution (based on Fourier decomposition):

1D $u_j^n = \rho^n e^{ikx_j}$



time

space

2D $u_{r,s}^n = \rho^n e^{ik_x x_r} e^{ik_y y_s}$

Bound if $\rho < 1$

Substitute assumption into FD scheme and find the condition for which the solution is bound at any time n

Performance evaluation of parallel algorithms

Execution time on p processors: $T(p)$

- $T(1)$ is the best time on a single CPU core
- $p T(p) \geq T(1)$, as otherwise running the parallel program on one core will give shorter time than $T(1) \Rightarrow$ contradiction

Speed up: $S(p) = T(1) / T(p)$

- clearly $S(p) \leq p$
- $S(p) < 1$ is possible, means computation on p cores is slower than on a single core. While not advisable we might be forced to still use such a parallel code for memory reasons.

Strong scaling

Strong scaling: defines how the solution time varies with increasing number of processors p for a **fixed total problem size** (i.e. fix workload is split among cores):

► **Speedup** for strong scaling:

$$S(p) = \frac{T(1)}{T(p)}$$

► **Efficiency** for strong scaling:

$$E_s(p) = \frac{S(p)}{p}$$


$T(1)$ = time of one thread to process the data

$T(p)$ = time of p threads to process the same data


Example:

System with fix problem size N (e.g. # particles)

Problem size:	Execution time:	Speedup:
$N = 100$	$T(1) = 12.0$ s	$S(1) = 1.0$
$N = 100$	$T(2) = 6.0$ s	$S(2) = 2.0$
$N = 100$	$T(4) = 4.0$ s	$S(4) = 3.0$
$N = 100$	$T(8) = 3.0$ s	$S(8) = 4.0$



Split of the fixed problem size over a) 1 core, b) 2 cores, c) 4 cores



Weak scaling

Weak scaling: defines how the solution time varies with the number of processors p for a fixed **problem size per processor**

► Speedup for weak scaling doesn't make sense

► **Efficiency** for weak scaling:

$$E_w(p) = \frac{T(1)}{T(p)}$$

$T(1)$ = time of one thread to process the data

$T(p)$ = time of p threads to process p times the data

Example:

Problem size:

$N = 100$
 $N = 200$
 $N = 400$
 $N = 800$





Execution time:

$T(1) = 12$ s
 $T(2) = 15$ s
 $T(4) = 20$ s
 $T(8) = 24$ s



Efficiency:

$E(1) = 1.0$
 $E(2) = 0.8$
 $E(4) = 0.6$
 $E(8) = 0.5$



Report

- Hand-in
 - Code
 - PDF with results and comments
 - Optional: movie of the simulation (VMD, OpenGL,...)
- Rule: feedback is proportional to the effort in reporting
 - Even if something does not work, tell us about it!